

Chromosome2026 Technical Record

Physics, Numerical Methods, Implementation, Validation, and
Development History

Keheng Zhu and project contributors

Release baseline documented: August 2026

Abstract

This document is the engineering and numerical record for Chromosphere2026. It describes the current `model_column` release, the physical and numerical models from which it evolved, the validation evidence supporting the implementation, the approaches that were rejected or retired, and the limitations that remain. It is not a journal article and it is not a chronological change log.

The active release is a one-dimensional field-aligned hydrogen hydrodynamic column with a common velocity and temperature, equilibrium Saha ionization, conserved hydrogen ionization energy, and a density- and temperature-dependent acoustic exponent Γ_1 supplied by a validated CRASH-derived table. The atmosphere is initialized from the Model C7 temperature profile using shape-preserving cubic interpolation and is reconstructed in Saha-consistent hydrostatic equilibrium. Hydrodynamics uses a nonuniform finite-volume MUSCL–Hancock method that reconstructs $(\log \rho, v, \log p)$ and advances an SWMF-style exact-Riemann Godunov flux evaluated at frozen composition. It is reference-free: no frozen hydrostatic state is stored or subtracted. Thermal conduction is nonlinear and implicit. Static local refinement resolves the upper transition-region structure; the mesh-scaled artificial conductivity is implemented but disabled in the release configuration, which uses physical conduction only. Every stored physical scalar — the conserved state, the mesh metrics, the physical constants, the Courant number, and the simulation time — is double precision, which is both a measured requirement of the resolved chromosphere and the precision the target coupling framework already uses; single-precision storage survives only as a default-off diagnostic build. The validated runtime supports deterministic OpenMP execution.

The repository also contains two-fluid, kinetic-ionization, radiative, flare-heating, three-temperature, loop, and multidimensional formulations, and a separate experimental two-temperature solver used for a controlled test study of the release’s single-temperature assumption. These remain technically useful, but they are not all part of the active `model_column` release. Status labels are used throughout to prevent historical or optional material from being mistaken for current production behavior.

Contents

Status Legend and Reading Guide	vii
I Project Overview and Release Baseline	1
1 Purpose and Scope	2
1.1 Scientific motivation	2
1.2 Software scope	2
2 Exact Release Baseline	3
2.1 Active calculation	3
2.2 Validated release configuration	4
2.3 Capabilities outside the release path	5
II Physical Models and Their Evolution	7
3 Active Release Governing Equations	8
3.1 Closed field-aligned system	8
3.2 Relation between area and magnetic weighting	9
3.3 Constant-field release specialization	9
4 General Plasma-Neutral Formulation	11
4.1 Two-fluid model	11
4.2 Field-aligned reduction and electron folding	11
4.3 Ion-neutral drag and thermal equilibration	12
5 Ionization and Radiation Models	13
5.1 Finite-rate hydrogen network	13
5.2 Radiative closures	13
6 Equilibrium-Ionization Thermodynamics	14
6.1 Saha composition	14
6.2 Caloric equation of state	14
6.3 γ_E and Γ_1	15
6.4 Caloric inversion	15

7 Thermal Conduction and Optional Energy Sources	16
7.1 Physical conductivity	16
7.2 Numerical conductivity	16
7.3 Volumetric heating and cooling stages	16
 III Current model_column Implementation	 18
8 Geometry, Domain, and Initial State	19
8.1 Geometry and gravity	19
8.2 Model C7 temperature data	19
8.3 Saha-consistent hydrostatic reintegration	19
9 Boundary Conditions	21
9.1 Two ghost cells	21
9.2 Release boundary summary	21
9.3 Lower truncation of the modeled domain	22
9.4 Upper hydrodynamic reservoir	23
9.5 Hydrodynamic and conductive temperatures	23
9.6 Boundary-condition development history	24
 IV Numerical Method	 27
10 Conserved and Primitive States	28
10.1 Authoritative release state	28
10.2 Storage precision of the conserved state	28
10.2.1 Measured necessity: the sub-ULP continuity increment	29
10.2.2 Why a <code>Vec</code> -only promotion silently did nothing	29
10.2.3 Consistency with the coupling target	30
10.2.4 What the precision cutover does and does not fix	31
10.2.5 Post-cutover validation	32
10.2.6 Would the coupling target hit this?	32
10.3 Retired conservative equilibrium projection	32
11 Finite-Volume Hydrodynamics	34
11.1 Nonuniform control volumes	34
11.2 Equilibrium-manifold MUSCL reconstruction	34
11.2.1 The MUSCL–Hancock step	35
11.3 Numerical flux	35
11.3.1 Reference flux: the mixture Roe characteristic linearization	36
11.3.2 Release flux: the SWMF-style exact-Riemann Godunov flux at frozen composition	37
11.3.3 Thermodynamic interpretation: frozen composition versus instantaneous Saha equilibrium	38
11.3.4 Remaining limitations of the release flux	39
11.4 Gravity and reference-free hydrostatic balance	40

12 Implicit Conduction and Time Integration	43
12.1 Nonlinear backward-Euler conduction	43
12.2 Timestep control	43
12.3 Active release update path	44
12.4 Legacy/full two-fluid update path	45
13 Static Refinement and Inactive Artificial Diffusion	46
13.1 Mesh construction	46
13.2 Historical local numerical-conduction scaling	46
V Implementation and Performance	47
14 Code Architecture	48
14.1 Primary modules	48
14.2 Runtime contracts	50
15 Performance Engineering	51
15.1 Serial bottlenecks and method-preserving optimization	51
15.2 Deterministic OpenMP	51
15.3 Long-run controls	52
VI Verification and Numerical Results	53
16 Verification by Numerical Risk	54
16.1 Thermodynamic consistency	54
16.2 State and closure semantics	54
16.3 Hydrostatic preservation and boundaries	55
16.4 Conduction conservation and nonlinear convergence	55
16.5 Irregular-grid consistency	55
16.6 Serial/OpenMP equivalence	56
16.7 CFL and duration validation	56
17 Development Experiments and Their Disposition	58
17.1 Constant- γ baselines	58
17.2 Kinetic-ionization experiments	58
17.3 Transition to Saha/ Γ_1	58
17.4 PCHIP correction	59
17.5 Lower split conduction-to-hydro artifact	59
17.6 Upper-boundary diagnosis	59
17.7 Face-flux diagnosis	59
17.8 Uniform resolution scan	60
17.9 Static outer refinement	60
17.10 Face-local numerical conduction	60
17.11 Historical N=1000/2000/4000 convergence evidence	60

18 Experimental Two-Temperature Test Study (model_column_2t)	62
18.1 Motivation and the question being asked	62
18.2 Governing equations	63
18.3 Thermodynamic closure	64
18.4 Acoustics: the release flux carries over unchanged	64
18.5 Separate heat conduction and collisional exchange	65
18.6 Boundary conditions from the characteristic budget	66
18.6.1 The hyperbolic budget: both temperatures must free-float at the top	66
18.6.2 The parabolic budget: one condition per channel per face	66
18.6.3 The controlled comparison	67
18.7 Numerics and implementation	67
18.8 Deliberately absent physics and inherited caveats	69
18.9 Validation status	70
18.10 Results	70
18.10.1 Run integrity and the measured CFL equivalence	70
18.10.2 Exact degeneracy of the closure at $T_e = T_i$	71
18.10.3 The measured decoupling: small, and concentrated in a startup transient	71
18.10.4 Why the temperatures lock: timescale separation	72
18.10.5 Two-temperature against the matched release run	72
18.10.6 The outer heavy-channel condition is immaterial — a robustness result	73
18.10.7 What this validates, and what it does not	73
 VII Current Status and Roadmap	 74
19 Trust Boundaries	75
19.1 Scientifically trustworthy components	75
19.2 Numerically validated but configuration-specific results	75
19.3 Provisional scientific results	76
20 Remaining Limitations and Priorities	77
20.1 Outer conductive boundary	77
20.2 Radiative balance	77
20.3 Equilibrium versus non-equilibrium ionization	77
20.4 Per-step Saha re-equilibration and the frozen Riemann solve	77
20.5 Unresolved corona and coupling	79
20.6 Numerical precision	79
20.7 Recommended order of future work	80
 VIII Appendices	 81
A General Two-Fluid Equations and Source Stages	82
A.1 Conservative variables	82
A.2 Backward-Euler drag stage	82
A.3 Backward-Euler reaction stage	82
A.4 Historical multidimensional formulation	83

B EOS Table Construction and Format	84
B.1 Generation	84
B.2 Runtime policy	84
C Implementation Status Matrix	85
D Release-Relevant Environment Variables	88
D.1 EOS, domain, and physics	88
D.2 Boundary and mesh	88
D.3 Reference-only overrides and diagnostics	89
D.4 Run and OpenMP controls	89
E Non-Release Two-Fluid Research Environment Variables	91
F Canonical Run Configurations and Diagnostics	93
F.1 Validated low-I/O reference run	93
F.2 Diagnostic output semantics	93
G Test Inventory	95
H Chronological Development Ledger	97
I Retired Scenarios and Archived Formulations	99
I.1 Retired active assumptions	99
I.2 Archived but useful formulations	99

Status Legend and Reading Guide

Label	Meaning
Active release path	Used by the validated gamma-table <code>model_column</code> release configuration described in this document.
Implemented but inactive	Present in the code and testable, but disabled in the reference release calculation.
Experimental	Implemented for controlled studies; validation or physical interpretation is not sufficient for release use.
Historical	Important to the design history or still used as a regression baseline, but superseded as the primary model.
Retired	Deliberately removed from the active path or rejected after testing.
Planned	A documented future capability that is not yet implemented or validated.

Two configurations must be distinguished. The *software defaults* are conservative fallbacks used when environment variables are absent. The *validated release configuration* is the specific 1600–2153 km, outer-refined, gamma-table calculation used for the current convergence, performance, and long-run evidence. A claim about one does not automatically apply to the other.

Part I

Project Overview and Release Baseline

Chapter 1

Purpose and Scope

1.1 Scientific motivation

The chromosphere and transition region mediate mass and energy exchange between the dense lower atmosphere and the corona. Conductive energy arriving from above can be distributed among sensible heating, hydrogen ionization, radiation, mechanical work, and upward enthalpy transport. A lower boundary that fixes thermodynamic variables cannot represent this response dynamically. Chromosphere2026 therefore develops a field-aligned column that can evolve the lower atmosphere and eventually provide time-dependent mass and energy fluxes to a global coronal model.

The project is designed as a future lower-atmosphere component for AWSoM within SWMF. **Planned** The present repository does not yet perform that coupling. Its current role is to establish a thermodynamically consistent and numerically controlled column model whose boundary fluxes can later be exposed to a global calculation.

The initial atmospheric data are based on Model C7 of Avrett and Loeser [2008]. The broader two-fluid lineage follows chromospheric plasma-neutral formulations such as Al Shidi et al. [2019], while the current release deliberately takes the strong-coupling equilibrium limit rather than retaining independent ion and neutral dynamics.

1.2 Software scope

The repository contains three overlapping layers:

1. a separate, non-release seven-row plasma-neutral state representation and source-stage framework;
2. several scenarios that exercise different physical closures and geometries; and
3. the current `model_column` release path, a separate single-fluid solver that evolves the three-variable equilibrium mixture directly.

These are now *three* separate solvers over one shared mesh, with three different conserved-state widths. The release solver advances the three-variable single-fluid state $(\rho, \rho v, E)$ and never touches the seven-row representation; the seven-row representation belongs exclusively to the historical two-fluid/multi-temperature/finite-rate-ionization research path; and a third, experimental solver advances the four-variable two-temperature state $(\rho, \rho v, E, E_e)$ for the controlled test study of chapter 18. Selecting a Γ_1 table selects one of the two Saha solvers and the scenario decides which; every solver rejects another's Grid or state width outright.

Chapter 2

Exact Release Baseline

2.1 Active calculation

**Active release
path**

The release baseline consists of the following coupled choices:

- one-dimensional field-aligned hydrodynamics on a prescribed straight field line;
- constant solar gravity along the vertical column;
- pure hydrogen with a common bulk velocity and common temperature, advanced as the three conserved variables $(\rho, \rho v, E)$;
- double-precision storage of the conserved state and of every other stored physical scalar, under one coherent precision knob (section 10.2);
- equilibrium ionization from the Saha relation;
- pressure and internal energy evaluated consistently from that equilibrium composition;
- explicit inclusion of hydrogen ionization energy in the conserved mixture energy;
- caloric inversion $e_{\text{int}}(\rho, T) \mapsto T$ and tabulated $\Gamma_1(\rho, T)$;
- Model C7 temperature initialization with PCHIP interpolation;
- Saha-consistent hydrostatic pressure and density reintegration;
- finite-volume MUSCL–Hancock time stepping with log-space reconstruction of $(\log \rho, v, \log p)$ and the MC3/Koren limiter, the predictor carrying the same momentum source as the corrector;
- the SWMF-style exact-Riemann Godunov numerical flux at frozen composition ($\gamma = 5/3$, ionization energy carried in a passive offset), with face-local Rusanov fallback;
- reference-free hydrostatic balance (no frozen equilibrium state);
- nonlinear backward-Euler physical thermal conduction;
- static local refinement near the upper transition region;
- no artificial or numerical conduction anywhere in the production operator;
- no TRAC, beam heating, volumetric coronal heating, or radiative cooling stage in the production timestep;
- a fixed lower hydrostatic reservoir and a fixed-pressure upper reservoir;
- separate hydrodynamic and conductive temperature data at the upper boundary;
- deterministic OpenMP execution and validated long-run controls.

The reference runtime table is `data/eos/gamma1_hydrogen_v1.dat`. Loading a Γ_1 table selects the single-fluid release solver, and the `model_column` scenario supplies the production table *unconditionally*; only its path is overridable. `model_column` therefore names exactly one solver and one production model. It refuses `ISO_GAMMA` outright, and the initial-condition builder additionally rejects every historical two-fluid research knob rather than silently accepting it: `ISO_TWO_FLUID`, `ISO_IONIZATION`, `ISO_COOLING`, `ISO_TRAC`, `ISO_CORONA`, `ISO_CHEAT`, `ISO_QFLUX`, `ISO_TBOOST`, and

`ISO_NUMERICAL_DIFFUSIVITY_MULT`. The release also rejects RK4, the explicit-only integrator, finite-rate ionization, and the legacy runtime settings `SINGLE_FLUID` and `ENABLE_TE` before evolution begins. A release user therefore needs none of those knobs, and no environment variable can reroute `model_column` into the historical solver. The historical fixed- γ two-fluid column is reached through the separate `model_gentle` scenario, which never loads a Γ_1 table.

2.2 Validated release configuration

The canonical release configuration is:

Item	Reference value
Domain	1600–2153 km, a 553 km truncated upper-chromosphere/lower-TR column
EOS	<code>data/eos/gamma1_hydrogen_v1.dat</code>
Coarse-equivalent resolution	<code>ISO_NS=500</code>
Static refinement	outer profile, factor 4, fully refined above 2100 km, graded over approximately 2080–2100 km
Actual cell count	661
Fine spacing	approximately 0.276 km at the upper boundary
Hydro top temperature	zero-gradient extrapolation of the live top cell
Conduction top temperature	fixed $T_{\text{face}} = 22000$ K at the physical outer face, a half top-cell width from the top-cell center
Outer pressure	fixed reference back pressure with a one-sided HSE ghost ladder
Outer velocity	live top velocity, capped at $0.1c_s$
Lower boundary	1600 km truncation: isothermal Saha-consistent trapezoidal discrete-HSE reservoir with $v = 0$
Energy physics	physical conduction only; artificial conduction, beam/-coronal heating, radiative cooling, and TRAC are structurally absent from the release operator, not merely disabled
Reconstruction/flux	$(\log \rho, v, \log p)$, MC3/Koren $\beta = 2$, SWMF exact-Riemann Godunov flux at frozen $\gamma = 5/3$
State storage precision	<code>double</code> , via the single scalar type <code>chromosphere::Real</code> (section 10.2); <code>-DCHROMO_STATE_FLOAT32=ON</code> is a diagnostic-only build and is not a production setting
Well balancing	none — reference-free (section 11.4)
Runtime & Lie-split explicit hydrodynamics + implicit conduction; OpenMP launcher defaults to 12 threads	
Validated production CFL	0.50 for this model shape and hardware; solver fallback remains 0.25

Selecting `model_column` supplies these model, grid, and EOS settings as override-preserving scenario defaults; the release has no source stages to configure. The production launcher adds the validated CFL, OpenMP, and low-I/O runtime policy but does not redefine the physical model.

2.3 Capabilities outside the release path

Capability	Status	Disposition
Fixed- γ column	Historical	Retained under the separate <code>model_gentle</code> scenario; $\gamma = 1.05$ was a useful nearly isothermal surrogate but is not the release thermodynamics. <code>model_column</code> rejects <code>ISO_GAMMA</code> .
Independent ion and neutral fluids	Implemented but inactive	Implemented in the separate two-fluid solver; rejected by the release solver.
Finite-rate collisional/photoionization	Implemented but inactive	Implemented for legacy two-fluid studies; replaced by equilibrium Saha ionization in the release path.
Separate electron temperature	Experimental	Two routes, neither in the release. Three-temperature exchange machinery exists in the two-fluid solver, and a separate four-row two-temperature solver ($\rho, \rho v, E, E_e$) is used for the controlled test study of chapter 18, reached only through the <code>model_column_2t</code> scenario. That study measured the release's common-temperature assumption on the canonical column and found it good to 0.2% in the evaporation observables, so the release remains a common-temperature mixture by measurement rather than by assumption; it rejects <code>ENABLE_TE</code> .
Optically thick/thin radiation	Implemented but inactive	Implemented and tested in the two-fluid solver. Removed from the release timestep; <code>model_column</code> rejects <code>ISO_COOLING</code> .
TRAC	Implemented but inactive	Implemented and tested in the two-fluid solver. Removed from the release conduction operator; <code>model_column</code> rejects <code>ISO_TRAC</code> .
Artificial (mesh-scaled) conduction	Historical	Retained inside the two-fluid Stage-D conduction operator for controlled legacy comparison. Removed from the release operator entirely; <code>model_column</code> rejects <code>ISO_NUMERICAL_DIFFUSIVITY_MULT</code> .
Imposed heat-flux boundary	Experimental	Alternative to the fixed-temperature conductive wall; not the reference boundary.
Volumetric coronal heating	Experimental	Implemented for resolved-corona studies in the two-fluid solver; not a release stage.
Nonthermal electron-beam heating	Experimental	Implemented for flare experiments in the two-fluid solver; not a release stage.
Loop, flare, gentle, canopy, and PFSS scenarios	Historical	Valuable scenario-specific work, not evidence for the current release column.

Multidimensional two-
fluid MHD formulation

Historical

Preserved as theoretical and architectural
background; not evolved by the current one-
dimensional release executable.

AWSOM/SWMF coupling

Planned

Intended endpoint; no production coupling in-
terface is yet validated.

Part II

Physical Models and Their Evolution

Chapter 3

Active Release Governing Equations

3.1 Closed field-aligned system

Active release path

The release is a one-dimensional hydrodynamic model on a prescribed magnetic field line. The coordinate s follows the field, and magnetic-flux conservation relates the prescribed field strength $B(s)$ to the flux-tube area $A(s)$:

$$A(s)B(s) = \text{const.} \quad (3.1)$$

The magnetic field is geometry, not an evolved component of the state. The release therefore has no induction equation or dynamically evolved magnetic-pressure or Lorentz-force term.

For a time-independent prescribed field and gravitational potential $\Phi(s)$, the continuum equations advanced by the release path are

$$\frac{\partial \rho}{\partial t} + B \frac{\partial}{\partial s} \left(\frac{\rho v}{B} \right) = 0, \quad (3.2)$$

$$\frac{\partial(\rho v)}{\partial t} + B \frac{\partial}{\partial s} \left(\frac{\rho v^2 + p}{B} \right) = pB \frac{\partial B^{-1}}{\partial s} - \rho \frac{\partial \Phi}{\partial s}, \quad (3.3)$$

$$\frac{\partial \mathcal{E}}{\partial t} + B \frac{\partial}{\partial s} \left[\frac{v(\mathcal{E} + p) + q_c}{B} \right] = 0. \quad (3.4)$$

Here

$$\mathcal{E} = e_{\text{int}} + \frac{1}{2} \rho v^2 + \rho \Phi, \quad (3.5)$$

$$q_c = -\kappa_{\text{phys}}(\rho, T) \frac{\partial T}{\partial s}, \quad \kappa_{\text{phys}} = \kappa_e + \kappa_n. \quad (3.6)$$

The active release has no volumetric heating, radiative cooling, TRAC modification, or artificial conductivity. Its energy input is the physical conductive flux established by the external temperature at the upper physical face.

The equilibrium-mixture closure completes the system. With $n_H = \rho/m_H$, proton and electron fractions $x = n_p/n_H = n_e/n_H$, and hydrogen ionization energy χ_H ,

$$\frac{x^2}{1-x} = \frac{1}{n_H} \left(\frac{2\pi m_e k_B T}{h^2} \right)^{3/2} \exp \left(-\frac{\chi_H}{k_B T} \right), \quad (3.7)$$

$$p = (1+x)n_H k_B T, \quad (3.8)$$

$$e_{\text{int}} = \frac{3}{2} p + x n_H \chi_H. \quad (3.9)$$

The equilibrium acoustic closure of the equation of state is

$$c_{\text{eq}}^2 = \Gamma_1(\rho, T) \frac{p}{\rho}, \quad (3.10)$$

where the production Γ_1 table supplies the equilibrium adiabatic derivative. This is the acoustic response *when the ionization balance re-equilibrates within the wave*. The release numerical flux and the release timestep instead use the frozen-composition response $c_{\text{fr}}^2 = (5/3) p/\rho$, on the physical grounds set out in section 11.3.3; both speeds follow from the same closure, and Eq. (3.10) remains the equilibrium derivative of the model. Equations (3.1)–(3.10), together with the prescribed $B(s)$, $\Phi(s)$, initial state, and boundary data, are the self-contained mathematical model. The later thermodynamic, boundary, numerical-flux, and implicit-conduction chapters describe how this closed system is evaluated and discretized.

3.2 Relation between area and magnetic weighting

Using $A \propto B^{-1}$, the continuity operator has the equivalent forms

$$\frac{1}{A} \frac{\partial}{\partial s} (A \rho v) = B \frac{\partial}{\partial s} \left(\frac{\rho v}{B} \right). \quad (3.11)$$

The area form of field-aligned momentum is

$$\frac{\partial(\rho v)}{\partial t} + B \frac{\partial}{\partial s} \left(\frac{\rho v^2}{B} \right) + \frac{\partial p}{\partial s} = -\rho \frac{\partial \Phi}{\partial s}. \quad (3.12)$$

Combining the advective and pressure terms inside one B -weighted flux introduces $pB\partial_s(B^{-1})$. Moving that contribution to the right gives exactly Eq. (3.3). The term is a pressure-area source caused by prescribed tube expansion, not a magnetic force from an evolved field.

The conductive part of Eq. (3.4) is likewise geometry aware:

$$-B \frac{\partial}{\partial s} \left(\frac{q_c}{B} \right) = B \frac{\partial}{\partial s} \left(\frac{\kappa_{\text{phys}}}{B} \frac{\partial T}{\partial s} \right). \quad (3.13)$$

This is the continuum operator represented by the face conductivities and $B_i/B_{i\pm 1/2}$ factors in the nonlinear conduction solve.

No separate gravitational source belongs in Eq. (3.4). If $\mathcal{E}_g = e_{\text{int}} + \rho v^2/2$, its equation contains gravitational work $-\rho v \partial_s \Phi$. Multiplying Eq. (3.2) by the time-independent Φ gives

$$\frac{\partial(\rho \Phi)}{\partial t} + B \frac{\partial}{\partial s} \left(\frac{\rho v \Phi}{B} \right) = \rho v \frac{\partial \Phi}{\partial s}. \quad (3.14)$$

Adding the two equations cancels gravitational work and produces the conservative total-energy equation written above.

3.3 Constant-field release specialization

Active release
path

The reported `model_column` calculation sets

$$B_i = B_{i\pm 1/2} = 1, \quad \frac{\partial B^{-1}}{\partial s} = 0. \quad (3.15)$$

The tube area is therefore constant, all B -metric factors reduce to unity, and Eqs. (3.2)–(3.4) become the Cartesian one-dimensional mass, momentum, and total-energy equations. The finite-volume state and the Roe reference eigensystem later abbreviate $\mathcal{E}$ as E ; they discretize this specialization of the same closed field-aligned model.

Chapter 4

General Plasma-Neutral Formulation

4.1 Two-fluid model

Historical The project began from a reacting plasma-neutral description in which charged hydrogen and neutral hydrogen have independent mass, momentum, and energy equations. In compact one-dimensional notation,

$$\frac{\partial \rho_s}{\partial t} + \frac{1}{A} \frac{\partial}{\partial s} (A \rho_s u_s) = S_{\rho,s}, \quad (4.1)$$

$$\frac{\partial(\rho_s u_s)}{\partial t} + \frac{1}{A} \frac{\partial}{\partial s} [A(\rho_s u_s^2 + p_s)] = \rho_s g_s + S_{m,s}, \quad (4.2)$$

$$\frac{\partial E_s}{\partial t} + \frac{1}{A} \frac{\partial}{\partial s} [A(E_s + p_s)u_s] = \rho_s u_s g_s + S_{E,s}, \quad (4.3)$$

where $s \in \{i, n\}$, $A \propto B^{-1}$ is the flux-tube area, and the source terms include ion-neutral drag, thermal equilibration, ionization, recombination, conduction, radiation, and imposed heating. Related source-consistent formulations are discussed by Meier and Shumlak [2012] and Khomenko et al. [2014].

The full model remains useful when charge-neutral drift, reaction time scales, or independent temperatures are the scientific target. It is not used by the current gamma-table release because that release assumes the strong-coupling, equilibrium-ionization limit.

4.2 Field-aligned reduction and electron folding

Historical The general formulation was reduced from multidimensional MHD to a prescribed field line by evolving parallel flow and treating the field strength as prescribed geometry. Quasi-neutrality gives $n_e = n_i$ for hydrogen. Electron mass is neglected in the total mass and momentum, while electron pressure and thermal conduction remain in the charged-fluid energetics.

In the common-temperature limit, the charged pressure is

$$p_i = (n_p + n_e)k_B T = 2n_i k_B T, \quad (4.4)$$

and the neutral pressure is $p_n = n_n k_B T$. The release mixture pressure is their sum. The field-aligned reduction remains conceptually important for later AWSOM coupling, but the active release uses a straight, constant-area column.

4.3 Ion-neutral drag and thermal equilibration

**Implemented but
inactive**

The legacy integrator advances drag through center-of-mass and relative-drift variables. This avoids an explicit collision-time restriction and preserves total momentum. The removed relative kinetic energy is returned to thermal energy. A separate point-implicit temperature stage equilibrates ion, neutral, and optionally electron temperatures.

These stages do not exist in the release solver. Their infinite-coupling limit is instead built into the release state itself: there is one momentum and one internal energy per cell, so a common velocity and a common temperature hold identically at every stage rather than being enforced after the fact.

Chapter 5

Ionization and Radiation Models

5.1 Finite-rate hydrogen network

**Implemented but
inactive**

The finite-rate path contains electron-impact ionization, radiative and three-body recombination, a multilevel collisional-radiative channel, and a frozen radiation-field photoionization profile. Representative rates draw on Voronov [1997], Hummer [1994], and the non-equilibrium chromospheric literature [Carlsson and Stein, 2002, Leenaarts et al., 2007].

The reaction stage solves the local backward-Euler composition equation rather than taking an explicit chemistry step. Ionization energy is removed from or returned to the thermal pool according to the reaction channel. This machinery was important in exposing the sensitivity of the upper boundary to energy balance. It is not part of the release solver, whose composition is the Saha equilibrium value of the EOS closure at every instant.

5.2 Radiative closures

**Implemented but
inactive**

Two radiative closures are implemented:

1. an optically thick chromospheric loss based on tabulated H I, Ca II, and Mg II escape-probability approximations following Carlsson and Leenaarts [2012]; and
2. an optically thin transition-region/coronal loss $Q_{\text{thin}} = n_e n_H \Lambda(T)$, smoothly activated near 2×10^4 K.

TRAC broadening modifies electron conductivity and optically thin losses below an adaptive cut-off while preserving their product, following Johnston et al. [2019, 2020]. Both closures and TRAC live in the two-fluid research solver; neither is a stage of the release timestep, and `model_column` rejects `ISO_COOLING` and `ISO_TRAC`. Their absence from the release results is therefore a physical limitation of those results, not a missing code capability.

Chapter 6

Equilibrium-Ionization Thermodynamics

6.1 Saha composition

**Active release
path**

The current physical core is a classical pure-hydrogen equilibrium mixture. With total hydrogen-nucleus density $n_H = \rho/m_H$ and ionization fraction $x = n_p/n_H = n_e/n_H$,

$$\frac{x^2}{1-x} = \frac{1}{n_H} \left(\frac{2\pi m_e k_B T}{h^2} \right)^{3/2} \exp \left(-\frac{\chi_H}{k_B T} \right). \quad (6.1)$$

The implementation evaluates the positive root in logarithmic form to avoid overflow, underflow, and cancellation in the neutral and ionized limits.

All physical thermodynamic quantities use the unclamped solution x_{Saha} . A small carrier fraction floor is applied only when packing the historical ion and neutral storage rows. Changing that floor changes the trace row but not pressure, temperature, conductivity, radiation, or sound speed.

6.2 Caloric equation of state

**Active release
path**

The mixture pressure and internal-energy density are

$$p(\rho, T) = (1 + x_{\text{Saha}}) n_H k_B T, \quad (6.2)$$

$$e_{\text{int}}(\rho, T) = \frac{3}{2} p + x_{\text{Saha}} n_H \chi_H. \quad (6.3)$$

The second term records the energy stored in ionizing hydrogen. In the ionization region, conductive input can therefore change x substantially without producing the temperature rise predicted by a constant- γ gas.

The derivative used by the nonlinear conduction solve is

$$C_V^{\text{eq}} = \left(\frac{\partial e_{\text{int}}}{\partial T} \right)_{\rho} = \frac{3}{2} (1 + x) n_H k_B + \left(\frac{3}{2} n_H k_B T + n_H \chi_H \right) \left(\frac{\partial x}{\partial T} \right)_{\rho}, \quad (6.4)$$

with

$$\left(\frac{\partial x}{\partial T} \right)_{\rho} = \frac{x(1-x)}{(2-x)T} \left(\frac{3}{2} + \frac{\chi_H}{k_B T} \right). \quad (6.5)$$

This effective heat capacity is positive over the production table and becomes large in the partial-ionization region.

6.3 γ_E and Γ_1

Two exponents have different roles and must not be conflated.

Caloric ratio $\gamma_E = 1 + p/e_{\text{int}}$ describes how the chosen zero of internal energy affects the ratio of pressure to stored energy. It enters interpretation and diagnostics, not the acoustic eigenvalue.

First adiabatic exponent $\Gamma_1 = (\partial \ln p / \partial \ln \rho)_S$ controls the equilibrium-mixture sound speed,

$$c_{\text{eq}}^2 = \Gamma_1 \frac{p}{\rho}. \quad (6.6)$$

Frozen translational exponent $\gamma = 5/3$, the monatomic value, controls the sound speed when the ionization state cannot change on the timescale of the disturbance, $c_{\text{fr}}^2 = (5/3)p/\rho$. It follows from the exact split $e_{\text{int}} = p/(5/3 - 1) + e_{\text{ion}}$ with e_{ion} held fixed.

Active release path

Γ_1 and $\gamma = 5/3$ are the equilibrium and frozen limits of the same closure, not competing approximations to one number. The release face Riemann problem and the release timestep use the frozen limit; Γ_1 remains the tabulated equilibrium derivative and is used for diagnostics, the outer-boundary Mach cap, and the reference Roe flux. Section 11.3.3 gives the physical argument and the literature; section 20.4 records the unresolved consistency question that follows from it.

Active release path

The runtime $\Gamma_1(T, n_{\text{H}})$ is bilinearly interpolated in $(\log T, \log n_{\text{H}})$ from the CRASH-derived table. The caloric closure itself is analytic. This division was deliberate: CRASH supplies the validated acoustic derivative, while pressure, energy, Saha composition, and caloric inversion remain transparent and independently testable.

An independent analytic identity used in validation is

$$\Gamma_1 = \chi_\rho + \frac{p\chi_T^2}{TC_V^{\text{eq}}}, \quad \chi_\rho = \left(\frac{\partial \ln p}{\partial \ln \rho} \right)_T, \quad \chi_T = \left(\frac{\partial \ln p}{\partial \ln T} \right)_\rho. \quad (6.7)$$

6.4 Caloric inversion

Active release path

Conserved evolution provides ρ and e_{int} , so every decode may require the solution of

$$e_{\text{int}}(\rho, T) = e_{\text{int}}^{\text{target}}. \quad (6.8)$$

The implementation brackets the root on the EOS-table temperature interval, attempts safeguarded Newton updates using C_V^{eq} , and falls back to bisection when a Newton step is unsafe or insufficiently reduces the bracket. Out-of-range density or energy is a production error, not a silent clamp. A finite previous temperature is only an accelerator; correctness does not depend on a cache.

Chapter 7

Thermal Conduction and Optional Energy Sources

7.1 Physical conductivity

Active release path

The common-temperature conductive flux is

$$q_{\text{phys}} = -(\kappa_e + \kappa_n) \frac{\partial T}{\partial s}. \quad (7.1)$$

The electron coefficient approaches Spitzer behavior in the ionized limit and includes electron-neutral collisions in weakly ionized gas [Spitzer, 1962, Vranjes and Krstic, 2013]. The neutral coefficient dominates much of the lower chromosphere. Physical diagnostic plots use only $\kappa_e + \kappa_n$; they explicitly exclude TRAC and artificial diffusion.

7.2 Numerical conductivity

Historical A stabilizing artificial temperature diffusivity was once folded into the conduction operator:

$$\chi_{\text{num},f} = C_{\text{num}} \Delta s_f, \quad (7.2)$$

$$K_{\text{num},f} = \chi_{\text{num},f} C_{V,f}^{\text{eq}}, \quad (7.3)$$

with the reference value $C_{\text{num}} = 2000 \text{ m s}^{-1}$ before an explicit diagnostic multiplier, and Δs_f the actual center-to-center face spacing. This term is artificial and was never reported as physical conductivity.

It has been removed from the release conduction operator entirely: the release face conductivity is $\kappa_e + \kappa_n$ and there is no $K_{\text{num},f}$ term to disable. The release is therefore physical-conduction-only structurally, not because a multiplier happens to be zero, and `model_column` rejects `ISO_NUMERICAL_DIFFUSIVITY_MULT`. Equation (7.3) remains implemented inside the two-fluid Stage-D operator so that the historical convergence studies recorded in this document can be reproduced.

7.3 Volumetric heating and cooling stages

Implemented but inactive

Beam heating, volumetric coronal heating, and radiative cooling are implemented as post-conduction source stages of the *two-fluid* solver. Each modifies only the thermal pool of a cell. They are not stages of the release timestep, which consists of the explicit hydrodynamic update and the implicit physical conduction solve and nothing else; `model_column` rejects `ISO_COOLING`, `ISO_CHEAT`, and the flare beam configuration.

Experimental

An imposed outer heat flux can replace the fixed-temperature wall, and a resolved-corona option can combine TRAC, radiation, and volumetric heating in the two-fluid solver. These configurations are useful for future physical studies but are not the current release baseline.

Part III

Current model_column Implementation

Chapter 8

Geometry, Domain, and Initial State

8.1 Geometry and gravity

Active release path The release column is vertical, straight, and constant area: $B_i = B_{i\pm 1/2} = 1$ and $\partial_s(B^{-1}) = 0$. The code retains general flux-tube metric factors, but they reduce to unity in the active configuration. Gravitational potential is stored at faces and centered for cell decoding. The acceleration magnitude is 273.95 m s^{-2} .

8.2 Model C7 temperature data

Active release path Model C7 provides the temperature profile and the base total hydrogen density. The complete repository table spans the photosphere through the corona, but the release initializer samples only the requested domain.

The historical linear temperature interpolation was continuous but had a piecewise-constant derivative, which seeded step-like conductive-flux structure at every sparse C7 knot. The gamma-table initializer now uses a Fritsch–Carlson/Fritsch–Butland shape-preserving cubic Hermite interpolant. It passes exactly through the C7 knots, remains monotone on monotone intervals, and is C^1 across internal knots. The fixed-gamma scenarios retain their historical interpolation for regression stability.

8.3 Saha-consistent hydrostatic reintegration

Active release path The release does not use the tabulated C7 electron and neutral densities directly away from the base. It retains $T(h)$, anchors the base with the C7 total hydrogen density, and reintegrates hydrostatic pressure according to the active Saha EOS.

Between cell centers,

$$\ln p_i = \ln p_{i-1} - \frac{\Delta h_i}{2} (H_{i-1}^{-1} + H_i^{-1}), \quad H^{-1} = \frac{m_H g}{(1+x)k_B T}. \quad (8.1)$$

Because x depends on the density implied by p , each cell solves a pressure-density fixed point. Non-convergence after 60 iterations is a hard error with cell, height, temperature, and pressure context.

The resulting cell state is packed directly on the equilibrium manifold. No frozen copy is retained: the release keeps no reference equilibrium (section 11.4), and `grid.eq_state` is populated only on the non-release two-fluid leg of this scenario.

Chapter 9

Boundary Conditions

9.1 Two ghost cells

Active release path

Every boundary provides two packed ghost states because the MUSCL reconstruction accesses two exterior layers. Each ghost is packed with the gravitational potential used at that boundary face. The ghosts are not arbitrary extrapolation buffers: they are part of the discrete boundary closure and must satisfy the same EOS as interior face states.

9.2 Release boundary summary

Active release path

The release has two distinct boundary mechanisms:

Hydrodynamic ghost states close the MUSCL reconstruction stencil and contain pressure, hydrodynamic temperature, EOS-derived composition, and packed energy. **Conductive physical-face data** close the nonlinear implicit conduction operator. The fixed 22000 K upper conductive datum is not a hydro ghost temperature, and the second upper hydro ghost has no conductive role.

Quantity	Lower boundary	Upper boundary	BC type / independent?	Numerical implementation
u	Fixed reservoir, $u = 0$	Live top-cell velocity with $ u_g \leq 0.1c_{eq}$	Lower independent; upper extrapolated and capped	Stored in both hydro ghosts used by the MUSCL stencil
p	Two-rung discrete-HSE continuation	First ghost fixed to p_{ref} ; second from one-sided HSE	Upper p_{ref} is the single external hyperbolic datum	Ghost pressure is used with T_{hydro} to close the EOS
T_{hydro}	Fixed reservoir temperature	Zero-gradient from the live top cell	Lower independent; upper extrapolated	Used only to construct hydro ghost density, composition, energy, and Riemann states
ρ	EOS recovery from HSE pressure and fixed T	EOS recovery from ghost pressure and T_{hydro}	Derived at both boundaries	$\rho = \rho(p, T)$ from the equilibrium hydrogen closure

Quantity	Lower boundary	Upper boundary	BC type / independent?	Numerical implementation
x	Saha equilibrium at the ghost (ρ, T)	Saha equilibrium at the ghost (ρ, T_{hydro})	Derived at both boundaries	Recomputed by the equilibrium face-state packer
Energy	Packed from the reservoir thermodynamic state	Packed from each EOS-derived hydro ghost state	Derived; not separately prescribed	Preserves the mixture energy convention, including ionization and gravitational energy
Conductive T	Fixed lower reservoir temperature	$T_{\text{face}} = 22000$ K at the physical outer face	Independent Dirichlet data	Lower state enters the left conduction face; upper state is separate from both hydro ghosts
Conductive heat flux	Determined by the live first cell and reservoir state	Determined by the live top cell and physical-face state	Derived, not independently imposed	$q_{\text{top}} = \kappa_{\text{face}}(T_{\text{face}} - T_{\text{top}})/(\Delta s_{\text{top}}/2)$

9.3 Lower truncation of the modeled domain

Active release path

The lower boundary at 1600 km is not the photosphere. It is a deliberate truncation inside the chromosphere, placed at the height above which this solver’s Saha/LTE equilibrium closure is intended to be used. The atmosphere below it is real and is simply outside the model. Three things must therefore be kept apart:

Physical boundary information at 1600 km. On the timescales resolved here, the chromosphere below the truncation acts as a quasi-static stratified reservoir. That is the entire physical content of the condition: it fixes a reservoir temperature T_0 , a hydrostatic pressure stratification, and $v = 0$.

Numerical ghost closure below it. MUSCL reconstruction needs two exterior layers, so two ghost states are packed one and two base-cell widths below cell 0. They are a stencil closure. They are *not* an assertion that LTE, or this model, describes the atmosphere at those heights, and no physical meaning should be read into their densities or ionization fractions.

Dependence on a global equilibrium. There is none. The reservoir reference is captured once from cell 0 of the initial condition and is a local base-cell quantity; the release does not hold a global hydrostatic profile (section 11.4).

Each ghost pressure is a trapezoidal, EOS-closed hydrostatic step taken downward from the rung above,

$$p_{k+1} = p_k + \Delta s_0 \frac{1}{2}(\rho_k + \rho_{k+1})g, \quad \rho_{k+1} = \rho_{\text{EOS}}(p_{k+1}, T_0), \quad (9.1)$$

iterated to a fixed point, with Δs_0 the actual base-cell width (on the refined mesh this is the fine spacing, not the coarse one). Ghost densities come from the equilibrium EOS rather than being copied flat, which keeps the pressure and logarithmic-density slopes mutually consistent and removes the Rusanov mass leak of the older flat-density boundary.

Retired Equation (9.1) replaces the earlier first-order ladder $p_{g,1} = p_0 + \rho_0 g \Delta s_0$, $p_{g,2} = p_0 + 2\rho_0 g \Delta s_0$, which used the cell-0 density for both rungs. That form carried an $\mathcal{O}((\Delta s_0/H)^2)$ pressure error and hence an $\mathcal{O}(\Delta s_0/H)$ error in the hydrostatic force balance of the lowest interior cells, and it was inconsistent with the trapezoidal ladder the outer face had already been using. With the predictor source of section 11.2.1 in place but the old ladder still present, the reference-free 20 s residual velocity peaked at 1607.7 km — in the first few interior cells. Equation (9.1) removed that layer: the base-band $\max |v|$ fell from 3.2×10^{-2} to 1.4×10^{-2} m s^{−1}, the global maximum moved

to the upper transition region, and the $t = 0$ face mass-flux defect fell by a further factor of 11.

Historical Because 1600 km is an artificial cut rather than a wall, a fixed $v = 0$ reservoir is in principle an acoustic mirror, and a gravity-aware characteristic alternative was built and measured. It imposes only the incoming (upward, $u + c$) acoustic invariant $\delta u + \delta p/Z = 0$ about the base reference state with impedance $Z = \rho c$, taking the outgoing invariant from the live base cell, which gives ghost deviations $v_g = -(p_0 - p_{0,\text{ref}})/Z$ and $p_g = p_{g,\text{ref}} - Zv_0$. It degenerates exactly to the fixed reservoir at equilibrium, so it costs nothing in hydrostatic balance. It also gained nothing: over 20 s hydrostatic, 100 s and 1000 s conduction-driven runs it left the $t = 0$ defect and the residual-velocity profile unchanged, altered the evaporation-window velocity by $\leq 0.4\%$, and left the detrended base-region oscillation content unchanged (3 sign changes in 150 s either way; peak-to-peak 0.31 vs 0.31 m s^{-1}). The disturbance the transition region sends down arrives as a slow quasi-static pressure adjustment, not as acoustic ringing, so this face is not an important reflector for this model. The characteristic condition was therefore not promoted, and its code was removed rather than left as a default-off switch. **Limitation:** the corollary is that mass cannot cross 1600 km, so the column cannot be replenished from below during evaporation — the characteristic variant recovered about 17% of the 1000 s mass loss this way. That is a modeling choice of the truncation, not a numerical defect, and it should be revisited if base mass supply becomes a quantity of interest.

The lower conduction condition is separate. The default is a fixed-temperature reservoir; `ISO_INNER_T_NEUMANN=1` instead imposes zero conductive flux. Either choice leaves the hydrodynamic closure at $v = 0$.

9.4 Upper hydrodynamic reservoir

Active release path

The upper boundary represents an external reservoir through one imposed hydrodynamic datum: a fixed reference back pressure. This is appropriate for subsonic outflow, which has one incoming characteristic.

At initialization, or after an optional relaxation stage, the code captures a reference pressure one one-sided hydrostatic step outside the top cell. That pressure is then held fixed. The second ghost is obtained from another one-sided HSE rung,

$$p_{k+1} = p_k - \frac{\Delta s_N g}{2}(\rho_k + \rho_{k+1}), \quad (9.2)$$

with the density closed by the equilibrium EOS at the hydrodynamic ghost temperature. The local fixed-point iteration is repeated twice, sufficient for the small hydrostatic correction.

The ghost velocity copies the live top-cell velocity but is clipped to

$$|v_g| \leq M_{\text{cap}} c_s, \quad M_{\text{cap}} = 0.1 \quad (9.3)$$

in the reference configuration. During an optional relaxation phase it is fixed to zero.

9.5 Hydrodynamic and conductive temperatures

Active release path

The upper boundary deliberately carries two temperatures:

Hydrodynamic ghost temperature zero-gradient extrapolation from the live top cell when `ISO_HYDRO_T_DECOUPLE=1`; this temperature determines ghost density, internal energy, and the Riemann state.

Conductive physical-face temperature a fixed external value, 22000 K in the release calculation; this appears only in the implicit conduction stencil.

This separation avoids imposing both pressure and temperature on the hydrodynamic subsystem while retaining a controlled thermal reservoir. The conductive Dirichlet datum is located at the physical outer face. For the top cell,

$$d_{\text{boundary}} = \frac{1}{2}\Delta s_{\text{top}}, \quad q_{\text{top}} = \kappa_{\text{face}} \frac{T_{\text{face}} - T_{\text{top}}}{d_{\text{boundary}}}, \quad T_{\text{face}} = 22000 \text{ K}. \quad (9.4)$$

The face conductivity is evaluated from the live top-cell state and an external Saha-equilibrium face state constructed from the fixed back pressure and T_{face} . The hydro solver still reconstructs through two ghost cells, but neither the first- nor second-ghost center defines the release conductive wall.

9.6 Boundary-condition development history

The upper boundary was revised by numerical question rather than by date:

Question	Controlled change	Result	Disposition
Could centered pressure extrapolation preserve HSE?	Anchor the first ghost from the second interior cell across the top cell.	In the non-hydrostatic upper TR it inverted the boundary pressure gradient and produced a standing downward force and last-cell reversal.	Retired

Could a live top-cell pressure anchor be transparent? Use the current top pressure and one-sided HSE outside it. Originally: removed back pressure; the subsonic column drained, accelerated to about 1.9 km s^{-1} , and failed near 12 s. Re-measured in the current release numerics (ISO_OUTER_P_NEUMANN=2), it does *not* fail — 4000 s, no fallbacks, mass drift -1.03% — but the unanchored pressure level drifts $+5.8\%$ and still rises at 4000 s while the evaporation signal falls by $\sim 19\times$ ($V_{\text{top}} = +0.49$ against $+9.50 \text{ m s}^{-1}$). Re-run as a wall-temperature study at 22 and 100 kK it keeps the exponent ($n = 2.80$ against 2.91) and loses everything else: $\sim 20\times$ suppression at both walls, top and base mass fluxes of opposite sign, and at 100 kK -22.7% column mass with the transition region migrated down to 1900 km ($T = 21\,931$ K there, against 6546 K anchored). **Retired**(for measurability, not stability)

Must the outer pressure closure be hydrostatic?	Impose a literal $dp/ds = 0$ at the outer face with the hydro ghost temperature pinned to the conduction wall (ISO_OUTER_T_HYDRO_WALL column mass by 1000 s, ISO_OUTER_P_NEUMANN=1)	Yes, and this is the decisive case: withholding the top cell's hydrostatic support collapses the column inward at the M_{cap} clamp, removes half the column mass by 1000 s, and replaces the chromosphere with a stationary 23.7 kK fully ionized downdraft — while still reporting a clean termination and zero Riemann fallbacks. Pinning the hydro temperature alone is inert here (V_{top} 9.501 vs 9.515 m s ⁻¹), so the entire effect is the pressure closure.	Retired
What supplies the incoming characteristic?	Capture and hold a fixed external reference pressure.	Stable reservoir pressure without gradient inversion or uncontrolled drainage.	Active release path
Was hydro temperature over-specified?	Keep the 22 kK conduction wall but set hydro ghost temperature from the top cell.	Removed the last-cell velocity reversal; the broader interior ripple remained.	Active release path
Was the remaining ripple a flux discontinuity?	Record central, diffusive and total face fluxes against the $t = 0$ baseline.	The effective conserved face flux was smooth; cell-centered ρv and Rusanov components were anti-correlated.	Diagnostic conclusion
Does the ripple converge?	Uniform and local resolution scans.	Ripple/mismatch decreases approximately first order and is efficiently reduced by upper-TR refinement.	Validated numerical interpretation
Was the old ghost-center conductive input converged?	N=1000/2000/4000 with local artificial conductivity.	Mass flux changed only 2% on the last pair, but total conductive input still changed 19%.	Historical Superseded by physical-face, physical-only conduction

Part IV

Numerical Method

Chapter 10

Conserved and Primitive States

10.1 Authoritative release state

Active release path

The release solver stores and advances exactly three conserved rows per cell,

$$U = (\rho, m, E)^\top, \quad m = \rho v, \quad E = e_{\text{int}}(\rho, T) + \frac{m^2}{2\rho} + \rho\Phi, \quad (10.1)$$

in the row order `mix::RHO`, `mix::MOM`, `mix::ENERGY` of `chromosphere.hpp`. This is the discrete form of Eqs. (3.2)–(3.5): nothing else is stored, and every stage of the timestep — reconstruction, numerical flux, geometric source, boundary ghosts, implicit conduction, and the timestep control — reads and writes these three rows directly.

Everything else is a derived thermodynamic quantity of (ρ, e_{int}) produced by the equilibrium closure:

$$x = x_{\text{Saha}}(n_{\text{H}}, T), \quad n_{\text{H}} = \rho/m_{\text{H}}, \quad (10.2)$$

$$n_e = x n_{\text{H}}, \quad n_{\text{HI}} = (1 - x) n_{\text{H}}, \quad (10.3)$$

$$\rho_i = x \rho, \quad \rho_n = (1 - x) \rho, \quad (10.4)$$

$$p = (1 + x) n_{\text{H}} k_B T, \quad p_e = x n_{\text{H}} k_B T. \quad (10.5)$$

They are single-valued functions of the conserved state, they are recomputed wherever needed, and none of them is an independent degree of freedom. In particular there is no ion/neutral mass split, no per-carrier momentum, and no electron energy row in the release state.

Historical Earlier releases stored the same physics in the seven-row carrier vector $(\rho_i, \rho_n, \rho_i v_i, \rho_n v_n, E_i, E_n, E_e)$ inherited from the two-fluid solver, with the ionization fraction clipped to a trace floor $f_{\text{min}} = 10^{-8}$ so that both mass rows stayed representable. That layout carried four redundant degrees of freedom for a three-degree-of-freedom physical system, which is why it required a per-step conservative repacking (below). The seven-row state now belongs solely to the non-release two-fluid solver of Chapter 4.

10.2 Storage precision of the conserved state

Active release path

The release stores U in `double`. Precision is a single knob. `chromosphere::Real` is the solver’s scalar storage type and `Vec` is `arma::Col<Real>`:

```

#ifdef CHROMO_STATE_FLOAT32
typedef float Real;    // DIAGNOSTIC ONLY
#else
typedef double Real;
#endif
typedef arma::Col<Real> Vec;

```

Every stored physical scalar is declared in terms of those two types: the conserved state of Eq. (10.1), the mesh metric arrays, the **Grid** physical constants $(m_i, m_n, m_e, k_B, q_e, \chi_H, g, \mu_0)$, the Courant number, the limiter β , the run controls, and **sim_time**. Nothing in the release path names `float` or `double` directly.

Implemented but inactive

`-DCHROMO_STATE_FLOAT32=ON` reverts `Real` to `float` and reproduces the pre-cutover storage. It exists only so that the round-off sensitivity described below stays reproducible, it is built into a separate tree (`build_omp_f32/`), it is not exercised by the release validation, and no quoted release number may come from it.

Retired The earlier `CHROMO_STATE_FLOAT64` option is gone, and CMake hard-errors on it rather than accepting it silently.

10.2.1 Measured necessity: the sub-ULP continuity increment

Active release path

The conservative update is $U^{n+1} = U^n + \Delta t R(U^n)$ with no compensated summation. It therefore stagnates in a row whenever the increment is smaller than half a unit in the last place (ULP) of the stored value, because the addition rounds back to the operand exactly. In the 4000s $N = 500$ release-shaped run the per-step continuity increment $\Delta t (-\partial_s F^\rho)$ is about $2.6 \times 10^{-17} \text{ kg m}^{-3} \text{ s}^{-1}$, while one single-precision ULP of $\rho \sim 9 \times 10^{-10} \text{ kg m}^{-3}$ is about 5.6×10^{-17} . **619 of 660 cells sat below half a ULP** (median ratio 0.30), so $\rho += \Delta t R_\rho$ returned ρ unchanged every step. The lower chromosphere of that run could not evolve its density at all; the observable symptom was a broad, static height gradient in the conserved face mass flux between about 1600 and 1850 km that never relaxed. This is permanent stagnation, not slow drift.

The controlling dimensionless number is the ratio of the local evolution timescale $\tau = |U|/|dU/dt|$ to the timestep, measured against the reciprocal unit round-off $1/u$ of the storage type. Stagnation occurs when $\tau/\Delta t > 1/u$:

This column	Δt	τ	$\tau/\Delta t$	$1/u$	margin
single precision	$5.6 \times 10^{-3} \text{ s}$	$1.9 \times 10^5 \text{ s}$	3.5×10^7	1.7×10^7	fails by $2 \times$
double (release)	$5.6 \times 10^{-3} \text{ s}$	$1.9 \times 10^5 \text{ s}$	3.5×10^7	9.0×10^{15}	safe by 2.6×10^8

`double` is therefore not a marginal fix but eight orders of magnitude of headroom, which matters for the decision: the choice is robust to a factor of ten in resolution, in run duration, or in how slow the physics of interest turns out to be. **What manufactures the problem is resolving the chromosphere.** A fine mesh and a $\sim 7 \text{ km s}^{-1}$ sound speed give a tiny Δt ; mass-loading a dense stratified layer gives a long τ . Very few solar models have both at once, which is why the failure mode has no upstream precedent (section 10.2.6).

10.2.2 Why a Vec-only promotion silently did nothing

Historical The first attempt at this promoted `arma::Col<float>` to `arma::Col<double>` and nothing else. Six sites in the release path still narrowed through `float`: the ghost/state packing (and

hence the initial condition), the MUSCL–Hancock predictor row, *the right-hand-side row — the increment at the centre of the whole problem*, the per-cell timestep, the packed-energy round entering the caloric domain, and the Grid physical constants $m_i, m_n, m_e, k_B, q_e, \chi_H$ that enter pressure and temperature everywhere.

One narrowing cast anywhere in the update chain reimposes the coarser representation on the entire chain, regardless of what the storage container holds. This was measurable rather than theoretical: with the container widened but the pack still casting through float, a state packed at an equation-of-state table endpoint arrived 4×10^{-8} relative below the solver’s acceptance bracket — exactly one single-precision ULP — which made the table-corner domain regression unrunnable under that option and forced it to be skipped. **The lesson to keep is that a precision knob must be coherent or it does nothing.** That is also how the target framework does it: one PRECISION variable in Makefile.conf, one default real kind, no per-array exceptions.

10.2.3 Consistency with the coupling target

Active release path

The second justification is independent of the first and is the one that matters for how this model is positioned. A standalone audit of the local SWMF tree (SWMF master 2026-03-16, GM/BATSRUS master 2026-03-18), read from source rather than assumed, establishes the following.

Persistent state precision. BATS-R-US declares its state with no kind parameter — `real`, `allocatable`, `target :: State_VGB(:, :, :, :, :)` at GM/BATSRUS/src/ModAdvance.f90:127, and likewise `StateOld_VGB (:131)`, `Source_VC (:151`, which holds $\Delta t(S - \nabla \cdot F)$, the flux arrays `(:204)` and `DtMax_CB (:140)`. Precision is thus entirely a compile-time decision, and that decision is double in *every* shipped configuration: `PRECISION = ${DOUBLEPREC}` with `DOUBLEPREC = -fdefault-real-8 -fdefault-double-8` in `share/build/Makefile.Darwin.gfortran:19--21` and, verified individually, in all 23 `share/build/Makefile.{Darwin, Linux}.*` templates — `gfortran`, `ifort`, `nagfor`, `nvfortran`, `pgf90`, `crayftn`, `xlfort90`, `flang`, `absoft`, `lf95` — *including* the `nvfortran/OpenACC` GPU builds. `share/Scripts/Config.pl:124` sets `$DefaultPrecision = 'double'`. A `-single` switch exists (`Config.pl:168`) but is legacy and unvalidated: the only match for “single” in the 109 kB regression makefile is the phrase “single-ion” in a test description.

Update precision. The conservative update is two statements at the same default real kind: `ModUpdateState.f90:372--380` forms `Source_VC = DtLocal*(Source_VC + (Flux_VXI(i) - Flux_VXI(i+1) + ...)/CellVolume_GB)`, and `:728--731` performs `State_VGB = StateOld_VGB + Source_VC(1:nVar, ...)`. The semi-implicit conduction stage — the direct analogue of the release’s backward-Euler stage — is the same plain accumulation (`ModHeatConduction.f90:1743--1762`).

No special mechanism exists. A search for compensated/Kahan/Neumaier summation over GM/BATSRUS/src/ and srcBATL/ returns zero hits. A search for explicit `Real8_`, `real*8`, `double precision` or `selected_real_kind` over GM/BATSRUS/src/*.f90 returns 21 hits, and every one is a time or an I/O buffer — `ModMain.f90:276,283` (`StartTime/EndTime`, absolute epoch seconds), CPU clocks, satellite datetimes, input time series, and the restart records. Zero on state, flux or source. There is no well-balanced or hydrostatic-deviation formulation in the cell update; the only deviation formulation in the tree is B_0/B_1 splitting for the magnetic field (`ModB0.f90:80`), a different problem. The one precision *philosophy* signal is the contrast between `real :: Dt` (`ModMain.f90:50`) and `real :: tSimulation`

(:270) on the one hand and `real(Real8_) :: StartTime (:276)` on the other: the framework selectively promotes the quantity carrying a large additive offset, and leaves the state to the global switch.

The framing this establishes is the important output of the audit. SWMF’s answer to sub-ULP secular accumulation is not a mechanism but a policy: the conservative state has enough precision, everywhere, by default. The single-precision state in this prototype was therefore the SWMF-*inconsistent* component. **This cutover does not introduce an unusual high-precision requirement into SWMF; it corrects a standalone prototype to the precision SWMF production already uses.** Coupling becomes strictly simpler afterwards, because `share/Library/src/ModMpiModified.f90:16--19` aliases `MPI_REAL` to `MPI_DOUBLE_PRECISION` whenever `nByteReal == 8`, so a double state matches the framework’s exchange type directly.

Retired Selective mixed precision — a double persistent state with single-precision scratch and geometry — was the earlier preference and is not the design. The arrays are small at $N \sim 500$ –2000, so the memory saved does not justify the architectural complexity; the framework is uniformly double, so a split knob is the *less* consistent choice for coupling; and a split knob is precisely the failure mode above, since the previous experiment was an accidental mixed-precision build that silently did nothing.

Planned Revisit a split knob only if profiling over many field-line threads demonstrates a real cost, and then as a performance optimization with its own evidence, never as a way to preserve single-precision storage.

10.2.4 What the precision cutover does and does not fix

Active release path

Three consequences must be kept separate.

Fixed. The broad, spurious 1600–1850 km height gradient in the conserved face mass flux, which was the lower chromosphere being unable to represent its own continuity increment; and the masking of the hydrostatic truncation trend, which is now measurable and second order (section 11.4).

Not fixed, and precision independent. The *localized* first-interior-face artifact. In double the base face is still dominated by the dissipative part of the flux ($f_{\text{central}} = -4.78 \times 10^{-15}$, $f_{\text{diff}} = +1.2935 \times 10^{-13}$, $f_{\text{total}} = 1.2457 \times 10^{-13} \text{ kg m}^{-2} \text{ s}^{-1}$), so what remains there is Godunov dissipation acting on a genuinely nonzero reconstructed left/right jump — a real inner-ghost and reconstruction defect, not bit noise. The ratio $|f_{\text{diff}}|/f_{\text{total}}$ at face 0 is 15.9% in double against 16.3% in single precision, and the cell-0 ρv overshoot is 1.275 against 1.280. **The precision cutover must never be described as solving the lower-boundary mass flux.** It remains open and must be diagnosed on its own.

Untouched, and a different limitation entirely. The $N = 500$ long-duration evaporation mass flux is still not established as grid-converged to better than 10%. That is a statement about *mesh resolution*, not about storage precision, and the two must not be merged.

Active release path

Two smaller consequences are recorded so nobody re-derives them. `sim_time` is now double; it was single precision, accumulating $\Delta t = 5.6 \times 10^{-3} \text{ s}$ toward 4000 s where one single-precision ULP is $2.4 \times 10^{-4} \text{ s}$ — about 23 ULPs per step, so not yet broken, but with only about 1.5 decades of margin before it would be on a longer run. The promotion removes a future hazard rather than fixing an observed defect. And the legacy `time_mult`-derived step cap in `run_control.hpp`, previously evaluated in single precision to reproduce a historical expression bit for bit, is now evaluated in `Real`; because the value is cast to an integer the resulting cap is unchanged for every sane `time_mult`, and `CHROMO_T_END/CHROMO_STEP_CAP` bypass the path entirely.

10.2.5 Post-cutover validation

Active release path

The test suite passes 11647/11647 in both the `double` release build and the single-precision diagnostic build, **with no tolerance loosened anywhere**. The 29 failures previously seen under the incoherent `CHROMO_STATE_FLOAT64` option do not occur once the knob is coherent — they were a symptom of the `Vec`-only promotion, not of double precision exposing loosened assumptions. The exact-zero assertions that had been suspected of being quantization artifacts turn out to be structurally exact (a $v = 0$ inner boundary row, heating windows evaluated outside their support, and an energy difference between two face states built with identical ϕ_g) and hold in both precisions. The table-corner domain regression is no longer skipped in either precision: its ULP walk now steps in the storage type and nudges against the solver’s own equation-of-state acceptance rather than a test-side recomputation, which makes the construction precision independent and is a stronger assertion than the original.

`scripts/release_validation.sh` passes in the `double` release build: the default path reproduces an explicitly selected `ISO_RIEMANN=swmf-godunov ISO_RECONSTRUCTION=lnrho-v-lnp` run bit for bit; the $N = 500$ 100s smoke reaches `termination=end_time` with 2.35×10^7 faces solved exactly and *zero* fallbacks; ϵ_p is at most 0.0002%; $Q_V = 0.00133$, $Q_M = 0.00073$, $Q_{\text{eff}} = 7.26 \times 10^{-4}$; the mean upflow is 14.070 ms^{-1} with no negative-velocity cell; and $F_{\text{eff}} = 1.0414 \times 10^{-9} \text{ kg m}^{-2} \text{ s}^{-1}$, $T_{\text{top}} = 21838.6 \text{ K}$, $p_{\text{top}} = 0.010288 \text{ Pa}$, boundary-layer thickness 11.870 km , $q_{\text{phys}} = 0.7728 \text{ W m}^{-2}$, mass drift -3.454×10^{-4} .

10.2.6 Would the coupling target hit this?

This subsection is an estimate about an external code, and is labelled as one. Production AWSoM does not resolve a dynamic chromosphere: `ModChromosphere.f90:21--22` and the `#CHROMOBC` command make it a boundary condition at $n = 2 \times 10^{17} \text{ m}^{-3}$, $T = 5 \times 10^4 \text{ K}$, with the unresolved transition region *broadened* by an extension factor rather than resolved. AWSoM-R places the chromosphere and transition region on one-dimensional threads below $r = 1 R_{\odot}$, where the density is *algebraically reconstructed* from a barometric closure (`ModThreadedLC.f90:446,456--467`) and never time-integrated at all. The finest base radial cell in the coronal grid is $\sim 300 \text{ km}$, about three orders of magnitude coarser than the refined chromospheric mesh here, so Δt is $\sim 10^3$ larger while coronal timescales are comparable or shorter. And the production solar-corona run is not time accurate: it is a local-time-stepping steady-state relaxation in which $\Delta t R < u U$ is the *convergence criterion* rather than a failure. Estimated $\tau/\Delta t \sim 10^3\text{--}10^4$ there, against thresholds of 1.7×10^7 and 9.0×10^{15} : **AWSoM would plausibly be unaffected even in single precision**. That is why this failure mode has no upstream precedent, and it is exactly the difference a resolved dynamic chromosphere introduces.

10.3 Retired conservative equilibrium projection

Retired Because the seven carrier rows were fluxed independently, an explicit hydrodynamic step left them slightly off the equilibrium manifold: the per-carrier momenta drifted apart and the mass split no longer matched x_{Saha} . A projection stage therefore ran after every explicit update. It formed

$$\rho = \rho_i + \rho_n, \quad m = m_i + m_n, \quad v_{\text{cm}} = m/\rho, \quad (10.6)$$

$$e_{\text{int}}^{\text{proj}} = E_i + E_n - \frac{m^2}{2\rho} - \rho\Phi, \quad (10.7)$$

inverted $e_{\text{int}}(\rho, T)$, evaluated Saha equilibrium, and rebuilt the rows, thermalizing the removed relative-drift kinetic energy $\frac{1}{2}(\rho_i \rho_n / \rho)(v_i - v_n)^2$ exactly once.

The three totals $(\rho, m, E_i + E_n)$ were invariant under that map — the projection was the identity on the physical state and acted only on the redundant rows. With the native three-variable state of Eq. (10.1) there is no redundancy to repair, so the stage has been removed entirely: it is not an active release stage, and the spurious per-step drift heating it used to thermalize no longer exists. The removal also eliminates one full caloric inversion per cell per step.

Chapter 11

Finite-Volume Hydrodynamics

11.1 Nonuniform control volumes

Active release path Cell widths, center separations, and face metrics are cached explicitly. For a prescribed area A , a conserved state obeys

$$\frac{dU_i}{dt} = -\frac{1}{A_i \Delta s_i} (A_{i+1/2} F_{i+1/2} - A_{i-1/2} F_{i-1/2}) + S_i. \quad (11.1)$$

The active straight column has constant area, but the metric-aware implementation is retained. Uniform meshes follow guarded legacy algebra to preserve byte-level regression behavior.

11.2 Equilibrium-manifold MUSCL reconstruction

Active release path The release branch decodes each cell through the equilibrium closure and reconstructs the primitive triple

$$W = (\log \rho, v, \log p). \quad (11.2)$$

Logarithmic density and pressure enforce positivity without independent species clipping, and the face temperature is recovered by inverting the same authoritative Saha closure, $T(\rho, p)$. Each face state is then rebuilt algebraically on the Saha equilibrium manifold.

Historical The release previously reconstructed $W = (\log \rho, v, \log T)$ and obtained the face pressure afterwards from the nonlinear map $p(\rho, T)$. Two independently limited variables then set one mechanical quantity, so a mechanically smooth (constant- p) state was not reproduced across the partial-ionization transition. The resulting artificial face-pressure mismatch was the dominant source of the late-time coarse-grid velocity ripple; switching to $\log p$ reduced the mismatch by two to three orders of magnitude under both numerical fluxes, so it is a reconstruction improvement rather than a solver-specific patch. The temperature form remains selectable as a reference configuration (`ISO_RECONSTRUCTION=lnrho-v-lnt`) and is not a release setting.

On nonuniform grids, slope ratios are formed from center-to-center gradients, and each extrapolation uses the owning cell's half width. The active limiter is the asymmetric MC3/Koren form with $\beta = 2$. A diagnostic override can select minmod or first-order reconstruction, but that override is not a release setting.

11.2.1 The MUSCL–Hancock step

Active release path

The hydrodynamic time discretization is the MUSCL–Hancock method [van Leer, 1979, Toro, 2009]; Berthon [2006] gives the form closest to what is implemented here, including the use of primitive rather than conservative gradients in the reconstruction. Naming matters because “MUSCL” alone describes only the spatial reconstruction. The three pieces are:

MUSCL reconstruction. Limited extrapolation of $W = (\log \rho, v, \log p)$ from each cell centre to its two faces, as described above.

Hancock predictor. Each cell is advanced to the half-time level using *only its own* two extrapolated fluxes, plus the momentum source:

$$U_i^p = U_i^n - \frac{\Delta t}{\Delta s_i} \left[F(W_{i+1/2}^L) - F(W_{i-1/2}^R) \right] + \Delta t S(U_i^n), \quad (11.3)$$

with $S = (0, p B \partial_s B^{-1} - \rho \partial_s \Phi, 0)^\top$. This stage needs no neighbour data and therefore no Riemann solve; only two of the four face reconstructions are built for it. The half state is decoded through the *caloric* EOS only — it never needs the acoustic Γ_1 — and legacy `cons2prim` is not used.

Corrector. The four face states are rebuilt from the time average $\frac{1}{2}(W^n + W^p)$, and the release Godunov flux of section 11.3.2 is evaluated on those time-centred states.

The source term in Eq. (11.3) is not optional. The bracketed flux difference is cell i ’s own pressure gradient, which in a hydrostatic column is balanced by gravity; omitting S therefore leaves the half state of a column at rest with a spurious velocity of order $\Delta t |g|$, and the corrector reconstructs that velocity into every face. Measured on the release mesh with no equilibrium reference, restoring S cut the $t = 0$ discrete hydrostatic face mass-flux defect from 7.5×10^{-10} to $1.4 \times 10^{-11} \text{ kg m}^{-2} \text{ s}^{-1}$, a factor of 52, and the 20 s conduction-off max $|v|$ from 0.71 to 0.032 m s^{-1} . Because $\rho \Phi$ is carried inside $\mathcal{E}$, gravity contributes no energy-row source in either stage, and for the constant-area release geometry $B \equiv 1$ so the flux-tube term vanishes identically; it is retained because the operator is written for general $B(s)$.

Historical A separate experiment restructured the momentum equation in the manner of SWMF’s `ModTransitionRegion`: remove the pressure from the momentum flux entirely and evaluate its gradient explicitly alongside gravity, using the acoustic star pressure of the reconstructed face states. On this constant-area, ultra-low-Mach column it reproduced the predictor-source result and nothing more, and it was rejected on that evidence alone. For $B \equiv 1$ the shared face pressure telescopes and the momentum update stays conservative, so conservation is not a reason to reject it here; the form is, however, not generally equivalent to the conservative area-weighted operator on a variable-area flux tube. The useful part of that experiment was the isolation of Eq. (11.3)’s missing source term.

11.3 Numerical flux

Active release path

Every face solver in this chapter acts on the authoritative equilibrium-mixture system

$$U = (\rho, m, E)^\top, \quad m = \rho v, \quad F(U) = \begin{pmatrix} m \\ m^2/\rho + p \\ (m/\rho)(E + p) \end{pmatrix}, \quad (11.4)$$

where E is the stored energy row of Eq. (10.1), including mixture internal, kinetic, and gravitational energy. Both reconstructed states at a face use the same face gravitational potential, so gravity does not enter the state jump seen by the characteristic dissipation.

11.3.1 Reference flux: the mixture Roe characteristic linearization

Implemented but inactive

Not the release flux. The mixture Roe characteristic flux was the release numerical flux until the Godunov cutover of section 11.3.2, and is retained as the controlled comparison solver, selected by `ISO_RIEMANN=roe-local`. It is documented here in full because it is the reference against which the release flux is quantified, and because selecting it also reverts the timestep to the equilibrium signal speed (section 12.2), so the comparison mode remains internally consistent.

Its defining property, for the present purpose, is that it linearizes the *equilibrium* system: its acoustic eigenvalues carry Γ_1 , so it is the discrete realization of the instantaneous-Saha-equilibrium acoustic closure discussed in section 11.3.3.

For a local equilibrium state define

$$b = \left(\frac{\partial p}{\partial e_{\text{int}}} \right)_\rho, \quad c^2 = \Gamma_1 \frac{p}{\rho}, \quad H = \frac{E + p}{\rho}. \quad (11.5)$$

The implemented right eigensystem is

$$\lambda_- = v - c, \quad r_- = \begin{pmatrix} 1 \\ v - c \\ H - vc \end{pmatrix}, \quad (11.6)$$

$$\lambda_0 = v, \quad r_0 = \begin{pmatrix} 1 \\ v \\ H - c^2/b \end{pmatrix}, \quad (11.7)$$

$$\lambda_+ = v + c, \quad r_+ = \begin{pmatrix} 1 \\ v + c \\ H + vc \end{pmatrix}. \quad (11.8)$$

The contact-wave energy component $H - c^2/b$ carries the non-ideal-EOS correction; it reduces to the usual ideal-gas form when the corresponding ideal-gas derivative is inserted.

At each face, the code evaluates $v_q = m_q/\rho_q$, $p_q = F_{m,q} - \rho_q v_q^2$, $H_q = (E_q + p_q)/\rho_q$, b_q , and $c_q^2 = \Gamma_{1,q} p_q/\rho_q$ for $q = L, R$. With $w_q = \sqrt{\rho_q}$, it then forms

$$\tilde{v} = \frac{w_L v_L + w_R v_R}{w_L + w_R}, \quad \tilde{H} = \frac{w_L H_L + w_R H_R}{w_L + w_R}, \quad (11.9)$$

$$\tilde{b} = \frac{w_L b_L + w_R b_R}{w_L + w_R}, \quad \tilde{c}^2 = \frac{w_L c_L^2 + w_R c_R^2}{w_L + w_R}. \quad (11.10)$$

All quantities below are evaluated at this averaged state, with tildes suppressed for readability.

For

$$\Delta U = (\Delta \rho, \Delta m, \Delta E)^\top, \quad (11.11)$$

the implemented decomposition first forms

$$\theta = \frac{b}{c^2} [\Delta E - v \Delta m + (v^2 - H) \Delta \rho], \quad \eta = \frac{\Delta m - v \Delta \rho}{c}, \quad (11.12)$$

and then

$$\alpha_- = \frac{1}{2}(\Delta\rho + \theta - \eta), \quad \alpha_0 = -\theta, \quad \alpha_+ = \frac{1}{2}(\Delta\rho + \theta + \eta). \quad (11.13)$$

Thus the characteristic dissipation and numerical flux are

$$D = \sum_{k \in \{-, 0, +\}} |\lambda_k| \alpha_k r_k, \quad F^{\text{Roe}} = \frac{1}{2} (F_L + F_R - D). \quad (11.14)$$

The three components of D are applied directly to the three conserved rows. There is no carrier-row redistribution of the characteristic dissipation and no subsequent repacking: the flux written by the solver is exactly the flux in the equation above.

This is a local mixture Roe characteristic linearization. It is not claimed to define an exact nonlinear general-EOS Roe matrix satisfying

$$F_R - F_L = \tilde{A}(U_R - U_L). \quad (11.15)$$

If either state has invalid density, pressure, enthalpy, sound speed, or b , or if the constructed dissipation is non-finite, only that face falls back to Rusanov.

Implemented but inactive

The Rusanov / local Lax–Friedrichs flux,

$$F_{i+1/2} = \frac{1}{2} [F(U_L) + F(U_R)] - \frac{a_{i+1/2}}{2} (U_R - U_L), \quad (11.16)$$

with

$$a_{i+1/2} = \max(|v_L| + c_{s,L}, |v_R| + c_{s,R}), \quad c_s^2 = \Gamma_1 p / \rho, \quad (11.17)$$

remains available as a reference solver (`ISO_RIEMANN=rusanov`) and as the per-face fallback of *both* other solvers. It was the release flux two numerical-method generations ago. Its dissipation is acoustic in scale, $\sim \frac{1}{2} c_s \Delta U$, while the evaporation flow satisfies $|v| \ll c_s$; the resulting numerical transport left a persistent cell-centered velocity and momentum ripple in the upper chromosphere and transition region that both the Roe and the Godunov flux remove. Note that $a_{i+1/2}$ retains the *equilibrium* speed $c_s^2 = \Gamma_1 p / \rho$ even when it serves as the fallback of the frozen-composition release flux: on a fallback face the exact solve has been abandoned and the physical wave speeds of the equilibrium system are the correct Rusanov coefficient. The release timestep is the smaller of the two bounds in any case, so the fallback is never under-resolved in time.

11.3.2 Release flux: the SWMF-style exact-Riemann Godunov flux at frozen composition

Active release path

This is the release numerical flux. It follows the construction SWMF uses in `util/CRASH/src/test_godunov.f90` with the exact ideal-gas Riemann solver of `share/Library/src/ModExactRS` (Toro’s exact two-rarefaction/two-shock pressure iteration, Toro, 2009), ported line for line to `src/single_fluid/exact_rs.hpp`. It requires no environment variable; `ISO_RIEMANN=roe-local` and `ISO_RIEMANN=rusanov` select the two reference solvers instead. The three are mutually exclusive, so the face Riemann solver is the sole variable the setting changes — reconstruction, the MUSCL–Hancock predictor of section 11.2.1, the momentum source, the boundary closures, the Saha closure and the conduction split are identical under all three — with the single deliberate exception of the timestep, which follows the flux in force (section 12.2).

The method solves an *ideal-gas* Riemann problem at a fixed $\gamma = 5/3$ and carries every non-ideal part of the energy in a passive specific offset attached to the material. For a face state (ρ, u, p, E) ,

$$E_0 = \frac{E - \frac{1}{2}\rho u^2 - p/(\gamma - 1)}{\rho} = \frac{e_{\text{int}} - p/(\gamma - 1)}{\rho} + \phi_g, \quad (11.18)$$

so that E_0 absorbs both the Saha ionization energy inside e_{int} and the gravitational potential this model carries in E . Because both sides of a face share ϕ_g , only the equation-of-state content is actually upwind. With $(\rho_\star, u_\star, p_\star)$ the exact self-similar solution sampled along $x/t = 0$, and E_0 taken from whichever side the contact declares upwind,

$$E_\star = \frac{p_\star}{\gamma - 1} + \frac{1}{2}\rho_\star u_\star^2 + \rho_\star E_0, \quad F_{i+1/2}^{\text{God}} = (\rho_\star u_\star, \rho_\star u_\star^2 + p_\star, (E_\star + p_\star)u_\star). \quad (11.19)$$

When the sampled state coincides with a side state this reproduces that side's total energy exactly, and a constant offset is advected without altering (ρ, u, p) at all. Any face whose exact solve fails — inadmissible input, the two-rarefaction vacuum guard, a non-positive star pressure, or an exhausted iteration budget — or whose sample is inadmissible falls back to the same face-local Rusanov flux the Roe path uses, and is counted and reported at the end of the run rather than absorbed silently.

The solver reproduces the five standard Riemann tests of Toro [2009] to the six significant digits their star pressures are tabulated to, and a first-order Godunov solve of Sod's problem with equation (11.19) converges at the expected L^1 rate while leaving a nonzero E_0 inert to round-off. On the conduction-off $N = 500/\text{R4}$ release mesh (661 cells, CFL 0.50, 3620 steps) the initialized-atmosphere face mass-flux defect is 1.08×10^{-12} against 1.33×10^{-12} for Roe, and the 20s velocity residual is $3.19 \times 10^{-2} \text{ m s}^{-1}$ against 3.02×10^{-2} . Both of those legs were run under the pre-cutover single-precision storage and both sat at its round-off floor, so that measurement does not separate the two fluxes and is not a current release hydrostatic number; the **double** re-measurement of the release Godunov leg is in section 11.4 (1.7782×10^{-13} and $5.7274 \times 10^{-3} \text{ m s}^{-1}$), the Roe counterpart has not been re-run in **double**, and the flux comparison therefore rests on the 100s conduction-driven observables below rather than on the hydrostatic defect. On the separate 100s release smoke configuration (17781 steps, identical for both solvers) every conduction-driven observable agrees to a few parts in 10^4 : mean upflow 14.96 against 14.96 m s^{-1} , effective face mass flux 1.1168×10^{-9} against $1.1165 \times 10^{-9} \text{ kg m}^{-2} \text{ s}^{-1}$, with identical T_{top} , p_{top} , boundary-layer thickness and conductive flux, and no negative-velocity cells in either. The exact solve succeeded at all 2.35×10^7 face evaluations with no fallback and with the star-pressure Newton iteration converging in a single step at every face, at a cost of +12.9% in the explicit right-hand side and +4% in measured whole-step work. The dissipative part of the Godunov mass flux runs 8–12% below Roe's during the evaporation phase; no causal attribution is claimed for that, because the measurement does not isolate the contact field and a competing effect — the frozen flux's faster waves — pushes the opposite way.

Active release path

Every release run reports what the exact solver did. The driver prints a **godunov.*** tally (faces attempted, faces solved exactly, fallbacks broken down by cause, and Newton-iteration statistics), warns on **stderr** if any face fell back, and **scripts/release_validation.sh** fails the release smoke on a nonzero fallback count. A silent degradation to Rusanov is therefore not possible.

11.3.3 Thermodynamic interpretation: frozen composition versus instantaneous Saha equilibrium

Active release path

The fixed $\gamma = 5/3$ of Eqs. (11.18)–(11.19) is the substantive modelling content of this flux, and it is a statement about the thermodynamics of the wave, not an ideal-gas approximation to the equation of state. This subsection states the distinction, because it is easy to state it wrongly.

The release caloric closure, Eq. (3.9), splits exactly into a translational term and an ionization reservoir,

$$e_{\text{int}} = \frac{p}{5/3 - 1} + e_{\text{ion}}, \quad e_{\text{ion}} = x n_{\text{H}} \chi_{\text{H}}, \quad (11.20)$$

so a monatomic translational gas plus a store of ionization energy is not an approximation of this equation of state — it is an identity satisfied by it. The two closures differ only in what happens to e_{ion} while a compressive disturbance passes.

Instantaneous Saha equilibrium. If the ionization balance re-equilibrates within the wave, the isentropic response is the equilibrium first adiabatic exponent Γ_1 of equation (3.10), and the acoustic speed is $c_{\text{eq}}^2 = \Gamma_1 p / \rho$. Compressing the gas then partly ionizes it, the ionization reservoir absorbs part of the compressional work, and Γ_1 falls well below the monatomic value: in the release column its median is 1.090 and it lies below 1.5 across 94.9% of cells.

Frozen composition. If the ionization state cannot change on the timescale of the wave, e_{ion} is inert, only the translational term of equation (11.20) responds, and the acoustic speed is the monatomic frozen value $c_{\text{fr}}^2 = (5/3) p / \rho$.

Neither exponent is universally correct; they are the two limits of the same closure, and the choice between them is set by the ratio of the ionization/recombination relaxation time to the dynamical time of the disturbance. For chromospheric hydrogen that ratio is decisively in favour of the frozen limit. Carlsson and Stein [2002] compute the relaxation timescale directly and find it dominated by the slow collisional leakage from the ground state to $n = 2$: about 1 s in the photosphere, rising to $\sim 10^5$ s in the mid-chromosphere and falling to $\sim 10^2$ s at the transition-region base, and still 10 – 10^3 s inside shocks, so that “the timescale is still too long for the ionization to reach its equilibrium value in the shock”. Leenaarts et al. [2007] reach the same conclusion in two-dimensional radiation-MHD and state the energetic consequence directly — in non-equilibrium, hydrogen ionization is “a less effective internal energy buffer”, giving shock temperatures of 10 000–13 000 K against ≤ 8000 K under LTE. Leenaarts [2020] puts the chromospheric hydrodynamic timescale at about one minute and summarizes the limit as: with a long relaxation time “an increase in internal energy leads directly to a temperature increase”, and “the gas in the chromosphere behaves somewhat like an ideal gas”.

The release therefore adopts the frozen-composition interpretation for the local Riemann interaction: the composition is held fixed while the waves cross the face, the ionization energy is carried inertly in E_0 , the translational gas is monatomic, and the exact Godunov solver evolves precisely that frozen-composition local problem. Γ_1 is not discarded. It remains the equilibrium derivative of the closure, it is the quantity reported in the diagnostic sidecar, it sets the outer-boundary Mach cap, and it is what the reference Roe flux of section 11.3.1 linearizes.

Two things this argument does *not* establish should be stated plainly. It does not show that the release result is more accurate: no observable measured on this model separates the two closures, and the numerical case for the change is neutral (section 11.3.2). And it does not resolve the consistency of the frozen Riemann solve with the equilibrium decode applied around it — see section 20.4.

11.3.4 Remaining limitations of the release flux

Active release path

Three limitations survive the promotion, and one was retired by it.

1. **The two closures have not been separated by measurement on this model.** At the Mach numbers of the release column, every physical observable agrees between the frozen and equilibrium fluxes to a few parts in 10^4 . The release choice rests on the relaxation-time argument above and on the SWMF methodology, not on measured evidence that it is better.
2. **Only smooth, low-Mach behaviour has been exercised.** The regime where the distinction would matter most — a strongly compressive, shock-forming event inside the partial-ionization zone — has been tested only in the ideal-gas shock tube, where γ is correct by construction. The flare and shock scenarios run the non-release two-fluid solver, which this flux does not touch.
3. **The offset is upwinded by $\text{sign}(u_*)$ alone,** a piecewise-constant switch rather than a smooth function of the state. This is SWMF’s construction and it is exact for a contact; no pathology was observed, because $u_* \rightarrow 0$ also drives the energy flux to zero, but it is not smooth.
Retired: the timestep inconsistency. Before the cutover the CFL was deliberately taken from the equilibrium Γ_1 speed so that the two comparison legs took identical steps, and this document previously recorded the resulting effective CFL as “about 0.62”. That figure was wrong, and the inconsistency is now fixed; see section 12.2.

11.4 Gravity and reference-free hydrostatic balance

Active release path Face states use a common face gravitational potential on both sides, removing an artificial gravitational-energy jump from the numerical dissipation. Beyond that the release carries *no* well-balancing mechanism: the *interior* hydrodynamic operator is given no hydrostatic state in advance and subtracts no reference residual. This is a statement about the interior discretization only. The boundary closures still capture fixed reservoir data once from the initial condition — the outer back pressure `kOuterPRef` (section 9.4) and the lower reservoir state `kInnerTRef`/`kInnerPres` (section 9.3) — as any truncated-domain problem must. What was removed is the dependence on a global hydrostatic *profile*, not every frozen quantity in the problem. Two classes of well-balanced scheme exist for the Euler equations with gravity — those that require the target equilibrium a priori [Berberich et al., 2019] and those that preserve a class of discrete equilibria to machine precision without knowing them [Käppeli and Mishra, 2014, Chandrashekar and Klingenberg, 2015]. **The release belongs to neither, and should not be described as a well-balanced scheme.** It removes the two discretization inconsistencies that dominated its hydrostatic imbalance and leaves a small, resolution-dependent residual, rather than constructing exact preservation of anything.

The defects were the missing predictor source of section 11.2.1 and the first-order lower ghost ladder of section 9.3. On the release mesh, measured as the numerical face mass flux of the initialized $u = 0$ atmosphere with no reference subtraction:

Configuration (single-precision storage)	$\max F^\rho $ at $t = 0$	20 s conduction-off $\max v $
Previous release	7.5×10^{-10}	0.71 ms^{-1}
+ predictor source	1.4×10^{-11}	0.032 ms^{-1}
+ trapezoidal lower ghosts	1.3×10^{-12}	0.030 ms^{-1}
single-precision floor $\epsilon \rho(v + c_s)$	1.0×10^{-12}	—

Historical That sequence was measured under the single-precision state storage then in use, and its endpoint sat at 1.07 times the storage round-off floor — that is, at the representation limit. On that evidence this document previously concluded that “no clean truncation-error trend survives over the resolutions tested”. **That conclusion was a property of the storage, not of the scheme, and it no longer holds.**

Active release path

Re-measured with the release `double` state (section 10.2), same configuration — $N = 500/\text{R4}$, 661 cells, 20 s, conduction off, release Godunov flux, CFL 0.50, `godunov.fallbacks=0`, `termination=end_time`:

$N = 500/\text{R4}$, Godunov, 20 s	single precision	double (release)
$t = 0 \max F^\rho $ [kg m ⁻² s ⁻¹]	1.0758×10^{-12}	1.7782×10^{-13}
location of that maximum	1601.1 km (base face)	2149.1 km (upper TR)
$t = 0$ base-face $ F^\rho $	1.0758×10^{-12}	1.2457×10^{-13}
$t = 0$ rms / median $ F^\rho $	3.163×10^{-13} / 2.033×10^{-13}	1.991×10^{-14} / 3.271×10^{-16}
$\max v $ at 20 s [m s ⁻¹]	3.1943×10^{-2}	5.7274×10^{-3}
location	2093.9 km (upper TR)	2153.0 km (top cell)
defect / storage round-off floor	1.07	9.5×10^7
$N = 500 \rightarrow N = 1000$, $\max F^\rho $	$\rightarrow 2.1518 \times 10^{-12}$, order -1.0	$\rightarrow 3.9112 \times 10^{-14}$, order 2.19
base-face $ F^\rho $	ratio 0.73	ratio 3.86, order 1.95
rms $ F^\rho $	ratio 0.49	ratio 6.64, order 2.73
$\max v $ at 20 s	2.4885×10^{-2} , order 0.36	1.5291×10^{-3} , order 1.90

Active release path

Two results, and they matter differently. First, the single-precision residual was the representation limit and the `double` residual is nowhere near its own: $1.07\times$ the single-precision floor against $\sim 10^8\times$ the `double` floor. The surviving residual therefore has *no representation component at all* and is entirely discretization and boundary closure. Second, **the reference-free discretization is second-order convergent on the hydrostatic fixed point, and this could not be measured before.** In `double`, refining $N = 500 \rightarrow N = 1000$ reduces the defect at order 2.19 in the maximum, 1.95 at the base face, 2.73 in the rms, and 1.90 in the 20 s velocity residual. In single precision the *same* refinement made the defect worse, at apparent order -1.0 , which is the signature of per-face ULP noise rather than truncation. This retires a real ambiguity in the reference-free argument: the residual is a genuine truncation error of a convergent scheme, and no longer needs to be hedged as possibly a storage artifact.

Active release path

That two-point estimate is not the strongest available evidence. A six-rung refinement ladder — five reference-free steps on the initialized column after `model_column_ic`, CFL 0.25, conduction off, `godunov.fallbacks=0` throughout — spans a $41\times$ range in mesh spacing:

ISO_NS	cells	Δs [km]	double $\max v $ [m s ⁻¹]	order	single precision	order
48	70	11.410	8.109×10^{-1}	—	8.098×10^{-1}	—
125	169	4.394	7.617×10^{-2}	2.48	7.702×10^{-2}	2.47
250	331	2.209	1.878×10^{-2}	2.04	1.790×10^{-2}	2.12
500	661	1.105	4.986×10^{-3}	1.91	5.845×10^{-3}	1.62
1000	1320	0.553	1.317×10^{-3}	1.92	3.381×10^{-3}	0.79
2000	2638	0.276	3.321×10^{-4}	1.99	4.590×10^{-3}	-0.44

In `double` the order is a clean, unbroken ≈ 2 across the entire ladder, with no flattening at the fine end, and $\max |v|$ falls monotonically from 8.11×10^{-1} to 3.32×10^{-4} m s⁻¹. Single precision tracks it only to ISO_NS=250, degrades to order 0.79 by 1000, and *reverses* at 2000, bottoming out near 3.4×10^{-3} m s⁻¹ — a round-off floor, reached and then climbed away from. Over the same ladder the $t = 0 \max |F^\rho|$ in `double` falls $3.32 \times 10^{-11} \rightarrow 1.31 \times 10^{-14}$ at orders 1.96–2.41 and the rms falls $8.62 \times 10^{-12} \rightarrow 5.46 \times 10^{-16}$ at orders 2.45–2.73. **The convergence is a property of the discretization, not of the release flux:** a double *Roe* ladder at the

identical configuration gives the same orders, 1.91–2.41. And the archived single-precision/Roe ladder reproduces digit-for-digit in the diagnostic build with `ISO_RIEMANN=roe-local`, which is what makes the comparison a controlled measurement rather than a redefinition. (Do not cross-substitute the two release-mesh numbers: $4.986 \times 10^{-3} \text{ m s}^{-1}$ is the 5-step value at `ISO_NS=500`; the 5.727×10^{-3} in the table above is the 20s value.)

Active release path

What survives is real, and is the open defect — “six times better” must not be read as “solved”. In `double` the base face is still dominated by the dissipative part of the flux ($f_{\text{central}} = -4.78 \times 10^{-15}$, $f_{\text{diff}} = +1.2935 \times 10^{-13}$, $f_{\text{total}} = 1.2457 \times 10^{-13}$), so the remaining base defect is Godunov dissipation responding to a genuinely nonzero reconstructed left/right jump: an inner-ghost and reconstruction defect, precision independent, and the same feature as the first-interior-face artifact of section 10.2.4. Likewise the `double` $\max|v|$ now sits on the *topmost* cell, where the value is nearly precision independent (5.27×10^{-3} against $5.73 \times 10^{-3} \text{ m s}^{-1}$), so that part is an outer-boundary-closure residual that precision does not touch, while the interior 2090 km peak collapsed. Both residuals remain two to three orders of magnitude below the 15–40 m s^{-1} evaporation flow. A 5-step release-mesh regression asserting this reference-free behaviour is part of the test suite. **The one quantity not re-measured** is the 20s relative density drift, which the six-significant-digit snapshot cannot resolve at the `double` level.

Retired The previous release stored the initialized equilibrium U^{eq} , evaluated its explicit residual once, and advanced

$$\frac{dU}{dt} = R(U) - R(U^{\text{eq}}), \quad (11.21)$$

a δ -form equilibrium-reference method that made that one chosen state an exact discrete fixed point. It was removed for three reasons. It bound the solver to a pre-known global equilibrium. The residual was evaluated once, at $t = 0$ and at the first step’s Δt , and then subtracted unchanged for the rest of the run — including after the transition region had completely restructured, and at every subsequent Δt , so the “exact fixed point” property held strictly only at $t = 0$. And, decisively, it no longer changed the answer: with the corrected discretization, enabling it altered the 100 s evaporation solution by 0.03% (0.15% at $N = 1000$) and the 1000 s solution by 0.4% in analysis-window velocity, 0.5% in peak velocity, and 0.9% in column mass drift. `Grid::eq_wb`, `eq_state` and `eq_residual` still exist for the two-fluid research solver, which never received the predictor source term; a release `Grid` sets `eq_wb=false` and the `ISO_EQ_WB` knob is gone.

Active release path

The $t = 0$ face mass flux F_{ref}^{ρ} remains useful as a *diagnostic* baseline, and boundary diagnostics still report $F_{\text{eff}}^{\rho} = F_{\text{total}}^{\rho} - F_{\text{ref}}^{\rho}$ — the change in transport since initialization. It is no longer anything the solver subtracts. A smooth F_{eff}^{ρ} can coexist with oscillatory cell-centered ρv because the central and diffusive flux components compensate.

Chapter 12

Implicit Conduction and Time Integration

12.1 Nonlinear backward-Euler conduction

Active release path

At fixed density, the conduction stage solves

$$e_{\text{int}}(\rho, T^{n+1}) - e_{\text{int}}(\rho, T^*) - \Delta t \nabla_s \cdot [K(T^{n+1}) \nabla_s T^{n+1}] = 0, \quad (12.1)$$

where

$$K = \kappa_e + \kappa_n, \quad (12.2)$$

and T^* is the temperature decoded from the post-hydrodynamic state U^* . Conductivity, Saha composition, and effective heat capacity are refreshed each nonlinear iteration.

The quasi-Newton matrix uses C_V^{eq} on the diagonal and treats the updated face conductivities in a Picard-like manner. Each iteration solves a tridiagonal system with the Thomas algorithm. Convergence is accepted only after evaluating the full nonlinear residual at the updated temperature. A small update that fails the residual test is reported as stagnation; reaching the iteration limit is an error.

On refined meshes, interior physical conductivity uses the width-weighted series resistance of the two half cells. Uniform meshes retain the historical arithmetic average. The production numerical-diffusivity multiplier is zero, so no artificial conductivity enters this solve. Closed-boundary tests verify that the weighted total-energy contribution telescopes on irregular grids.

12.2 Timestep control

Active release path

The hydrodynamic timestep is

$$\Delta t_{\text{ac}} = C_{\text{CFL}} \min_i \frac{\Delta s_i}{|v_i| + c_i}, \quad c_i^2 = \frac{\gamma_{\text{sig},i} p_i}{\rho_i}, \quad (12.3)$$

where γ_{sig} is the acoustic index of the *numerical flux in force*, not unconditionally the equilibrium one:

$$\gamma_{\text{sig}} = \begin{cases} \max(\Gamma_1, 5/3), & \text{release Godunov flux (section 11.3.2),} \\ \Gamma_1, & \text{Roe and Rusanov reference fluxes.} \end{cases} \quad (12.4)$$

The rule exists so that C_{CFL} always means the CFL of the scheme actually running. The release flux propagates the frozen waves of section 11.3.3, which in the partial-ionization interior travel about 1.24 times faster than the equilibrium waves; sizing the release step from Γ_1 would leave that factor as an undeclared margin. The max form, rather than a bare constant, keeps the bound conservative should the table ever return an index above the monatomic value, and it keeps the Roe comparison mode internally consistent rather than merely unchanged. Optional cooling, heating, or electron-energy physics can add smaller limits in non-release configurations. In the validated reference calculation, every step was acoustic limited.

Active release path

Measured effect on the canonical release configuration: none, and the reason is worth recording. The CFL-limiting cell of the $N = 500/\text{R4}$ column is the topmost cell (2153 km, ~ 22 kK, hydrogen fully ionized, $\Delta s = 270$ m), where the tabulated Γ_1 is already 5/3. The large frozen/equilibrium speed gap lives in the partial-ionization interior, whose sound speed is far lower and which is nowhere near the CFL limit. Measured on both the $t = 0$ and the $t = 100$ s states, the two rules select the same cell, give Δt ratios of $1 - 10^{-6}$, and produce the identical step count (17781 over 100 s) and end time. Measured against the frozen wave speeds, the *equilibrium*-sized step gave a maximum effective CFL of exactly 0.500. An earlier version of this document reported an effective CFL of about 0.62 for this flux; that figure assumed the limiting cell carried $\Gamma_1 \approx 1.09$ and is withdrawn.

The change is nonetheless required, as a guarantee rather than a correction: on a different mesh, a deeper domain, a different outer temperature, or a refinement that places a partial-ionization cell at the CFL limit, only Eq. (12.4) bounds the release scheme correctly.

The executable fallback is $C_{\text{CFL}} = 0.25$. The low-I/O production launcher supplies 0.50, which has been validated for the canonical coarse-equivalent $N = 500$, R4, 661-cell physical-conduction-only configuration on the measured Apple Silicon system. That sweep was run under the Roe flux and the equilibrium signal speed; it transfers to the release flux because Δt and the step count are unchanged on this exact configuration, as measured above, and for no more general reason. It is not a universal stability guarantee.

12.3 Active release update path

Active release path

One step on the three-variable state of Eq. (10.1) — explicit hydrodynamics Lie-split from implicit conduction — is

$$U^n \xrightarrow{\text{explicit MUSCL-Hancock/Godunov hydrodynamics}} U^* \xrightarrow{\text{implicit physical conduction}} U^{n+1}.$$

More explicitly:

1. decode U^n once through the equilibrium closure, caching $(T, x, p, \Gamma_1, e_{\text{int}})$ per cell and for the two ghost layers;
2. reconstruct $(\log \rho, v, \log p)$ to cell i 's own two faces with the MC3/Koren limiter, and take the Hancock half step of Eq. (11.3), which applies the same momentum source as the corrector; decode the half state through the caloric EOS;
3. rebuild all four face reconstructions from the time average $\frac{1}{2}(W^n + W^p)$, close them into equilibrium face states, and form the SWMF exact-Riemann Godunov fluxes of section 11.3.2, using face-local Rusanov only if the exact solve or its sample is inadmissible;
4. evaluate the explicit hydrodynamic RHS — flux divergence plus momentum source, with no reference subtraction — and advance density, momentum, and energy to the intermediate state U^* ;

5. solve the nonlinear backward-Euler physical-conduction problem at fixed ρ and m , writing back $E^{n+1} = e_{\text{int}}^{n+1} + m^2/2\rho + \rho\Phi$.

The two stages are composed by *first-order Lie splitting*: the conduction solve takes its initial temperature from the completed hydrodynamic state U^* , and the composition is not symmetrized about it. The MUSCL–Hancock construction supplies second-order space and time centering for the *homogeneous* flux evolution, but the corrector evaluates the momentum source at U^n rather than at the half-time state; in the source-only limit $dU/dt = S(U)$ the update degenerates to forward Euler, so the complete sourced hydrodynamic stage is not formally second order in time either. Composed with backward-Euler conduction, the step is first order in time overall. **No second-order temporal accuracy is claimed for the release, for the hydrodynamic stage or for the composition.** Raising the order of the source quadrature and symmetrizing the splitting are two separate questions, neither addressed here. This is acceptable here because the timestep is set by the acoustic CFL condition while the physics of interest evolves at $|v| \ll c_s$, so the splitting error per step is far below the acoustic scale that sets Δt ; whether a Strang composition would measurably improve the long-duration evaporation solution is a separate numerical question that has not been tested and is not addressed here. There is no projection, repacking, or manifold-repair stage between the two: U is already the physical state, and Saha equilibrium enters only through the EOS closure that decodes it. Those two stages are the entire release timestep. The release performs no finite-rate ionization source step, and artificial conduction, TRAC, beam or coronal heating, radiative cooling, drag, separate thermal equilibration, and floor/slaving stages are not merely disabled here — they are not implemented in this solver at all, so no default-off flag stands between the reader and the algorithm. They belong to the two-fluid research solver described in the next section. For the same reason the release timestep control has only one constraint, the acoustic CFL condition: there is no volumetric source term to impose a heating-timescale cap, and the implicit conduction solve is unconditionally stable.

12.4 Legacy/full two-fluid update path

Implemented but
inactive

The older, non-release sequence is

```
hydrodynamics → drag → thermal equilibration
→ conduction → ionization → heating/cooling
→ floors/slaving
```

It operates on the independent seven-row ion, neutral, and electron state and uses fixed- γ pressure updates. Since the release was decoupled it is a genuinely separate solver: it owns its own conserved-state width, its own packing convention, its own boundary construction, and its own source stages, and its entry points reject a Grid that carries a Γ_1 table. It remains valuable for two-fluid drift, three-temperature, and finite-rate-ionization experiments (`model_gentle`, `model_c7`, `model_flare`, `analytic_canopy`, `pfss_field_line`), but it must not be used to describe the active release algorithm.

Chapter 13

Static Refinement and Inactive Artificial Diffusion

13.1 Mesh construction

Active release path Refinement is static, not adaptive mesh refinement. The face grid is built once before the initial state, and no regridding, remapping, or runtime adaptation occurs.

ISO_NS specifies the coarse-equivalent cell count. Refinement only adds cells. Two profiles exist: **lower** a finite refined band followed by a geometric grade back to the coarse spacing;

outer a coarse region, a geometric coarse-to-fine grade, and a fine band extending to the outer face.

The adjacent cell-width ratio is capped at 1.1. The requested transition width is a minimum; it is widened when necessary to satisfy the ratio cap. Endpoints and the fully refined-region boundary are exact faces.

For the release outer profile, coarse-equivalent $N = 500$ produces 661 actual cells: coarse spacing about 1.105 km and fine spacing about 0.276 km.

13.2 Historical local numerical-conduction scaling

Historical The original implementation stored one global diffusivity proportional to the coarse spacing. Local refinement then reduced the wall stencil spacing without reducing the artificial conductivity, increasing the total conductive drive. A later face-local formulation stored C_{num} and evaluated $\chi_{\text{num},f} = C_{\text{num}}\Delta s_f$ independently at every face. A factor-four refined band therefore received approximately one quarter of the coarse-region diffusivity.

That change removed the refinement-induced increase in evaporation flux and preserved the uniform-grid result exactly, making it a useful controlled diagnostic. It has since been superseded for production: the artificial term was removed from the release conduction operator altogether, so the release advances physical conduction only by construction. The formulation survives only inside the two-fluid Stage-D operator, for historical and diagnostic comparisons, and is not release stabilization.

Part V

Implementation and Performance

Chapter 14

Code Architecture

14.1 Primary modules

File or module	Current responsibility
<code>scenarios/model_column.cpp</code>	Release-mode activation, domain and mesh parameters, PCHIP/Saha-HSE initial state, lower and upper boundaries, equilibrium reference. The same file also builds the shared two-fluid initial state for <code>model_gentle</code> , whose source controls are rejected in release mode.
<code>scenarios/ model_c7.cpp/.hpp</code>	Model C7 tables, historical interpolator, gamma-mode PCHIP temperature, photoionization support for legacy kinetic studies.
<code>scenarios/mesh.cpp/.hpp</code>	Uniform and static refined face-grid construction, profile parsing, geometric grading, ratio-cap enforcement.
<code>src/eos.cpp, eos.hpp</code>	<i>Shared.</i> Saha closure, caloric EOS, derivatives, table parser/interpolator, temperature inversion, mixture decode and equilibrium face mapping.
<code>src/grid.cpp, chromosphere.hpp</code>	<i>Shared.</i> Grid: mesh geometry, prescribed $B(s)$, gravity, physical constants, metric caches, runtime flags, both state-row index sets, both ghost-buffer sets, diagnostic captures. Declares neither solver's entry points.
<code>src/single_fluid/mixture.hpp</code>	<i>Release solver.</i> Governing equations, software contract, and the complete release API.
<code>src/single_fluid/mixture.cpp</code>	<i>Release solver.</i> Three-variable state decode and packing, metric-aware $(\log \rho, v, \log p)$ MUSCL reconstruction, equilibrium face states, SWMF exact-Riemann Godunov flux with face-local Rusanov fallback (Roe and Rusanov reference solvers alongside), MUSCL-Hancock predictor/corrector with a source-carrying half step, geometric and gravitational momentum source, acoustic timestep.
<code>src/single_fluid/integrator.cpp</code>	<i>Release solver.</i> Nonlinear backward-Euler <i>physical</i> conduction and the two-stage release timestep.

<code>src/two_temp/two_temp.hpp</code>	<i>Experimental two-temperature solver</i> (chapter 18). Governing equations, four-row $(\rho, \rho v, E, E_e)$ state, closure and collision declarations, caller-owned workspace, guard-counter struct. Non-release; reached only by <code>model_column_2t</code> .
<code>src/two_temp/closure.cpp</code>	<i>Experimental two-temperature solver</i> . Electron caloric closure as exact differences of the release EOS, T_e inversion, decode/pack, algebraic face states, κ_e/κ_n , and the g_{ei} exchange conductance.
<code>src/two_temp/hydro.cpp</code>	<i>Experimental two-temperature solver</i> . Four-slot $(\log \rho, v, \log p, \log T_e)$ MUSCL reconstruction, MUSCL–Hancock predictor with the electron compression-work source, the SWMF exact-Riemann Godunov flux on the three release rows with contact upwinding of $\mathcal{E}_e/\rho$, and the frozen acoustic timestep.
<code>src/two_temp/conduction.cpp</code>	<i>Experimental two-temperature solver</i> . Coupled implicit (T_e, T_i) conduction and collisional exchange on a 2×2 -block tridiagonal Newton solve, and the characteristic-derived boundary conditions of section 18.6.
<code>src/two_fluid/two_fluid.hpp</code>	<i>Two-fluid research solver</i> . Seven-row solver API: packed-state helpers, fluxes, source stages, integrators.
<code>src/two_fluid/state.cpp</code>	<i>Two-fluid research solver</i> . Packed seven-row indexing, ghost-layer access, <code>cons2prim/prim2cons</code> , limiter algebra.
<code>src/two_fluid/flux.cpp</code>	<i>Two-fluid research solver</i> . Fixed- γ physical fluxes, spectral radii, and explicit sources.
<code>src/two_fluid/rhs.cpp</code>	<i>Two-fluid research solver</i> . Seven-row reconstruction, predictor/corrector hydro operator, drag and conduction RHS.
<code>src/two_fluid/integrators.cpp</code>	<i>Two-fluid research solver</i> . Timestep limits, drag, thermal equilibration, conduction (including the historical artificial diffusivity), finite-rate ionization, TRAC, heating/cooling, floors, and the two-fluid update sequence.
<code>physics.hpp</code>	Inline formula library: collision frequencies, conductivities, radiative losses, heating rates, TRAC broadening, and the ionization/recombination rate coefficients. Used by the two-fluid solver and by EOS diagnostics; the release solver does not include it.
<code>chromo_main.cpp</code>	Command-line contract, fail-closed mode checks, run controls, output scheduling, diagnostics, termination reporting.
<code>parallel.hpp</code>	Static-schedule OpenMP loops and deterministic lowest-cell exception propagation.
<code>scripts/run_chromo_omp.sh</code>	Generic 12-thread OpenMP launcher.
<code>scripts/ run_chromo_realtime.sh</code>	Validated low-I/O 12-thread launcher with CFL 0.50.
<code>tests/chromo_tests.cpp</code> and CTest	Unit, regression, process-level activation, table, output, and mesh validation.

14.2 Runtime contracts

Release mode is deliberately narrow, and it is selected by loading a Γ_1 table — no toggle chooses between the Saha and the two-fluid solvers. The `model_column` scenario supplies the production table unconditionally, so the scenario name and the solver identity are the same fact: `model_column` is always the release solver, the historical two-fluid column is the separate `model_gentle` scenario, and the experimental two-temperature solver of chapter 18 is reached only through the separate `model_column_2t` scenario, whose initial condition is what selects it. The retired `model_isentropic` alias has been removed. The driver requires the semi-implicit `full` integrator and an explicit `no-ionization` command-line selection. Any legacy two-fluid setting supplied to a release run — `SINGLE_FLUID`, `ENABLE_TE`, `ISO_TWO_FLUID`, `ISO_IONIZATION`, `ISO_GAMMA`, `ISO_COOLING`, `ISO_TRAC`, `ISO_CORONA`, `ISO_CHEAT`, `ISO_QFLUX`, `ISO_TBOOST`, `ISO_NUMERICAL_DIFFUSIVITY_MULT` — is rejected outright rather than silently ignored, because none of them has meaning for a single-fluid common-temperature mixture advanced by hydro and physical conduction alone. Incompatible requests fail before the initial state is built.

Release state snapshots carry the three conserved rows (ρ, m, E); legacy two-fluid snapshots carry seven, and experimental two-temperature snapshots carry four. The header line records the width, and the driver validates it against the solver the scenario selected before evolution begins. A companion `.gamma_diag` stream reports physical mixture density, velocity, temperature, Saha fraction, number densities, total pressure, Γ_1 , and conductivity. Because release conduction is physical only, its `kappa_solver` column is identically `kappa_physical`; it is retained for format stability with the existing analysis scripts. The `.outercond` sidecar likewise now reports a single physical face conductivity and flux rather than a physical/artificial split. Optional `.faceflux` and `.outercond` files are default-off read-only diagnostics and do not feed the solver.

Chapter 15

Performance Engineering

15.1 Serial bottlenecks and method-preserving optimization

The first gamma-table implementation spent most of its time repeatedly decoding the same cell, solving caloric inversions, allocating temporary vectors, and rebuilding thermodynamic quantities already known to a caller. The major optimization stages preserved the numerical method:

Historical The timings in the following table used the superseded 2638-cell configuration and document the optimization history rather than current release throughput.

Stage		Material change	Measured effect
EOS/cache restructuring		Fast zero/one-update caloric inversion, decoded-cell caches, known-temperature packing, scratch reuse, fewer redundant table lookups.	About $4.25\times$ relative to the early gamma-table implementation.
Serial cleanup	production	Algebraic pressure-density inversion, fused EOS evaluation, fewer repeated logarithms and Γ_1 calls, reduced temporary work.	124.7 s to about 78.1 s for the same 20 s, 2638-cell case; about $1.60\times$.
Final serial follow-up		Additional cache and call-path cleanup without changing accepted states.	Median about 76.1 s, 2.67 ms per step.
OpenMP		Parallel per-cell decode, face construction, and a persistent conduction region.	16-thread median 36.35 s; $2.17\times$ same-session speedup.
Thread-count selection		Compare 12 and 16 threads on Apple Silicon.	12 threads within about 1% of maximum throughput and selected as the default.

These numbers are hardware- and configuration-specific. They document engineering impact, not an algorithmic complexity guarantee.

15.2 Deterministic OpenMP

**Active release
path**

OpenMP is opt-in at CMake configuration. The runtime uses static scheduling, index-disjoint writes, fixed per-thread reduction slots, and deterministic exception selection by the lowest physical cell. The generic and production launchers set

`OMP_DYNAMIC=FALSE, OMP_MAX_ACTIVE_LEVELS=1, OMP_NUM_THREADS=12.`

Serial, one-thread OpenMP, and 12-thread OpenMP outputs were byte-identical in the validated checks.

The conduction team is capped at four threads because the remaining work alternates parallel cell/face evaluations with global barriers and a serial tridiagonal solve. The Thomas algorithm and nonlinear synchronization are the principal limits to further scaling. Replacing them would be a numerical-algorithm project, not a minor threading edit.

15.3 Long-run controls

**Active release
path**

`CHROMO_T_END` requests an exact physical end time. When it is set without `CHROMO_STEP_CAP`, the legacy time-multiplier step cap is removed. `CHROMO_FRAME_DT` schedules snapshots in physical time without cumulative drift; the initial and final states are written exactly once. Every run reports its termination cause, requested and final physical times, final step, and effective step cap.

The current $N = 500$, R4, 661-cell $\log p$, physical-conduction-only run (measured under the Roe flux) reached its requested 4000 s end time in 712314 steps with 12 OpenMP threads. Every timestep was acoustic limited, which is now the only limiter the release timestep can report. The nonlinear conduction solve averaged about 2.02 Newton updates and used at most four.

Part VI

Verification and Numerical Results

Chapter 16

Verification by Numerical Risk

16.1 Thermodynamic consistency

Active release
path

The runtime table contains 501 temperature values and 57 density values over

$$3.2 \times 10^3 \leq T \leq 10^8 \text{ K}, \quad 10^{12} \leq n_{\text{H}} \leq 10^{26} \text{ m}^{-3}.$$

It was generated as classical pure hydrogen with excitation, Fermi, and Coulomb corrections disabled. The parser rejects missing or unknown metadata, malformed dimensions, nonrectangular axes, nonpositive Γ_1 , and unsupported schemas.

The validation program compares all 28557 production nodes with the independent C++ Saha and analytic caloric closure. Key maximum discrepancies are:

Check	Maximum discrepancy
C++ Saha vs. CRASH transition relative difference	1.74×10^{-3}
Analytic internal energy vs. CRASH relative difference	1.14×10^{-6}
Analytic effective heat capacity vs. CRASH relative difference	8.26×10^{-4}
Independent analytic Γ_1 vs. raw CRASH absolute difference	5.18×10^{-4}
Runtime table node vs. raw CRASH Γ_1	1.40×10^{-14}
Coarse bilinear interpolation vs. direct CRASH at all 28000 cell midpoints	5.37×10^{-3} absolute
Caloric inversion against direct CRASH energy	3.50×10^{-7} relative

The table approaches 5/3 in neutral and fully ionized limits. Continuity probes across every internal table grid line found a maximum jump of about 1.93×10^{-11} . Small adiabatic perturbations through 10%, 50%, and 90% ionization reproduced $c_s^2 = \Gamma_1 p / \rho$ within 5×10^{-4} relative.

16.2 State and closure semantics

Tests pack (ρ, v, T) into the release conserved state, decode it back, and verify that density, temperature, internal energy, pressure, the Saha fraction, and the derived carrier densities n_e , n_{HI} , ρ_i , ρ_n are all recovered from the three stored rows alone, including the identity $e_{\text{int}} - \frac{3}{2}p = xn_{\text{H}}\chi_{\text{H}}$.

Non-physical states — negative density, an energy below the gravitational plus kinetic carry, or a density outside the Γ_1 table — are required to throw rather than be silently repaired. Separate tests assert that the legacy two-fluid entry points refuse a release Grid and that the release entry points refuse a Grid without a Γ_1 table.

16.3 Hydrostatic preservation and boundaries

The Saha-HSE initializer is checked against the trapezoidal logarithmic-pressure relation. All cells and four ghosts are decoded through the release EOS and checked for consistent pressure, composition, energy, and gravitational potential. Five reference-free release-mesh steps on the initialized column retain $\max |v| = 5.8 \times 10^{-3} \text{ m s}^{-1}$, $\max |\rho v| = 4.2 \times 10^{-13}$ and $\|\Delta U\|/\|U\| = 5.1 \times 10^{-8}$; the assertion is tied to the release mesh because, with no frozen reference, the bound is resolution dependent. That dependence is resolution dependence *in the ordinary second-order sense*, not a floor: the six-rung ladder of section 11.4 shows $\max |v|$ falling at unbroken order ≈ 2 from 8.11×10^{-1} to $3.32 \times 10^{-4} \text{ m s}^{-1}$ over a $41\times$ refinement, so the coarse-grid value quoted for contrast (0.76 m s^{-1} on a $14\times$ coarser grid) is simply the convergent trend evaluated far from the release mesh, and the release-mesh bound sits on that trend. Under the retired single-precision storage the same ladder flattened and then reversed, which is what made the earlier reading — “the bound is resolution dependent, and that is the real cost of retiring `eq_wb`” — a measurement of a round-off floor rather than of the scheme. Refined-mesh tests confirm that HSE ghosts use the actual local boundary spacing.

The hydro-temperature-decoupled upper boundary is tested to ensure that the hydrodynamic ghosts carry the live top temperature while the conduction solver continues to use the fixed external wall. Diagnostics independently reconstruct the accepted outer conductive face coefficient and flux.

16.4 Conduction conservation and nonlinear convergence

Closed-domain uniform and irregular-grid problems verify conservation of the physically weighted energy $\sum_i E_i \Delta s_i / B_i$, and mass and momentum are required to be bit-preserved by the stage. The accepted packed state is checked against a separately recomputed full backward-Euler residual. A strong chromosphere-to-corona profile, which drives $\kappa_e \propto T^{5/2}$ through five decades, must converge without leaving the EOS table. A dedicated regression asserts that enabling every retired two-fluid stage flag on a release Grid reproduces the release step bit for bit, which is the executable statement that those stages are absent from the release timestep rather than merely disabled.

16.5 Irregular-grid consistency

Manufactured and regression tests cover metric caches, reconstruction gradients, pressure operators, acoustic CFL scaling, series-resistance conduction, HSE boundaries, local artificial diffusivity, and the uniform-grid compatibility branch. Mesh tests verify positive widths, exact endpoints, refinement-only cell addition, the 1.1 adjacent-width cap, and outer refinement extending through the last face.

16.6 Serial/OpenMP equivalence

The same release case has been run with a serial build, an OpenMP build restricted to one thread, repeated 12-thread runs, and the selected CFL. The main state and gamma diagnostic outputs were byte-identical after excluding the output path embedded in the header. This establishes deterministic execution for the tested path.

16.7 CFL and duration validation

CFL values 0.125, 0.25, 0.35, 0.40, 0.45, and 0.50 were compared on matched physical times. All candidates passed non-finite-state, EOS-domain, Newton-convergence, limiter, and termination gates. Thermodynamic and conserved-field errors at CFL 0.50 remained orders of magnitude inside the predefined tolerances relative to CFL 0.125, and medium-duration differences remained smaller than the already accepted 0.25-versus-0.125 gap.

Historical That sweep was run with the pre-cutover single-precision state, and it recorded a serious standing finding: the residual velocity field *did not converge in time at all*. Refining the timestep made the level-to-level difference *larger*, and measured against the finest run every CFL from 0.125 to 0.50 sat at the same $L^1|\Delta v| \approx 1.0\text{--}1.2 \text{ ms}^{-1}$, so a velocity-relative comparison against the finest run could not discriminate between CFL values at all. It was hypothesised to be single-precision cancellation and left open as a production risk.

Active release path The hypothesis was correct, and the precision cutover resolves it. Re-measured on the release mesh ($N = 500/\text{R4}$, 661 cells, 20 s, conduction on, release Godunov flux), with every run interpolated in time to exactly $t = 20.000 \text{ s}$ from a 0.02 s snapshot cadence:

Δt -halving pair	double $L^1 \Delta v $	order	double $L^\infty \Delta v $	order	single $L^1 \Delta v $	order
0.50 vs 0.25	5.340×10^{-3}	—	6.398×10^{-2}	—	1.476×10^{-1}	—
0.25 vs 0.125	1.409×10^{-3}	1.92	2.610×10^{-2}	1.29	3.727×10^{-1}	−1.34
0.125 vs 0.0625	4.537×10^{-4}	1.64	1.200×10^{-2}	1.12	8.284×10^{-1}	−1.15

Differenced against the finest run (CFL 0.0625), double $L^1|\Delta v|$ is monotone from 7.157×10^{-3} at CFL 0.50 down to $4.537 \times 10^{-4} \text{ ms}^{-1}$ at CFL 0.125, while single precision stays pinned at $1.03\text{--}1.20 \text{ ms}^{-1}$ at every CFL — this section’s own historical finding, reproduced quantitatively on the current release configuration. **The single-precision control at the identical configuration is what makes the attribution causal rather than circumstantial.** The mechanism is visible directly in the fields at $t = 20 \text{ s}$: across an $8 \times$ step-count range double is stable (mean $|v|$ $2.4015 \rightarrow 2.4084 \text{ ms}^{-1}$, max $40.641 \rightarrow 40.661$) while single precision *degrades as the step count grows* (mean $|v|$ $2.418 \rightarrow 3.271$, max $40.71 \rightarrow 44.39 \text{ ms}^{-1}$). **Error that grows with the number of additions is round-off accumulation, not truncation.** Repeated on the original 2638-cell mesh the conclusion holds and the single-precision runaway is worse (mean $|v|$ $2.698 \rightarrow 6.304 \text{ ms}^{-1}$), with double orders 1.45 and 1.14 in L^1 .

Active release path Bounded precisely: what is retired is the mechanism, not the caution. Three limits on how far this result may be carried.

1. **The CFL acceptance matrix has not been re-run in double.** CFL 0.50 therefore still stands on its original single-precision evidence for this model shape and hardware, and the gate failures recorded in that sweep are not re-adjudicated here.

2. **No positive noise floor has been established for m s^{-1} -scale chromospheric velocities.** What is gone is the $\approx 1 \text{ m s}^{-1}$ storage-imposed floor; quantitative claims at that scale still require their own measurement.
3. **This is not a bit-level replay.** The historical sweep’s exact configuration is not reproducible on a release build — it used `ISO_NUMERICAL_DIFFUSIVITY_MULT`, which `model_column` now rejects, and the Roe flux, now demoted to a reference override. These are the current release configuration on two meshes with matched single-precision controls. The time alignment was also load-bearing rather than cosmetic: runs overshoot `CHROMO_T_END` by up to one Δt , and at $|dv/dt| \approx 0.14$ that misalignment alone injects $\sim 7 \times 10^{-4} \text{ m s}^{-1}$ — the size of the entire `double` signal.

The consequence that does follow is methodological: in `double` a velocity self-convergence ladder is once again a valid discriminator between CFL values, which it was not before.

Chapter 17

Development Experiments and Their Disposition

Each experiment below records the question, the controlled change, the result, and the final disposition. Metrics are limited to those needed to support the conclusion.

17.1 Constant- γ baselines

Question. Could hydrodynamics, boundaries, and conduction be developed before the equilibrium EOS was complete?

Controlled change. Use fixed $\gamma = 5/3$, then $\gamma = 1.05$ as a nearly isothermal surrogate.

Result. These modes enabled isolation of hydrostatic and boundary errors, and $\gamma = 1.05$ reduced excessive pressure response. Neither includes hydrogen ionization energy or the correct acoustic derivative.

Disposition. Historical Retained for regression and closure comparison; not a physical substitute for the release EOS.

17.2 Kinetic-ionization experiments

Question. Could finite-rate ionization and recombination support a dynamic chromosphere with independent fluids?

Controlled change. Activate collisional, photoionization, radiative, and three-body channels in the legacy source sequence.

Result. The experiments exposed strong energy-balance and upper-boundary sensitivity, including collapse when ionization cooling was not balanced by a coronal energy source. They also motivated robust trace-species and drift handling.

Disposition. Implemented but inactive Technically preserved; not the release closure. Non-equilibrium ionization remains scientifically important but requires a separately validated model.

17.3 Transition to Saha/ Γ_1

Question. How can ionization energy and acoustic response be included without an inconsistent partial EOS conversion?

Controlled change. Stage the work through table validation, analytic caloric closure, safeguarded inversion, manifold-consistent face reconstruction, variable-EOS conduction/sources, and Saha-HSE boundaries.

Result. Every active operator now consumes the same equilibrium thermodynamics; incompatible legacy modes fail closed.

Disposition. Active release path Current physical core.

17.4 PCHIP correction

Question. Why did the initial conductive flux contain steps at C7 knots?

Controlled change. Replace piecewise-linear temperature interpolation with shape-preserving C^1 cubic Hermite interpolation, then reintegrate HSE.

Result. The initialization retained every C7 knot while removing derivative discontinuities and producing a smooth initial physical heat flux.

Disposition. Active release path Gamma-table initialization; fixed-gamma interpolation intentionally unchanged.

17.5 Lower split conduction-to-hydro artifact

Question. Was a small low-chromosphere mass-flux bump a physical evaporation signal?

Controlled change. Hold the well-balancing configuration of the time fixed (equilibrium-reference, then the release default) and refine only the lower 0–700 km band.

Result. The bump scaled approximately linearly with local cell width and extrapolated close to zero. It arose when the explicit hydro step differentiated a pressure field modified by the split conduction stage.

Disposition. Historical Numerical coupling artifact; local lower refinement or tighter hydro-conduction coupling is the remedy. It is distinct from the upper-TR release study.

17.6 Upper-boundary diagnosis

Question. Was the top-cell reversal caused by an over-specified hydrodynamic boundary?

Controlled change. Decouple hydro ghost temperature from the fixed conductive wall while preserving fixed back pressure and velocity cap.

Result. The last-cell reversal disappeared, but the 2130–2150 km ρv ripple remained. This separated a boundary over-specification problem from an interior truncation-error problem.

Disposition. Active release path Decoupled boundary retained; ripple investigated separately.

17.7 Face-flux diagnosis

Question. Did the ρv ripple violate discrete continuity?

Controlled change. Record the central, Rusanov-diffusive, total, reference, and effective face mass fluxes without changing the state.

Result. F_{eff} remained smooth. Central and diffusive components were strongly anti-correlated, while cell-centered ρv oscillated around the conserved flux.

Disposition. Use F_{eff} , not cell-centered ρv , for quantitative transport at the upper boundary.

17.8 Uniform resolution scan

Question. Does the visible upper-TR ripple converge?

Controlled change. Run uniform $N = 500, 1000, 2000$ meshes with the same EOS, wall, boundary, limiter, and well balancing.

Result. Ripple and cell-to-face mismatch decreased at approximately first order. The mean effective mass flux did not converge because the realized conductive input also changed with resolution.

Disposition. Ripple classified as spatial truncation error; evaporation magnitude remained provisional.

17.9 Static outer refinement

Question. Can the upper ripple be reduced without globally refining the whole column?

Controlled change. Add a factor-two and factor-four outer refined band at coarse-equivalent $N = 1000$.

Result. The normalized ripple amplitude fell from 0.1366 to 0.0514 to 0.0094 while the factor-four mesh used 1320 rather than 4000 cells. The grading interface introduced no localized artifact.

Disposition. Active release path Outer static refinement adopted. The first study's globally coarse-scaled artificial conductivity was superseded.

17.10 Face-local numerical conduction

Question. Why did refinement increase the evaporation rate in the first outer-refinement study?

Controlled change. Replace the global coarse-scaled artificial diffusivity with $\chi_{\text{num},f} = C_{\text{num}} \Delta s_f$.

Result. Uniform results were unchanged; the refined outer diffusivity became one quarter of the coarse value; the refinement-induced flux increase disappeared; the artificial fraction of the outer conductive input fell with resolution.

Disposition. Historical Retained for controlled legacy comparisons; superseded in production by physical conduction only.

17.11 Historical N=1000/2000/4000 convergence evidence

Historical This superseded local-scaling sequence used the old nonzero artificial-conduction coefficient and ghost-center thermal boundary. It remains useful for explaining the transition to local scaling, but it is not release evidence:

Metric at 500 s	N=1000	N=2000	N=4000
Actual cells	1320	2638	5275
Fine width [m]	138.02	69.10	34.55
Total outer conductive input [W m^{-2}]	1.4768	0.9808	0.7938
Artificial fraction	0.2308	0.1289	0.0685
Mean F_{eff} [$\text{kg m}^{-2} \text{s}^{-1}$]	1.0566×10^{-9}	6.1004×10^{-10}	6.2250×10^{-10}
Cell-to-face mismatch amplitude	0.0250	0.0101	0.00253

The mass flux changed by -42.3% and then $+2.04\%$; the sequence is nonmonotone and does not support an apparent order or Richardson limit. The difference between the last two values is 35.8 times smaller than the difference between the first two, which is evidence of approaching convergence. However, total conductive input still changed by -19.1% on the last refinement. The limiting evaporation rate therefore remains provisional.

The outer thermal layer was not a one-cell feature at the highest resolution: it was about 10 km thick and spanned roughly 300 fine cells. The maximum temperature gradient lay inside the domain rather than at the wall. In that retired implementation, the wall gradient and conductive input remained first-order dominated by the ghost-center Dirichlet stencil. The release now places the datum at the physical face and uses the half-cell distance in equation (9.4).

Chapter 18

Experimental Two-Temperature Test Study (`model_column_2t`)

Experimental Everything in this chapter is a test study. None of it is release physics, none of it changes a release default, and no `model_column` command line can reach it. It is a third solver, `src/two_temp/`, selected only by the separate scenario `model_column_2t`. The release solver of `src/single_fluid/` is untouched by it and produces identical results with or without it compiled in. Its validation is a smoke test and its physical interpretation is not established; nothing here may be quoted as a property of the release.

18.1 Motivation and the question being asked

The release advances a single-fluid equilibrium mixture with *one* temperature shared by electrons, protons, and neutral hydrogen (chapter 3). That is an imposed closure, not a derived one: nothing in the release solver measures how well it holds. The upper part of the modelled column is exactly where one would expect it to be worst — the conductive energy that drives the calculation arrives from above through the electron channel, the gas there is fully ionized and tenuous, and the electron–heavy equilibration time grows as the density falls. This study asks one question and is scoped to it:

How far do the electron and heavy-particle temperatures actually drift apart in the modelled 1600–2153 km column, and is the release’s single-temperature assumption safe there?

The design is a controlled experiment. `model_column_2t` takes the release scenario’s atmosphere, mesh, gravity, gravity gauge, Γ_1 table, refinement, CFL, boundary geometry, and initial temperature profile *unchanged*, by calling the same initial-condition and boundary-update implementation; the only controlled variable against `model_column` is the temperature split. Every cell and every ghost is initialized with $T_e = T_i$ equal to the release temperature, so any later separation is generated by the solver rather than imposed by the initial data.

A common initial temperature is also the physically correct starting state, and the same measurement fixes the expected regime. Taking the electron–heavy energy-equilibration time as $\tau_{eq} = 1/(\nu_{ei} + \nu_{en})$ on the canonical release column (`util/tau_ei_equilibration.py`, $N = 500/R4$, 661 cells):

Location	1600 km (base)	1850 km (mid)	2153 km (top)
τ_{eq} at $t = 0$ [s]	9.19×10^{-5}	1.62×10^{-4}	4.24×10^{-3}
τ_{eq} at $t = 4000$ s [s]	9.20×10^{-5}	1.63×10^{-4}	3.92×10^{-3}
$\tau_{\text{eq}}/\Delta t$	0.016	0.029	0.70

against the release/two-temperature acoustic step $\Delta t \approx 5.5\text{--}5.6 \times 10^{-3}$ s. Two things follow. The exchange is stiff — by one to two orders of magnitude through the bulk of the column — so it must be treated implicitly if the acoustic CFL condition is to remain the only timestep constraint. And it is *not* uniformly stiff: in the fully ionized topmost cells τ_{eq} is a substantial fraction of the step, which is precisely the region where a measurable decoupling is plausible and where the release assumption is least self-evidently safe. The purpose of the study is to replace that expectation with a measurement, and section 18.10 reports it — including the confirmation that the top cell is indeed where the largest decoupling occurs.

18.2 Governing equations

Experimental One fluid, one velocity, two temperatures, on the same prescribed flux tube of cross-section $A(s)$ with $A(s)B(s) = \text{const}$ as in equation (3.1). Electrons are massless and comoving — quasi-neutrality together with zero field-aligned current — so there is still exactly one momentum equation and no ambipolar term. Protons and neutral hydrogen are lumped into a single *heavy* population at temperature T_i ; they exchange energy by charge exchange and elastic collisions on a timescale far shorter than anything resolved here, and the release has no ion–neutral drift either.

The conserved state is four rows,

$$U = (\rho, \rho v, \mathcal{E}, \mathcal{E}_e)^\top, \quad (18.1)$$

and the system advanced is

$$\frac{\partial \rho}{\partial t} + \frac{1}{A} \frac{\partial}{\partial s}(A \rho v) = 0, \quad (18.2)$$

$$\frac{\partial(\rho v)}{\partial t} + \frac{1}{A} \frac{\partial}{\partial s}[A(\rho v^2 + p)] = p \frac{\partial \ln A}{\partial s} - \rho \frac{\partial \Phi}{\partial s}, \quad (18.3)$$

$$\frac{\partial \mathcal{E}}{\partial t} + \frac{1}{A} \frac{\partial}{\partial s}\{A[v(\mathcal{E} + p) + q_e + q_h]\} = 0, \quad (18.4)$$

$$\frac{\partial \mathcal{E}_e}{\partial t} + \frac{1}{A} \frac{\partial}{\partial s}(A \mathcal{E}_e v) = -p_e \frac{1}{A} \frac{\partial}{\partial s}(A v) - \frac{1}{A} \frac{\partial}{\partial s}(A q_e) + Q_{ei}. \quad (18.5)$$

The reported runs use the constant-area specialization $A = \text{const}$, $B \equiv 1$, exactly as the release does.

$\mathcal{E}$ keeps its release meaning: it is the TOTAL energy. Mass, momentum, and total energy therefore stay in exact conservation form, the release numerical flux applies to those three rows verbatim, and the heavy internal energy is a *derived* quantity,

$$e_h = \mathcal{E} - \mathcal{E}_e - \frac{1}{2} \rho v^2 - \rho \Phi, \quad T_i = \frac{2}{3} \frac{e_h}{n_H k_B}, \quad (18.6)$$

which makes the heavy temperature purely algebraic — no caloric inversion at all. Only T_e requires an inversion, of $\mathcal{E}_e(\rho, T_e)$. This “total energy plus electron energy” split is the standard two-temperature MHD formulation of the target framework [van der Holst et al., 2010, Sokolov et al., 2021] and is the same split this project’s historical two-fluid **E_E** row already used. Its practical consequence is that the electron equation, which is *not* in conservation form, cannot damage the total: any splitting error is absorbed by the heavy pool by construction.

18.3 Thermodynamic closure

Experimental With $n_H = \rho/m_H$ the heavy-particle density — protons and neutrals together always number n_H , independent of the ionization fraction — the closure is

$$x = x_{\text{Saha}}(n_H, T_e), \quad n_e = xn_H, \quad n_{\text{HI}} = (1 - x)n_H, \quad (18.7)$$

$$p_e = n_e k_B T_e, \quad p_h = n_H k_B T_i, \quad p = p_e + p_h = n_H k_B (T_i + x T_e), \quad (18.8)$$

$$\mathcal{E}_e = \frac{3}{2} p_e + x n_H \chi_H, \quad e_h = \frac{3}{2} p_h, \quad \mathcal{E} = e_h + \mathcal{E}_e + \frac{1}{2} \rho v^2 + \rho \Phi. \quad (18.9)$$

Here x_{Saha} is the same Saha relation as equation (3.7), evaluated at the *electron* temperature.

Two modelling choices are made in equations (18.7) and (18.9), and they are the same choice twice. Hydrogen ionization and recombination in this regime are electron-impact processes: the energy χ_H per ionization is taken from, and returned to, the electron thermal reservoir. The ionization reservoir is therefore assigned to the electron pool, and for the same reason the ionization balance is evaluated at T_e rather than at T_i or at some mixture temperature.

Exact degeneracy to the release closure. Setting $T_e = T_i = T$ in equations (18.8) and (18.9) gives $p = (1 + x)n_H k_B T$ and $e_h + \mathcal{E}_e = \frac{3}{2} p + x n_H \chi_H = e_{\text{int}}$, which are exactly equations (3.8) and (3.9). The two-temperature model is a strict generalization of the release closure, not a different one.

The implementation makes that degeneracy structural rather than coincidental. The release caloric closure $e_{\text{int}} = \frac{3}{2}(1 + x)n_H k_B T + x n_H \chi_H$ has as its “1” precisely the heavy translational term $\frac{3}{2} n_H k_B T$, so the electron pool and its heat capacity are formed as exact *differences* of the public release EOS functions,

$$\mathcal{E}_e(\rho, T_e) = e_{\text{int}}(\rho, T_e) - \frac{3}{2} n_H k_B T_e, \quad \frac{\partial \mathcal{E}_e}{\partial T_e} = C_V^{\text{eq}}(\rho, T_e) - \frac{3}{2} n_H k_B, \quad (18.10)$$

while the heavy capacity is the constant $\frac{3}{2} n_H k_B$ — the heavy pool carries no ionization energy. The two solvers therefore cannot drift apart in their thermodynamics, and at $T_e = T_i$ the two pools re-sum to e_{int} to within one floating-point subtraction and addition. $\mathcal{E}_e$ is strictly increasing in T_e (both $\frac{3}{2} p_e$ and $x n_H \chi_H$ are), so the EOS table’s temperature range is an exact bracket and the T_e inversion is unconditionally bracketable. It uses safeguarded Newton on equation (18.10) with *guaranteed bisection progress*, whose derivative is dominated by the dx/dT_e ionization term across the partial-ionization zone. That guarantee is not a formality; see section 18.7.

18.4 Acoustics: the release flux carries over unchanged

Experimental Both species are monatomic. At frozen composition the total internal energy therefore still splits exactly as

$$e_{\text{int}} = e_h + \mathcal{E}_e = \frac{3}{2} p_h + \frac{3}{2} p_e + e_{\text{ion}} = \frac{p}{5/3 - 1} + e_{\text{ion}}, \quad e_{\text{ion}} = x n_H \chi_H, \quad (18.11)$$

independently of how the thermal energy is partitioned between the electron and heavy pools. Equation (18.11) is equation (11.20) with the translational term written as a sum; splitting the temperature does not touch it. Three consequences follow, and they are a result of the formulation rather than a convenience:

1. **The release numerical flux applies unchanged.** The SWMF-style exact-Riemann Godunov flux of section 11.3.2 — an exact ideal-gas Riemann solve at fixed $\gamma = 5/3$ with all non-ideal

energy carried in the passive specific offset E_0 of equation (11.18) — describes the acoustic waves of the two-temperature system exactly as it describes those of the release system. The same port of `ModExactRS.f90` is used, with the same face-local Rusanov fallback and the same per-face fallback accounting.

2. **The frozen sound speed and the acoustic CFL condition are unchanged.** $c_{\text{fr}}^2 = (5/3)p/\rho$ depends only on the *total* pressure, so the timestep rule is byte-for-byte the release Godunov rule of equations (12.3) and (12.4). The stiff collisional exchange imposes no explicit limit because it is solved implicitly (section 18.5).
3. **The extra physics rides on a doubled contact characteristic.** The eigenvalues of the 4×4 system are $\{u - c, u, u, u + c\}$: the acoustic pair is untouched, and the second $\lambda = u$ carries the total entropy *and* the electron/heavy energy partition. That is the entire additional hyperbolic content of the model, and it is what fixes the boundary conditions of section 18.6.

Correspondingly, the numerical flux upwinds the specific electron energy $\mathcal{E}_e/\rho$ on the contact and multiplies it by the *exact* mass flux $\rho_* u_*$. One upwind decision, $\text{sign}(u_*)$ with SWMF’s tie-break to the right state at $u_* = 0$, serves both the passive offset E_0 and $\mathcal{E}_e/\rho$, so the total-energy flux and the electron-energy flux can never disagree about which side’s composition crossed the face.

The electron row is advanced in the advective form of equation (18.5): a mass-consistent flux plus a compression-work source $-p_e A^{-1} \partial_s(Av)$ evaluated on the face velocities of the stage that is calling. This is the standard non-conservative electron-pressure treatment of the target framework [van der Holst et al., 2010, Sokolov et al., 2021]. Because the total-energy row stays conservative and e_h is obtained by difference (equation (18.6)), total energy is conserved exactly and the splitting error lands in the heavy pool.

18.5 Separate heat conduction and collisional exchange

Experimental The two conductive channels are separated by *which temperature they act on*, and by nothing else:

$$q_e = -\kappa_e(n_e, n_{\text{HI}}, T_e) \frac{\partial T_e}{\partial s}, \quad q_h = -\kappa_n(n_e, n_{\text{HI}}, T_i) \frac{\partial T_i}{\partial s}. \quad (18.12)$$

κ_e is the same Spitzer-limit electron coefficient with electron-neutral collisions that the release uses [Spitzer, 1962, Vranjes and Krstic, 2013], now evaluated at T_e ; κ_n is the same neutral-hydrogen coefficient, now evaluated at T_i . The expressions are transcribed unchanged so that `src/two_temp/` depends only on the Grid and the shared EOS.

Spitzer proton conduction is deliberately omitted, and the omission is what makes the experiment controlled. For hydrogen the ion Spitzer conductivity is $\sim 2.3\%$ of κ_e — it scales as $(m_e/m_p)^{1/2}$ — so including it would change little physically. Omitting it, however, makes $\kappa_e + \kappa_n$ reproduce the release total conductivity *exactly* at $T_e = T_i$. The experiment therefore changes only which temperature each existing channel acts on, and any difference from the release cannot be attributed to a new transport coefficient. It is a real, if small, omission and is listed as such in section 18.8.

The two pools are coupled pointwise by

$$Q_{ei} = g_{ei}(T_i - T_e), \quad g_{ei} = \frac{3}{2}n_e k_B (\nu_{ei} + \nu_{en}), \quad \tau_{\text{eq}} = \frac{1}{\nu_{ei} + \nu_{en}}, \quad (18.13)$$

positive Q_{ei} heating the electrons. ν_{ei} is the Spitzer/NRL electron–ion energy-equilibration frequency with the Coulomb logarithm evaluated per cell and clamped to $[5, 30]$ [Spitzer, 1962]; ν_{en}

is the electron–neutral rate built on the Vranjes & Krstic cross-section [Vranjes and Krstic, 2013] with its $2m_e/m_H$ energy-transfer factor. Both are the same expressions the two-fluid research solver uses. The conductance is built on the *translational* electron heat capacity $\frac{3}{2}n_e k_B$ and not on the ionization-inflated $\partial\mathcal{E}_e/\partial T_e$, because the frequencies are defined against $dT_e/dt = \nu(T_i - T_e)$.

Both collision channels are kept, and in this model that matters. The lower chromosphere is only weakly ionized: at 1600 km the release column has $x \approx 0.27$, so $n_{\text{HI}} \gg n_e$ and electron–neutral collisions carry a large share of the electron thermalization. A purely Coulombic exchange would be the wrong closure for the bottom two-thirds of the domain. Conversely, at the fully ionized top $n_{\text{HI}} \rightarrow 0$ and equation (18.13) reduces to the Coulomb rate alone — which is why τ_{eq} rises by a factor of ~ 25 between 1850 km and the top in section 18.1.

18.6 Boundary conditions from the characteristic budget

Experimental Adding a temperature adds boundary data, and where that data may be imposed is not a free choice. The hyperbolic and the parabolic parts of the operator have *separate* boundary budgets, and conflating them is the standard way to over-specify a two-temperature outflow face. This section sets out both explicitly.

18.6.1 The hyperbolic budget: both temperatures must free-float at the top

From section 18.4 the eigenvalues of the 4×4 system are $\{u - c, u, u, u + c\}$. Both temperatures ride on $\lambda = u$; neither rides on an acoustic characteristic. At the outer face this column is a slow subsonic *outflow*, $0 < u \ll c$, so three characteristics leave the domain and exactly one, $u - c$, enters. **Exactly one condition may be imposed from outside, and the release already spends it on the fixed reservoir back pressure p_{ref}** (section 9.4). No hydrodynamic temperature datum is available at that face.

Both temperatures are therefore zero-gradient extrapolated from the live top cell into the two outer hydro ghosts, and the ghost density is obtained by inverting the two-temperature pressure closure $p = n_{\text{H}} k_B (T_i + x(n_{\text{H}}, T_e) T_e)$ at that pair of temperatures — monotone in n_{H} , so the EOS table’s density axis is an exact bracket. The ghost velocity is the live top velocity Mach-capped at $0.1c_s$ as in the release, with c_s taken here as the *frozen* speed $\sqrt{(5/3)p/\rho}$, which is the speed this solver’s flux and timestep actually use. Were the top flow to reverse to subsonic inflow, three characteristics would enter and both temperatures *would* have to be supplied alongside the pressure; the Mach-capped outflow closure is what keeps that from happening.

18.6.2 The parabolic budget: one condition per channel per face

The conduction stage is a separate operator with its own, parabolic, boundary requirement: one condition for each temperature at each face. Here the physics chooses them.

Outer face, electron channel — Dirichlet at 22000 K. The downward conductive flux from the transition region and corona is *electron*-conducted. Above the domain top hydrogen is fully ionized, so $\kappa_n \propto n_{\text{HI}}$ collapses while κ_e is the Spitzer $T^{5/2}$ channel. Assigning the release’s 22000 K conductive wall to the electron channel is what makes the heating chain of this model explicit:

TR conduction $\rightarrow$ electrons $\rightarrow$ collisional exchange $\rightarrow$ heavies $\rightarrow$ pressure $\rightarrow$ evaporation.

As in the release, the datum sits at the *physical* outer face, a half top-cell width from the top-cell centre (equation (9.4)), and the wall material state is rebuilt at 22000 K from the

imposed back pressure. At that temperature the gas is fully ionized and $T_e = T_i$, so using the equilibrium pressure inverse to build the wall state is exact there rather than an approximation.

Outer face, heavy channel — Neumann, zero conductive flux. There is no heavy-particle conductive flux to import from a fully ionized corona: the neutral channel has no carrier, and Spitzer proton conduction, which this model does not carry, would be only $\sim 2.3\%$ of κ_e . Imposing a heavy Dirichlet datum instead would manufacture a spurious neutral-channel flux at a face where n_{HI} has already gone to zero. This is the default.

Inner face, both channels — Dirichlet at the truncation temperature. The 1600 km reservoir is dense and cool and τ_{eq} there is 9.2×10^{-5} s (section 18.1), so $T_e = T_i = T_0$ is a statement about the reservoir rather than an extra modelling assumption. The release knob `ISO_INNER_T_NEUMANN` keeps its meaning and makes the base insulating for *both* channels.

Face	T_e	T_i	Rationale
Outer, hydrodynamic	Zero-gradient from top cell	Zero-gradient from top cell	One incoming characteristic, already spent on p_{ref}
Outer, conduction	Dirichlet, 22000 K at the physical face	Neumann, zero flux (default)	Coronal conduction is electron-carried; $\kappa_n \propto n_{\text{HI}} \rightarrow 0$ above the top
Inner, hydrodynamic	Reservoir T_0 , $v = 0$	Reservoir T_0 , $v = 0$	Quasi-static stratified reservoir (section 9.3)
Inner, conduction	Dirichlet T_0	Dirichlet T_0	$\tau_{\text{eq}} \approx 0.09$ ms: the base is collisionally equilibrated

18.6.3 The controlled comparison

The outer heavy-channel condition is the one genuinely arguable choice above, so it is exposed as the study’s controlled variable through the environment variable `TT_OUTER_TI`, which takes `neumann` (the default and the physically motivated case) or `dirichlet` (the comparison). It is the only setting varied anywhere in this study.

The comparison is deliberately the *naive* closure rather than a near-copy of the default: under `dirichlet` the same 22000 K wall temperature that the electron channel receives is imposed on the heavy channel as well, so the leg asks what would have happened had the heavy species been allowed to feel the coronal reservoir directly. Handing the heavy channel the extrapolated ghost temperature instead would have made the two legs differ only marginally and the test would have measured nothing. The result is section 18.10.6.

18.7 Numerics and implementation

Experimental The timestep is the release’s two-stage first-order Lie split of section 11.2.1 with the single scalar conduction solve replaced by a coupled 2×2 -block one:

$$U^n \xrightarrow{\text{explicit MUSCL-Hancock/Godunov hydrodynamics}} U^* \xrightarrow{\text{implicit coupled } (T_e, T_i) \text{ conduction+exchange}} U^{n+1}.$$

Explicit stage. The reconstruction adds exactly one slot to the release primitive set,

$$W = (\log \rho, v, \log p, \log T_e), \quad (18.14)$$

limited with the same asymmetric MC3/Koren form at $\beta = 2$ on the same nonuniform-mesh slope metrics. **The total pressure remains slot 2 and remains the authoritative mechanical variable**, so the new slot changes only the electron/heavy energy *partition* at a face, not the mechanics — the same reasoning that made $\log p$ the release reconstruction variable (section 11.2.1). Face states are then fully algebraic: $x = x_{\text{Saha}}(n_{\text{H}}, T_e)$, $p_e = x n_{\text{H}} k_B T_e$, $p_h = p - p_e$, $T_i = p_h / (n_{\text{H}} k_B)$, with no temperature inversion anywhere. If a reconstruction leaves $p_e \geq p$, which would imply $T_i \leq 0$, the *electron* temperature is reduced until $p_e = 0.999999 p$, because the total pressure is the authoritative variable and is preserved exactly; the event is counted, and a clamped face is admissible but no longer exactly on the Saha manifold. The MUSCL–Hancock predictor applies the same sources as the corrector — flux-tube pressure term, gravity, and the electron compression work — built on the stage’s own face velocities, exactly as equation (11.3) requires for the momentum source.

Implicit stage. The coupled backward-Euler problem is

$$\mathcal{E}_e(\rho, T_e) - \mathcal{E}_e^* = \Delta t [\nabla_s \cdot (\kappa_e \nabla_s T_e) + g_{ei} (T_i - T_e)], \quad (18.15)$$

$$e_h(\rho, T_i) - e_h^* = \Delta t [\nabla_s \cdot (\kappa_n \nabla_s T_i) - g_{ei} (T_i - T_e)], \quad (18.16)$$

solved together by quasi-Newton on the cell temperatures. The two conduction operators never couple across cells — each channel conducts on its own temperature — so the only inter-species coupling is the pointwise exchange, which lands entirely in the diagonal block. The Jacobian is therefore 2×2 -*block* tridiagonal with diagonal off-diagonal blocks, and is solved by a block Thomas algorithm; κ_e , κ_n , and g_{ei} are lagged, exactly as the release solve lags $d\kappa/dT$ (equation (12.1)). Both residuals are normalized on the heavy-energy scale, which is the larger of the two pools everywhere in this column, so the tolerance means the same thing for both rows; a vanishing Newton step with a live residual is reported as stagnation and reaching the iteration limit is an error, as in the release. The *accepted* energies are the backward-Euler targets rather than an EOS round-trip of the converged temperatures, so the exchange cancels exactly between the two pools, both conductive divergences are re-summed into the total-energy row, and the stage conserves total energy by construction while leaving mass and momentum untouched.

Caloric inversion near full ionization: a documented pitfall. The T_e inversion of equation (18.10) is safeguarded Newton in the `rtSAFE` sense: the iteration keeps a bracket in which $\mathcal{E}_e$ is monotone, and it falls back to bisection both when a Newton step would leave the bracket *and* when it fails to at least halve the bracket. **The second test is the one that matters, and omitting it is a real observed failure rather than a hypothetical.** The electron caloric curve has a very sharp knee at full ionization: as $x \rightarrow 1$ the capacity $\partial \mathcal{E}_e / \partial T_e$ collapses to the small translational value, while just below the knee the dx/dT_e term makes it larger by orders of magnitude. A Newton step taken from the flat side therefore overshoots the knee by thousands of kelvin. An inversion safeguarded *only* against leaving the bracket then oscillates between the two bracket ends indefinitely — each iterate lands strictly inside by $\sim 10^{-10}$ K, so the out-of-bracket test never fires and the bracket shrinks by $\sim 10^{-10}$ K per pass. This aborted a 4000 s run during the present study at $\rho = 8.26 \times 10^{-11} \text{ kg m}^{-3}$, $x = 0.999965$, and was fixed by adding the bracket-halving test. It is recorded here because any closure that puts an ionization reservoir in one pool inherits the same knee, and a bracket-only safeguard on it is not sufficient.

Guards, counters, and quotability. Every guard the solver can trip is counted rather than silently applied, and the driver prints a `twotemp.*` tally at the end of every run: `max_relative_decoupling` (the headline measurement $\max |T_e - T_i|/T_i$ over all cells and all steps), `godunov_faces` and `godunov_fallbacks` for the exact Riemann solve, the three clamp counters (`face_pe`, `decode_Te`, `decode_Ti`) and their sum, and the maximum Newton pass count of the conduction solve. **A run with a nonzero clamp count is not quotable** and the driver says so on `stderr`; a nonzero Godunov fallback count likewise raises a warning. The two-temperature solver keeps its own tally and never touches `Grid::godunov_stats`, so the release `godunov.*` block is not printed for it. One caveat on the headline number: it is accumulated inside the state decode, which is also applied to the MUSCL–Hancock predictor half-state, so the reported maximum includes intermediate predictor states and is a conservative upper bound on the decoupling of accepted states.

Output. With diagnostics enabled the run writes a `.twotemp` sidecar in the same header/frame-marker format as the release `.gamma_diag` stream, so the analysis scripts share a reader. Its fourteen columns are ρ , v , T_e , T_i , x , n_e , n_{HI} , p , p_e , p_h , κ_e , κ_n , Q_{ei} , and the local τ_{eq} that the implicit exchange stage actually used — which puts the stiffness of the coupling directly in the output. The main snapshot carries the four rows of equation (18.1) and its header line records the width, as for the other two solvers.

Isolation from the release. `src/two_temp/` shares only `chromosphere.hpp` (the Grid), `eos.hpp`, and the standalone SWMF Riemann-solver port `single_fluid/exact_rs.hpp`; it includes no other release or two-fluid header, and the release includes none of it. The solver’s entire mutable workspace is caller-owned and passed to each stage, so nothing about this experiment lives in the shared Grid beyond four ghost buffers and two flags. Loading a Γ_1 table still selects a Saha solver; *which* one is decided by the scenario’s initial condition, and the state width is validated against the selected solver before evolution begins. `model_column_2t` rejects `ISO_GAMMA` for the same reason `model_column` does, and additionally rejects `ISO_RIEMANN` and `ISO_RECONSTRUCTION`, because this solver has exactly one flux and one primitive set and a silently ignored setting would be worse than a refused one.

18.8 Deliberately absent physics and inherited caveats

Experimental The model is deliberately narrow, so that what it measures is attributable. Absent, and *not* behind a default-off flag:

- finite-rate ionization — the composition is Saha-equilibrium at T_e ;
- radiative cooling, TRAC, nonthermal beam heating, and volumetric coronal heating — as in the release, these are structurally absent rather than disabled;
- ion–neutral drift — protons and neutrals share one velocity and one temperature;
- an anisotropic or separate proton temperature — there is no $T_{\parallel}/T_{\perp}$ split and no third temperature;
- Spitzer proton conduction — omitted on purpose, see section 18.5;
- electron inertia, ambipolar diffusion, and field-aligned current — electrons are massless and comoving;
- Alfvén waves and any wave-pressure or turbulent-heating term.

The release’s largest open consistency question is inherited, not resolved. Section 20.4 records that the release freezes the composition inside the local Riemann problem while

the equation of state re-imposes instantaneous Saha equilibrium after every step, an operator-split instantaneous-relaxation approximation that is not justified when the physical relaxation time is long. **The two-temperature study carries exactly the same structure**, now with the equilibrium projection taken at T_e : a Saha-equilibrated state, a frozen-composition hydrodynamic step, and a Saha-equilibrated state. Separating T_e from T_i does not address that question and was not intended to; if anything it sharpens it, because assigning the ionization reservoir to the electron pool makes explicit that the quantity being instantaneously relaxed is an electron-impact process with its own rate. Any decoupling this study measures is measured *within* that approximation.

Two further limitations of scope should be stated. The study inherits the release’s resolution caveat — the $N = 500$ long-duration evaporation mass flux is not established as better than 10% grid converged — and it inherits the truncated-domain caveat that the 22000 K wall and the fixed back pressure are prescribed external data that do not respond to the column.

18.9 Validation status

Experimental A 5 s smoke run of the default (`neumann`) configuration on the canonical $N = 500/R4$ mesh completed with a maximum relative decoupling $\max |T_e - T_i|/T_i = 3.28 \times 10^{-3}$, zero clamps of any kind, and zero Godunov fallbacks. **That was a smoke run, not a result**, and it is recorded here only because it establishes that the solver runs on the release configuration without tripping a guard and that the exact Riemann solve is usable on the four-row state. The study is measured on the 4000 s production runs of section 18.10. As it happens the smoke run’s maximum survived into them unchanged, for a reason that is itself part of the result: the largest decoupling of the whole calculation occurs in the first few tens of milliseconds.

18.10 Results

Experimental Three 4000 s runs, all on the canonical $N = 500/R4$ mesh (661 cells) at CFL 0.50 with double storage and physical conduction only, all reporting `termination=end_time` at exactly $t = 4000$ s:

1. the two-temperature default, `TT_OUTER_TI=neumann`;
2. the two-temperature controlled comparison, `TT_OUTER_TI=dirichlet` (section 18.6.3);
3. a matched single-temperature release reference, run with the current binary at identical mesh, cadence, and CFL.

The release reference must be the matched run and no other. The earlier `lnp_godunov_N500_4000s_f64` output is the retired, incoherent `CHROMO_STATE_FLOAT64` container-only build (section 10.2) and reports a top velocity of 13.58 m s^{-1} , against 13.18 for the withdrawn single-precision figure and 9.50 for the true `double` release; it is not a valid comparand for anything.

18.10.1 Run integrity and the measured CFL equivalence

Both two-temperature legs are clean by the criteria of section 18.7: `twotemp.godunov_fallbacks=0` out of 9.413×10^8 face solves, `twotemp.clamps=0` with all three counters (`face_pe`, `decode_Te`, `decode_Ti`) individually zero, and `twotemp.conduction_iterations_max=4` — the coupled 2×2 -block Newton solve never needed more passes than the release’s scalar solve does. The runs are therefore quotable in full.

The step counts are 712031 for the two-temperature runs and 712029 for the release, a difference of two steps in 7.1×10^5 , or 3×10^{-6} . **This is the measured confirmation of the CFL-equivalence prediction of section 18.4:** the frozen signal speed depends only on the total pressure, so splitting the temperature cannot change Δt , and over four thousand seconds of evolution it does not.

18.10.2 Exact degeneracy of the closure at $T_e = T_i$

At $t = 0$, where the initial condition sets $T_e = T_i$ everywhere by construction, the two-temperature state and the release state are **bit-identical in all three release rows**: the maximum absolute difference is exactly 0.0 in ρ , ρv , and $\mathcal{E}$ across all 661 cells. The fourth row spans $\mathcal{E}_e = 0.0435$ – 0.367 J m^{-3} .

This is the strongest available check on the claim of section 18.3 that the two-temperature closure is a strict generalization of the release closure rather than a different one. The two solvers construct ρ , ρv , and $\mathcal{E}$ through entirely separate code paths, and at $T_e = T_i$ they land on the same doubles.

18.10.3 The measured decoupling: small, and concentrated in a startup transient

Global maximum. Over every cell and every one of 712031 steps,

$$\max_{\text{cell, step}} \frac{|T_e - T_i|}{T_i} = 3.28 \times 10^{-3}. \quad (18.17)$$

This is the same value the 5 s smoke run reported, which by itself proves that the global maximum occurs within the first five seconds and that nothing later in the 4000 s evolution approaches it. As noted in section 18.7, the counter also samples MUSCL–Hancock predictor half-states, so equation (18.17) is a conservative upper bound on the decoupling of accepted states rather than a tight one — which only strengthens the conclusion, since the bound is already small.

Where and when. A dense-cadence rerun (CHROMO_FRAME_DT=0.02) localizes the peak to $t \approx 0.022 \text{ s}$ in the topmost cell at 2153 km, where $T_e = 22339.3 \text{ K}$ sits 63.4 K *below* $T_i = 22402.8 \text{ K}$. **The sign is diagnostic and confirms the boundary construction of section 18.6.** The initial top cell is at 22674.8 K while the imposed electron wall is *colder*, at 22000 K, so on the first steps the electron channel alone conducts heat out of the top cell while the heavy channel is insulated — exactly the asymmetry the electron-Dirichlet/heavy-Neumann pair creates. The excess then decays roughly as $1/t$ as collisions re-lock the two pools: 43.1 K at 0.044 s, 18.0 K at 0.10 s, 5.0 K at 0.26 s, and 1.7 K at 0.5 s.

Quasi-steady state. Once the transient has passed, the decoupling is smaller by more than two orders of magnitude and still falling: $\max |T_e - T_i|/T_i = 2.05 \times 10^{-5}$ at $t = 20 \text{ s}$ and 9.01×10^{-6} at $t = 4000 \text{ s}$, the latter corresponding to $\max |T_e - T_i| = 0.0715 \text{ K}$ at 2138.4 km. **After startup the two temperatures agree to about one part in 10^5 .**

Final-frame structure, and it is physically interpretable. At 4000 s the profile of $T_e - T_i$ is smooth and monotone with a single physically meaningful sign change, at 2137.8–2138.4 km:

Below the crossing — through the partial-ionization zone, where x rises from 0.62 to 0.99 — $T_e < T_i$, by up to $1.1 \times 10^{-2} \text{ K}$. *The electron pool is paying for ionization:* it owns the χ_{H} reservoir (section 18.3), so every increment of ionization is drawn from the electrons and leaves them marginally cool relative to the heavies.

Above the crossing — where $x = 1$ and there is nothing left to ionize — $T_e > T_i$, by up to 7.1×10^{-2} K. *The electrons are being conductively heated from the boundary and are handing the energy on to the heavies*, which is the middle of the heating chain written out in section 18.6. The crossing is the point at which the ionization sink gives way to the conductive source, and the two-temperature model resolves the handover. An earlier reading of this profile as a cell-to-cell “grid-scale oscillation” was an artifact of the sidecar’s print floor and is withdrawn: there are two sign changes across 203 cells, not one per cell.

18.10.4 Why the temperatures lock: timescale separation

The mechanism is the one anticipated in section 18.1, and the 4000s run quantifies it on its own state rather than on the release column’s. Measured in the two-temperature run, the electron-heavy equilibration time runs from $\tau_{\text{eq}} = 9.25 \times 10^{-5}$ s at the 1600 km base to 3.97×10^{-3} s in the 2153 km top cell — within a per cent of the release-column values tabulated in section 18.1, which is itself a consistency check — against a timestep of $\Delta t \approx 5.6 \times 10^{-3}$ s and against a conduction/evaporation timescale of order 10^3 s. **That is five to seven orders of magnitude of separation between the collisional coupling and the physics of interest, and it is the whole explanation of the result.** Any energy imbalance the conduction or the ionization sink creates between the two pools is erased by collisions long before the evaporation flow has evolved at all; the only place a decoupling can survive is where something drives the pools apart faster than τ_{eq} , which in this model happens only at the very first steps, in the top cell, from the discontinuity between the initial temperature and the imposed wall.

18.10.5 Two-temperature against the matched release run

Comparing the `neumann` leg with the matched release reference at $t = 4000$ s:

Observable at 4000 s	Two-temperature	Release	Difference
Top-of-domain velocity [m s^{-1}]	9.4841	9.5008	−0.176 %
Top-of-domain mass flux ρv [$\text{kg m}^{-2} \text{s}^{-1}$]	2.69498×10^{-10}	2.69974×10^{-10}	−0.176 %
Top density	—	—	−0.0006 %
Total pressure, max relative difference over the column	—	—	3.11×10^{-5}

The evaporation observables agree to about 0.2 %, and the mechanical state agrees to three parts in 10^5 everywhere.

Two quantities look larger and must not be misread as bulk disagreement. The temperature difference reaches $\max |T_i - T_{\text{release}}| = 72.91$ K (8.05×10^{-3} relative), and $\max |T_e - T_{\text{release}}| = 72.92$ K — the two two-temperature temperatures track each other far more closely than either tracks the release value — with *both* maxima localized at 2138.7 km. Similarly, $\kappa_e + \kappa_n$ differs from the release total conductivity by at most 8.43×10^{-2} in relative terms, localized at 2138.1 km. **Both are the same feature: 2138 km is the ionization front and transition-region onset, where the vertical temperature gradient is extremely steep, so a sub-cell displacement of the front produces a large *local* difference in any quantity that varies sharply across it.** Away from the front the agreement is far tighter, and the 3×10^{-5} pressure agreement across the whole column is the direct evidence that the mechanical state is essentially identical. For the conductivity in particular the identity is exact at the formula level, not merely close: `two_temp_kappa_e` and `two_temp_kappa_i` are character-for-character the release `physical_kappa_e` and `physical_kappa_n` expressions (with $n_i = n_e$ for pure hydrogen), so $\kappa_e + \kappa_n$ reproduces the release conductivity exactly wherever $T_e = T_i$, by construction (section 18.5).

18.10.6 The outer heavy-channel condition is immaterial — a robustness result

With the corrected `dirichlet` leg — the full 22000 K wall imposed on T_i as well as on T_e — the two boundary treatments agree **to round-off**:

Quantity at 4000 s	Max. relative difference
T_e	4.7×10^{-10}
T_i	8.7×10^{-10}
ρ	3.5×10^{-10}
v	1.0×10^{-9}
p (total)	exactly 0

The top-cell velocity is 9.4841 ms^{-1} in both legs, to four decimals.

The physical reason is clean, and it is the argument of section 18.6 confirmed quantitatively rather than asserted. At the 22000 K wall the gas is fully ionized, so $n_{\text{HI}} \rightarrow 0$ and the neutral-hydrogen conductivity collapses: in the top cell $\kappa_e/\kappa_n = 7.6 \times 10^7$, i.e. the heavy channel carries about 1.3×10^{-8} of the electron channel’s conductivity. **There is simply no heavy conductive flux at that face for a boundary condition to control**, so it makes no difference what datum is imposed on it.

This should be read as a *positive* result rather than a failed experiment. The one genuinely arguable choice in the two-temperature boundary closure turns out not to matter at all for this model, which means the electron-conducted transition-region closure is robust to it, and it isolates the electron channel as the only one that carries energy into the domain from above. It also retires the concern that a naive closure — letting the heavy species feel the coronal reservoir directly — would have distorted the result: here it does not, because there is nothing for it to distort.

18.10.7 What this validates, and what it does not

For this model — a 1600–2153 km chromospheric column with physical conduction only and a 22000 K electron-conducted outer reservoir — the release’s single-temperature equilibrium-mixture assumption is validated, to about 0.2 % in the evaporation observables, with a worst-case transient decoupling of 0.33 % lasting tens of milliseconds and a quasi-steady decoupling of order 10^{-5} . The closure degenerates exactly at $T_e = T_i$ (section 18.10.2), the acoustic CFL equivalence predicted from the frozen-composition argument is confirmed to 3×10^{-6} , and the boundary choice that could have mattered demonstrably does not.

The scope of that statement must be stated as plainly as the result. This study measured one configuration. It does *not* license the single-temperature assumption in a flare, beam-heated, substantially hotter, or substantially more rarefied regime, where τ_{eq} rises — it scales as $T_e^{3/2}/n_e$ in the Coulomb channel — while the electron heat flux, scaling as $T_e^{5/2}$, becomes far larger. Both trends widen the same gap that this column closes in tens of milliseconds. A two-temperature treatment may well be required there, and the solver documented in this chapter exists partly so that the question can be asked again in that regime rather than assumed away.

One measurement limitation, recorded honestly. The `.twotemp` sidecar prints ten significant digits, so below about 2080 km the printed $T_e - T_i$ sits at the print floor — roughly 10^{-6} K at 6.6 kK. The quoted interior decoupling in that region is therefore an *upper bound* of order 10^{-6} K, not a measurement. This does not affect the headline number of equation (18.17), nor the quasi-steady maxima, which are accumulated inside the solver in full double precision over every cell and every step rather than read back from the sidecar.

Part VII

Current Status and Roadmap

Chapter 19

Trust Boundaries

19.1 Scientifically trustworthy components

The following statements are supported by code-level consistency checks and controlled tests:

- the release state lies on the intended pure-hydrogen Saha equilibrium manifold;
- pressure, internal energy, ionization energy, effective heat capacity, and the equilibrium acoustic speed are mutually consistent (the release flux and CFL use the frozen speed of section 11.3.3, which is exact for the split of Eq. (11.20) and needs no table);
- the CRASH-derived Γ_1 table is parsed strictly and agrees with independent Saha derivatives within documented error;
- the Model C7/PCHIP temperature profile is reintegrated in the active hydrostatic EOS;
- the reference-free hydrodynamic operator preserves that state with a residual that is a second-order-convergent truncation error, roughly 10^8 above the round-off floor of the `double` stored state and therefore free of any representation component (section 11.4);
- the conserved state, mesh metrics, physical constants, Courant number and simulation time are all stored in `double` under one coherent precision knob, and the test suite and release validation pass in both that build and the single-precision diagnostic build with no loosened tolerance (section 10.2);
- nonlinear conduction satisfies its backward-Euler residual and conserves energy in closed tests;
- nonuniform reconstruction, conduction, CFL, and HSE operators use actual mesh metrics;
- the effective face mass flux is a valid conserved transport diagnostic;
- serial and OpenMP execution are identical for the validated path;
- the selected long-run controls terminate and report correctly.

19.2 Numerically validated but configuration-specific results

The 12-thread default and CFL 0.50 validation apply to the coarse-equivalent $N = 500$, R4, 661-cell, physical-conduction-only release shape on the tested Apple Silicon workstation. A different mesh, scenario, source set, or hardware should repeat the acceptance sweep.

Active release path

The release’s *single-temperature* assumption now belongs in this category rather than in the untested one. The experimental two-temperature study of chapter 18 measured it directly on this exact configuration and found the evaporation observables reproduced to 0.2%, a quasi-steady $\max |T_e - T_i|/T_i$ of order 10^{-5} , and a worst-case decoupling of 3.3×10^{-3} confined to a startup transient of tens of milliseconds. The assumption is therefore validated *for this column*,

and for the same reason as the CFL: it was measured on one model shape. It should not be carried into a flare, beam-heated, hotter, or more rarefied regime, where the collisional coupling weakens and the electron heat flux grows (section 18.10.7).

19.3 Provisional scientific results

The release calculation develops a coherent upward effective mass flux that is nearly uniform with height across the upper-chromosphere diagnostic window. The $N = 500$ run (Roe flux, single-precision state) reaches 4000 s with low face-flux roughness, but this is not a formally converged evaporation rate because:

1. the matched $N = 1000$ calculation was stopped at 2080 s rather than completed to 4000 s;
2. long-duration $N = 500$ mass flux is therefore not established as grid-converged to better than 10%;
3. radiation is absent in the reference convergence calculation;
4. the upper atmosphere and corona are not resolved self-consistently;
5. that run predates the precision cutover, so its lower chromosphere (about 1600–1850 km) could not represent its own per-step continuity increment (section 10.2). The `double` repeat closes the column mass budget that the single-precision run violated, and in doing so moves the top-of-domain evaporation observables down by $\sim 28\%$, so the single-precision figures quoted above are superseded rather than merely provenance-labelled and must be re-extracted from the `double` run.

No absolute evaporation rate from the current calculation should yet be presented as a converged physical prediction.

Chapter 20

Remaining Limitations and Priorities

20.1 Outer conductive boundary

The release discretization question has been resolved: $T_{\text{face}} = 22000$ K is imposed at the physical outer face and the top-cell distance is $\Delta s_{\text{top}}/2$. The remaining limitation is physical rather than geometric. A fixed external temperature and back pressure do not respond to the evolving column and cannot replace a self-consistent coronal energy balance. Future coupling should replace these prescribed reservoir data with exchanged characteristic and conductive information.

20.2 Radiative balance

The reference evaporation calculation includes conductive heating but no radiative cooling. This isolates the numerical problem but omits a leading chromospheric/TR sink. The implemented radiation and TRAC modules need a separate convergence campaign before quantitative energy partition or quasi-steady evaporation is claimed.

20.3 Equilibrium versus non-equilibrium ionization

Saha equilibrium is a controlled release closure, not a universal chromospheric approximation. Dynamic hydrogen ionization can lag temperature changes substantially [Carlsson and Stein, 2002, Leenaarts et al., 2007]. Future work must compare reaction time scales with hydrodynamic and conductive time scales and validate a non-equilibrium model without reintroducing inconsistent pressure or energy bookkeeping.

20.4 Per-step Saha re-equilibration and the frozen Riemann solve

Active release path This is, in the author's assessment, the largest unresolved physical-consistency question in the released model, and it was brought into focus rather than created by the numerical-flux cutover.

The release now treats the composition as *frozen* inside the local Riemann problem, for the relaxation-time reasons of section 11.3.3. But the equation of state is unchanged: every decode, every reconstruction and every ghost closure evaluates the instantaneous Saha fraction $x_{\text{Saha}}(n_{\text{H}}, T)$ of

Eq. (3.7). The composition is therefore reset to equilibrium at the end of every timestep. Schematically, one step is

$$\underbrace{\text{Saha-equilibrated state}}_{U^n} \longrightarrow \underbrace{\text{frozen-composition hydrodynamic step}}_{\text{exact Godunov flux}} \longrightarrow \underbrace{\text{Saha-equilibrated state}}_{U^{n+1}}. \quad (20.1)$$

These are two separate modelling choices, and the first does not imply the second. It is coherent to freeze the composition over the Riemann fan while relaxing it over the step, and it is coherent to do neither; what is not coherent is to justify the first by a long relaxation time and the second by an implicitly short one.

Read charitably, Eq. (20.1) is an *operator-split instantaneous-relaxation* approximation: hydrodynamics at frozen composition, followed by a relaxation substep taken in the limit $\tau_{\text{ion}} \rightarrow 0$. Operator splitting of a stiff relaxation term against transport is standard and, in the genuinely stiff limit, correct: if the ionization state truly equilibrated on a timescale short compared with the step, projecting onto the Saha manifold each step would be the right thing to do, and the frozen Riemann solve would then be a numerical device that happens not to matter because the relaxation immediately erases its composition error. The difficulty is that this is the opposite of the regime invoked to justify the frozen solve.

Three timescales must be kept distinct, and conflating any two of them produces a wrong argument.

Numerical timestep $\Delta t \approx 5.5 \times 10^{-3}$ s in the release configuration. **This is not the relevant comparison.** Saha equilibrium is not invalidated by taking small steps; a numerical timestep is not a physical timescale, and an argument of the form “ $\tau_{\text{ion}} \gg \Delta t$, therefore Saha fails” is simply an error. The step size controls only the accuracy of whatever splitting is used.

Physical dynamical/compressive time the time over which a fluid element’s thermodynamic state actually changes: of order the acoustic crossing time of the structure of interest, a shock-front crossing at the short end and ~ 60 s for the chromospheric dynamical timescale quoted by Leenaarts [2020].

Ionization/recombination relaxation time τ_{ion} , computed for chromospheric hydrogen by Carlsson and Stein [2002]: $\sim 10^3$ – 10^5 s through the mid-chromosphere, $\sim 10^2$ s at the transition-region base, and 10 – 10^3 s even inside shocks.

The physically meaningful comparison is the second against the third. On that comparison τ_{ion} exceeds the dynamical time by one to three orders of magnitude through most of the modelled column, which is precisely the statement that hydrogen ionization is *not* in instantaneous equilibrium. Leenaarts and Wedemeyer-Böhm [2006] make the magnitude of the resulting error concrete: in their simulations the equilibrium ionization degree varies by more than twenty orders of magnitude between shocked and unshocked gas, while the actual non-equilibrium degree stays nearly constant. Leenaarts et al. [2007] report the thermodynamic consequence — LTE ionization acts as an energy buffer that suppresses shock temperatures by several thousand kelvin — and Leenaarts [2020] describes the artificial “preferred temperature” plateaus (near 6, 10 and 22 kK) that an LTE equation of state imposes and that vanish under non-equilibrium ionization.

The honest statement of the release model is therefore: the equilibrium decode is a known approximation whose error is *unquantified in this code*. It is not defended here as physically accurate for rapidly varying flow; it is a closure that makes the three-variable conserved state determinate and cheap, it is applied consistently everywhere in the solver, and it is what every quantitative result in this document was computed with. The quasi-static evaporation problem the release actually reports is the most favourable case for it — the flow is smooth, subsonic and slowly varying, so the composition error relative to a finite-rate treatment is smallest — and the least favourable case, a

strong compressive event inside the partial-ionization zone, is exactly the regime the release does not claim.

Planned A physically more complete model would evolve the ionization state rather than enforce it: carry x (or the ion and neutral populations) as an additional evolved variable with finite-rate collisional and radiative ionization/recombination source terms, so that the frozen Riemann solve and the composition update are consistent by construction and Eq. (20.1) is replaced by a genuine relaxation with a physical τ_{ion} . The non-release two-fluid solver already implements such a network (section 5.1), which makes this a matter of reconciling two existing formulations rather than of new physics. Two consequences should be anticipated before it is attempted: the conserved state widens and the caloric inversion is no longer a function of (ρ, e_{int}) alone, and the release’s careful separation between stored and derived quantities — currently one of its strongest correctness properties — would have to be redesigned rather than extended. Nothing in the present cutover attempts it, and it should not be attempted as a side effect of a numerical-method change.

20.5 Unresolved corona and coupling

The truncated domain treats the corona as an external pressure and thermal reservoir. It does not evolve a coronal energy balance, loop geometry, wave heating, or global mass inventory. **Planned** AWSoM coupling should replace prescribed reservoir data with exchanged characteristic, conductive, mass, momentum, and energy information while retaining a well-posed number of incoming boundary conditions.

20.6 Numerical precision

Active release path **Resolved.** The packed state, the mesh metrics, the physical constants, the Courant number and the simulation time are all **double**, under the single scalar type **chromosphere::Real** (section 10.2). The storage-imposed noise floor that once limited m s^{-1} -scale velocity studies is gone on all three measurements that were used to characterize it: the hydrostatic residual now sits $\sim 10^8$ above the round-off floor, it converges at unbroken second order over a $41\times$ mesh-refinement ladder (section 11.4), and the residual velocity field now converges in time as well, at observed order ≈ 1.9 in L^1 where it previously diverged. This section previously recommended evaluating a double-precision conserved state or a perturbation-variable formulation; the first was done, and the second is unnecessary as a result.

Active release path Two limitations that were formerly confounded with precision remain, and are *not* precision problems. The localized first-interior-face mass-flux artifact is precision independent and open (section 10.2.4). And the $N = 500$ long-duration evaporation mass flux is still not established as grid-converged to better than 10 %, which is a mesh-resolution limitation entirely.

Active release path Two cautions survive on the quantitative side. The CFL acceptance matrix has *not* been re-run in **double**, so CFL 0.50 stands on its original single-precision evidence, and no positive noise floor has been established for m s^{-1} -scale velocities — the storage-imposed floor is retired, the caution is not. And the 4000 s $N = 500$ reference calculation *has* now been repeated in **double**, with a material result: the gross continuity violation of the single-precision run disappears (base-to-top f_{total} ratio $2.811 \rightarrow 1.003$, the column mass budget closing at predicted/observed 0.994 against -214.9 , and no frozen cell), but the top-of-domain evaporation observables move by $\sim 28\%$ — top-cell velocity $13.18 \rightarrow 9.50 \text{ m s}^{-1}$ and top-face $f_{\text{total}} 3.753 \times 10^{-10} \rightarrow 2.702 \times 10^{-10} \text{ kg m}^{-2} \text{ s}^{-1}$.

The physical reading is coherent: with the lower chromosphere unable to drain, single precision sustained a mass supply the column does not have and so over-predicted the evaporation flux. **Every release evaporation number quoted from a single-precision run is therefore high by of order 30 % and must be re-measured before it is presented as a current release result.** The upper-chromosphere window diagnostics have now been extracted for both legs on the identical window and flux: the mean effective mass flux is $2.4880 \times 10^{-10} \text{ kg m}^{-2} \text{ s}^{-1}$ in **double** against 3.7390×10^{-10} in single precision, i.e. 33 % lower, with the mean window velocity falling $5.848 \rightarrow 4.741 \text{ ms}^{-1}$ (range 3.180–11.866 to 1.884–8.767), $Q_{\text{eff}} 1.48 \times 10^{-3} \rightarrow 5.25 \times 10^{-4}$, the maximum reconstructed face-pressure mismatch $0.0006 \% \rightarrow 0.0001 \%$, the relative column mass drift $-2.361 \times 10^{-4} \rightarrow -6.053 \times 10^{-5}$, and the thermal boundary layer thickening $12.42 \rightarrow 15.73 \text{ km}$. $2.49 \times 10^{-10} \text{ kg m}^{-2} \text{ s}^{-1}$ is the current release value; it remains subject to the unchanged resolution limitation below. A separate $\sim 13 \%$ monotone f_{total} decline within 1600–1850 km survives in **double** and is still relaxing at 4000 s; it is not the first-face artifact, and it is independent evidence for the standing long-duration caveats rather than a precision problem.

20.7 Recommended order of future work

1. re-extract the upper-chromosphere evaporation diagnostics from the completed 4000 s $N = 500$ **double** calculation, which supersede the single-precision figures by $\sim 30 \%$, then complete a matched long-duration $N = 1000$ release-numerics, physical-conduction-only comparison; re-run the CFL acceptance matrix in **double** alongside it;
2. add radiative balance and TRAC under the same local-refinement framework;
3. quantify equilibrium-ionization error against a controlled finite-rate calculation;
4. define a characteristic coupling contract for AWSoM/SWMF;
5. only then broaden geometry, heating, and multidimensional scope.

Part VIII

Appendices

Appendix A

General Two-Fluid Equations and Source Stages

A.1 Conservative variables

Historical A representative full state is

$$U = (\rho_i, \rho_n, \rho_i V, \rho_n U, E_i, E_n, E_e)^T. \quad (\text{A.1})$$

For fixed γ , the historical row energies are

$$E_i = \frac{p_i}{\gamma - 1} + \frac{1}{2} \rho_i V^2 + \rho_i \Phi, \quad (\text{A.2})$$

$$E_n = \frac{p_n}{\gamma - 1} + \frac{1}{2} \rho_n U^2 + \rho_n \Phi, \quad (\text{A.3})$$

$$E_e = \frac{p_e}{\gamma - 1}. \quad (\text{A.4})$$

The active release path replaces these caloric relations with the mapping in equation (6.3).

A.2 Backward-Euler drag stage

The legacy drag update uses the center-of-mass velocity v_{cm} and drift $w = V - U$. For a relaxation rate λ ,

$$w^{n+1} = \frac{w^*}{1 + \lambda \Delta t}. \quad (\text{A.5})$$

Total momentum is unchanged. The decrease in relative kinetic energy is deposited as heat. This stage is stable in the strong-coupling limit; the release solver has no such stage because it carries only one momentum.

A.3 Backward-Euler reaction stage

Writing the ion fraction as f , the local kinetic network produces polynomial backward-Euler equations because electron-impact and three-body channels depend on powers of the electron density. The implementation selects the physically admissible root in $[f_{\text{min}}, 1 - f_{\text{min}}]$, updates carrier masses conservatively, and accounts for ionization energy according to the reaction channel. This source stage belongs to the two-fluid solver and does not exist in the release.

A.4 Historical multidimensional formulation

Historical The repository documentation also preserves the parent two-fluid MHD system with induction, Lorentz forces, and divergence control. It motivated the conserved layout and field-aligned reduction, but no current release result is obtained by evolving that multidimensional system. It should be treated as architectural background rather than implemented release functionality.

Appendix B

EOS Table Construction and Format

B.1 Generation

The production Γ_1 table is generated by `util/eos/tabulate_gamma.f90`. The CRASH calculation is restricted to classical pure hydrogen:

- excitation disabled;
- Fermi-gas correction disabled;
- Coulomb correction disabled;
- statistical-weight convention aligned with the coefficient-one Saha relation;
- constants aligned with the C++ EOS.

The full validation output contains additional CRASH columns, while the tracked runtime table contains only $\log_{10} T$, $\log_{10} n_{\text{H}}$, and Γ_1 , plus machine-readable metadata.

B.2 Runtime policy

The loader requires an exact table identifier, axis names, ordering, dimensions, and physics metadata. Production interpolation is bilinear in log coordinates. Out-of-domain states throw; clamping exists only through an explicit debug option. A checksum test guards the tracked table, and the runtime diagnostic header records the actual loaded path and SHA-256 value.

Appendix C

Implementation Status Matrix

Element	Status	Current interpretation
Saha equilibrium	Active release path	Physical composition closure
CRASH Γ_1 table	Active release path	Acoustic derivative only
Analytic p, e_{int}, C_V	Active release path	Caloric EOS and conduction derivative
Seven-row carrier storage	Retired	Removed from the release; now exclusive to the non-release two-fluid solver
Independent ion-neutral drift	Implemented but inactive	Legacy two-fluid path
Finite-rate hydrogen chemistry	Implemented but inactive	Legacy path and controlled experiments
Common-temperature conduction	Active release path	Nonlinear backward Euler
Separate electron conduction equation	Experimental	Two non-release routes: the two-fluid three-temperature path, and the four-row two-temperature test-study solver of chapter 18 (κ_e on T_e , κ_n on T_i , coupled implicitly with the collisional exchange) <code>src/two_temp/</code> , ($\rho, \rho v, E, E_e$); selected only by <code>model_column_2t</code> ; test study of the release's single-temperature assumption, not a release capability. Completed: closure degenerates exactly at $T_e = T_i$, evaporation observables within 0.2% of the matched release run, quasi-steady decoupling $\sim 10^{-5}$ (chapter 18)
Four-row two-temperature state	Experimental	
Optically thick/thin radiation	Implemented but inactive	Two-fluid solver only; not a release stage

TRAC	Implemented but inactive	Two-fluid solver only; not a release stage
Beam heating	Experimental	Two-fluid solver only; flare studies
Volumetric coronal heating	Experimental	Two-fluid solver only; resolved-corona studies
SWMF-style exact-Riemann Godunov flux	Active release path	Release numerical flux: exact ideal-gas Riemann solve at frozen $\gamma = 5/3$ with the ionization energy in an upwinded passive offset, face-local Rusanov fallback, per-run fallback tally
Mixture Roe characteristic flux	Implemented but inactive	Reference/comparison solver (ISO_RIEMANN=roe-local); the previous release flux; equilibrium Γ_1 linearization
Flux-consistent CFL signal speed	Active release path	$\gamma_{\text{sig}} = \max(\Gamma_1, 5/3)$ under the release flux, Γ_1 under the reference fluxes
double conserved state	Active release path	Release storage precision, one coherent knob (chromosphere::Real) covering state, mesh metrics, constants, CFL and sim_time; required because the per-step continuity increment is sub-ULP in single precision, and consistent with SWMF production (section 10.2)
Single-precision state build	Implemented but inactive	Diagnostic only (-DCHROMO_STATE_FLOAT32=ON, separate tree); reproduces the pre-cutover round-off sensitivity; never a production or quotable configuration
CHROMO_STATE_FLOAT64 option	Retired	The incoherent container-only promotion; CMake hard-errors on it
Mixed-precision state/scratch split	Retired	Considered and rejected: negligible memory saving at this size, less consistent with a uniformly double framework, and the exact failure mode of the container-only promotion
Static outer refinement	Active release path	Release mesh option
AMR/regridding	Planned	Not implemented
Face-local artificial diffusion	Historical	Removed from the release conduction operator; retained inside the two-fluid Stage-D operator for historical diagnostics
Fixed-pressure outer reservoir	Active release path	One incoming hydrodynamic condition
Hydro/conduction temperature decoupling	Active release path	Release upper boundary

Characteristic outer bound- ary	Planned	Needed for coupling/generalization
AWSOM coupling	Planned	Not yet implemented

Appendix D

Release-Relevant Environment Variables

This chapter lists only variables that a canonical `model_column` release run actually accepts. Variables belonging to the historical two-fluid research path — which `model_column` rejects outright rather than silently ignoring — are listed separately in Appendix E. The distinction is exact: a variable appears here if and only if the release solver reads it.

D.1 EOS, domain, and physics

Variable	Meaning
<code>GAMMA_TABLE</code>	Path to the strict Γ_1 table; selects the release solver. The <code>model_column</code> scenario supplies the production path unconditionally, so this only redirects it.
<code>ISO_H_BASE</code>	Absolute lower height in km.
<code>ISO_DH</code>	Domain length in km.
<code>ISO_NS</code>	Coarse-equivalent cell count; refinement may add cells.
<code>ISO_HEAT_FLUX</code>	Enables the conduction-driven stage.
<code>ISO_T_TOP</code>	Absolute external conductive temperature at the physical outer face in K.

D.2 Boundary and mesh

Variable	Meaning
<code>ISO_HYDRO_T_DECOUPLE</code>	Use live top-cell temperature in hydro ghosts while retaining the fixed physical-face conduction datum.
<code>ISO_INNER_T_NEUMANN</code>	Replace lower fixed-temperature conduction condition by zero flux.
<code>ISO_VCAP</code>	Outer velocity cap in local sound-speed units.
<code>ISO_TJUMP_A</code> , <code>ISO_TJUMP_B</code>	Conductive-wall and historical second-ghost temperature multipliers.

ISO_REFINE_PROFILE	<code>lower</code> or <code>outer</code> ; overrides generic mesh alias.
ISO_REFINE_FACTOR	Fine-to-coarse resolution ratio; 1 disables refinement.
ISO_REFINE_S_LO_KM	Start of fully refined region measured from the inner face.
ISO_REFINE_S_HI_KM	End of lower-profile refined region; unused by outer profile.
ISO_REFINE_TRANSITION_KM	Minimum geometric-grade width.
GRID_REFINE_*	Generic aliases used when corresponding ISO_REFINE_* values are absent.
ISO_RELAX_TIME	Optional initial no-conduction relaxation interval.

D.3 Reference-only overrides and diagnostics

The release solver accepts the variables below, but none of them is a production setting. The first block selects a reference numerical method for controlled comparison against the release default; the second block only enables read-only sidecar output and does not alter the solution.

Variable	Meaning
ISO_RIEMANN	Reference flux override: <code>roe-local</code> selects the mixture Roe characteristic flux of section 11.3.1 — the equilibrium- Γ_1 linearization that was the previous release flux — and <code>rusanov</code> the local Lax–Friedrichs flux. The release default is the SWMF-style exact-Riemann Godunov flux of section 11.3.2; no variable is needed to select it. Either override also reverts the CFL signal speed to Γ_1 (section 12.2).
ISO_RECONSTRUCTION	Reference reconstruction override: <code>lnrho-v-lnt</code> restores the temperature-based form. The release default is $(\log \rho, v, \log p)$; no variable is needed to select it.
ISO_LIMITER	Diagnostic override: <code>mc3</code> , <code>minmod</code> , or <code>first</code> ; not for production.
ISO_MC3_BETA	MC3 beta under the diagnostic limiter override.
CHROMO_FACE_FLUX_DIAG	Enable read-only face-flux sidecar.
CHROMO_OUTER_COND_DIAG	Enable accepted outer-conduction sidecar.
CHROMO_GAMMA_DIAG	Enable decoded EOS-aware sidecar.

D.4 Run and OpenMP controls

Variable	Meaning
CHROMO_CFL	Runtime CFL override.
CHROMO_T_END	Requested physical end time.
CHROMO_STEP_CAP	Explicit positive step cap.
CHROMO_FRAME_DT	Snapshot cadence in physical seconds.

CHROMO_OUTPUT	Enable/disable main snapshots.
CHROMO_PROFILE	Enable timing/profile summary.
OMP_NUM_THREADS	OpenMP team size; launchers default to 12.
OMP_DYNAMIC	Launchers set false.
OMP_MAX_ACTIVE_LEVELS	Launchers set 1.

Appendix E

Non-Release Two-Fluid Research Environment Variables

Every variable in this chapter configures the historical seven-row two-fluid research solver, reached through the separate `model_gentle` scenario (and the other two-fluid scenarios). None of them has meaning for the single-fluid release, whose timestep is the explicit hydrodynamic update followed by the implicit physical conduction solve and nothing else. A canonical `model_column` run therefore *rejects* the master switches below before the initial state is built rather than silently ignoring them; their dependent sub-parameters are inert without those switches. This documentation is retained because the two-fluid path remains implemented, tested, and used for the historical studies recorded in this document.

Variable	Meaning
ISO_GAMMA	Historical fixed- γ index of the two-fluid solver; $\gamma = 1.05$ was the nearly isothermal surrogate. Mutually exclusive with <code>GAMMA_TABLE</code> .
ISO_TWO_FLUID, ISO_IONIZATION	Independent ion/neutral fluids and finite-rate ionization.
ISO_COOLING	Optically thick/thin radiative sink.
ISO_TRAC, ISO_TRAC_TCHROM, ISO_TRAC_TCMAXFRAC	TRAC enable/disable and cutoff controls.
ISO_CORONA, ISO_CORONA_TOP_KM	Extend the domain into a resolved corona.
ISO_QFLUX, ISO_QFLUX0, ISO_QFLUX_ENHANCE, ISO_QFLUX_T_ON, ISO_QFLUX_RAMP	Imposed outer-flux boundary and ramp.
ISO_CHEAT, ISO_CHEAT_*	Volumetric coronal-heating configuration.
ISO_TBOOST, ISO_TBOOST_TC, ISO_TBOOST_W	One-time initial-condition coronal superheat and its selector parameters.
ISO_NUMERICAL_DIFFUSIVITY_MULT	Multiplier on the historical C_{num} of the two-fluid Stage-D conduction operator. The release conduction operator has no such term.
SINGLE_FLUID, ENABLE_TE	Legacy runtime solver settings of the two-fluid path; also rejected by a release run.

Experimental One further non-release variable belongs to neither list. `TT_OUTER_TI` configures the outer-face heavy-particle conduction condition of the experimental two-temperature solver and is read only by the `model_column_2t` scenario; it takes `neumann` (default) or `dirichlet`, and any other value is rejected. It is documented in section 18.6.3. That scenario additionally rejects `ISO_GAMMA`, `ISO_RIEMANN`, and `ISO_RECONSTRUCTION`.

Appendix F

Canonical Run Configurations and Diagnostics

F.1 Validated low-I/O reference run

```
CHROMO_T_END=1000 \  
scripts/run_chromo_realtime.sh \  
  outputs/model_column/model_column_release_1000s.txt \  
  full no-ionization model_column - 20.0 no-cooling
```

The scenario supplies the production Gamma table, 1600–2153 km domain, $N = 500/R4$ mesh with 661 actual cells, physical-face 22000 K boundary, physical conduction only, and release/ $\log p$ numerics. The launcher supplies 12 threads, CFL 0.50, and low-I/O runtime defaults. Explicit environment assignments override these settings and define a new validation target.

F.2 Diagnostic output semantics

Main snapshot. The three packed release rows (ρ, m, E) plus geometry metadata; authoritative for restart/raw conserved state. A two-fluid snapshot carries seven rows instead and an experimental two-temperature snapshot four, and the header line records the width.

Gamma sidecar. Decoded physical mixture quantities. `kappa_solver` is identically `kappa_physical` because release conduction is physical only; the column is kept for format stability.

Face-flux sidecar. Central, diffusive, and total mass fluxes at selected faces, from which the diagnostic $F_{\text{eff}} = F_{\text{total}}^{\rho} - F_{\text{ref}}^{\rho}$ is formed against the $t = 0$ baseline. The mass flux is the `mix::RHO` continuity flux; the release state has no carrier rows to sum. Read-only.

Outer-conduction sidecar. Accepted final-Newton physical-face temperature, half-cell boundary distance, physical face conductivity, and the resulting face flux. There is no artificial-conduction column, because the operator has no such term. Read-only.

Two-temperature sidecar. Experimental Written only by `model_column_2t` (section 18.7). Fourteen decoded columns per cell, including both temperatures, both partial pressures, both conductivities, the exchange rate Q_{ei} , and the local τ_{eq} . Same header/frame-marker format as the gamma sidecar. Read-only.

Console/profile log. Termination reason, timestep-limit counts, Newton/inversion counters, OpenMP team information, and region timing. A two-temperature run additionally prints

the `twotemp.*` tally, whose clamp counters must all be zero for the run to be quotable.

Appendix G

Test Inventory

The release-relevant test groups are:

- strict table metadata, axes, interpolation, limits, continuity, and checksum;
- Saha, caloric energy, effective heat capacity, temperature inversion, and acoustic derivatives;
- release state round trips, derived-carrier identities, and face-state identities;
- release activation, solver isolation, and rejection of incompatible scenarios, integrators, and every historical two-fluid knob;
- a bit-for-bit check that enabling the retired two-fluid stage flags on a release Grid changes nothing, and a `model_gentle` smoke test that the retained seven-row path still builds and advances;
- PCHIP nodal and derivative behavior;
- Saha-HSE integration and all four ghost states;
- hydro/conduction temperature decoupling;
- nonlinear conduction residual, conservation, and strong-profile convergence;
- uniform and irregular-grid reconstruction, CFL, HSE, and conduction;
- lower and outer static refinement, grading edge cases, and the two-fluid local numerical diffusivity;
- reference-free hydrostatic preservation on the release mesh, and evaporation regression windows;
- face-flux and outer-conduction diagnostic neutrality;
- the release SWMF exact Riemann solver against the five standard reference star states, its $x/t = 0$ sampling, the degenerate stationary-contact and uniform-flow cases, and a shock-tube check that the energy offset of equation (11.18) is inert to the hydrodynamics; and, on the production right-hand side, that the release flux is conservative, differs measurably from both reference solvers, and counts rather than hides a fallback;
- that the CFL rule of Eqs. (12.3)–(12.4) uses the frozen signal speed on the release path and the equilibrium speed under the Roe reference flux, that the release step never exceeds the reference step, that no cell exceeds the requested CFL measured against the frozen waves, and that the frozen/equilibrium speed gap is genuinely large somewhere in the column, so the rule is not a disguised no-op;
- output scheduling, step-cap validation, and termination classification;
- serial/OpenMP process-level and output equivalence checks;
- storage-precision coherence: the equation-of-state table-corner domain regression walks its endpoints in the storage type `Real` and asserts against the solver’s own acceptance predicate rather than a test-side recomputation, so it is precision independent and runs unskipped in both the `double` release build and the single-precision diagnostic build (section 10.2).

The exact assertion count is intentionally omitted because it changes as coverage grows. Release

claims should be tied to named risks and passing test groups rather than a stale aggregate number.

Appendix H

Chronological Development Ledger

Phase	Main question	Lasting outcome
Early project	Can a plasma-neutral chromosphere be evolved in a field-aligned code?	Seven-row two-fluid architecture, source stages, C7 scenarios, conduction and radiation experiments; retained as a separate non-release solver.
Boundary/HSE work	Why do nominally static columns drift or reverse at the boundaries?	Discrete HSE reservoirs, log reconstruction, fixed upper back pressure, two-ghost consistency; equilibrium-reference well balancing, later superseded by completing the MUSCL–Hancock source treatment and the trapezoidal lower ghost ladder.
Constant- γ studies	Can conduction-driven response be isolated before a complete EOS?	Fixed 5/3 and 1.05 baselines; useful diagnostics but inadequate ionization thermodynamics.
Gamma stages 0–2	Is a CRASH-derived acoustic table reproducible and Saha consistent?	Strict pure-H table generation, metadata, checksums, whole-grid validation.
Gamma stages 3–4	Can conserved energy be inverted consistently?	Analytic caloric EOS, safeguarded inverse, conservative common-fluid mapping.
Gamma stages 5–6	Can hydro faces remain on the equilibrium manifold?	Initial $\log \rho, v, \log T$ reconstruction, common Γ_1 wave speed, and fail-closed activation; later superseded by $\log p$ reconstruction and Roe dissipation, and that in turn by the frozen-composition Godunov flux.
Gamma stages 7–9	Can sources, conduction, IC, and boundaries use the same EOS?	Nonlinear mixture conduction, total-energy source updates, PCHIP/Saha-HSE initialization, release sign-off tests.
Upper-TR diagnosis	Is the top ripple a boundary instability or flux error?	Hydro/conduction T decoupling, smooth effective face flux, first-order truncation interpretation.
Refinement	Can the ripple be resolved efficiently?	Static outer profile with bounded grading and actual-cell accounting.

Artificial conduction	Why did refined evaporation change incorrectly?	Face-local mesh scaling isolated the error; production later removed artificial conduction entirely.
Release numerics	How can low-Mach face pressure and dissipation be controlled?	$(\log \rho, v, \log p)$ reconstruction, local mixture Roe flux, and face-local Rusanov fallback.
Release flux cutover	Which acoustic response should the face Riemann problem use?	Frozen composition ($\gamma = 5/3$) via the SWMF exact-Riemann Godunov flux, with the CFL signal speed made consistent with it.
Precision cutover	Why would the lower chromosphere not relax?	The per-step continuity increment was sub-ULP in single precision in 619 of 660 cells. double storage under one coherent knob, matching SWMF production precision; the hydrostatic residual became second-order convergent and measurable, and the localized first-face artifact was isolated as precision independent.
Performance	Can the release run at useful wall time without changing the method?	EOS/cache optimization, deterministic OpenMP, 12-thread default, validated CFL 0.50 and long-run controls.
Current	Is the evaporation rate converged?	Stable $N = 500$ evolution; matched long-duration $N = 1000$ completion remains necessary for a sub-10% claim.

Appendix I

Retired Scenarios and Archived Formulations

I.1 Retired active assumptions

The following should not reappear as current release descriptions:

- analytic-isentrope initialization for `model_column`; the current IC is Model C7/PCHIP with Saha-HSE reintegration;
- independent finite-rate ionization in the release solver;
- a common fixed- γ pressure law in the release EOS;
- centered upper pressure extrapolation across the top cell;
- a live top-pressure anchor with no external back pressure;
- a single upper ghost temperature serving both hydro and conduction in the validated configuration;
- globally coarse-scaled artificial diffusivity on a refined mesh;
- interpretation of cell-centered ρv ripple as a discontinuous conserved mass flux;
- single-precision storage of the conserved state, or of any other stored physical scalar, in a production run;
- attribution of the reference-free hydrostatic residual to storage round-off, or the claim that no truncation-error trend survives between $N = 500$ and $N = 1000$;
- presentation of a last-pair near-flat mass flux as formal convergence.

I.2 Archived but useful formulations

The full two-fluid equations, kinetic rate network, three-temperature exchange, TRAC, chromospheric and coronal radiation, beam heating, coronal heating, loop geometries, and multidimensional MHD derivations remain valuable. They support future work and explain past decisions. Their equations and references are retained in this record, but their status labels must accompany any reuse.

Bibliography

- Qusai Al Shidi, Ofer Cohen, Paul Song, and Jiannan Tu. Time-dependent two-fluid magnetohydrodynamic model and simulation of the chromosphere. *Solar Physics*, 294(9):124, 2019. doi: 10.1007/s11207-019-1513-8.
- E. H. Avrett and R. Loeser. Models of the solar chromosphere and transition region from SUMER and HRTS observations: Formation of the extreme-ultraviolet spectrum of hydrogen, carbon, and oxygen. *The Astrophysical Journal Supplement Series*, 175:229–276, 2008. doi: 10.1086/523671.
- Jonas P. Berberich, Praveen Chandrashekar, Christian Klingenberg, and Friedrich K. Röpke. Second order finite volume scheme for Euler equations with gravity which is well-balanced for general equations of state and grid systems. *Communications in Computational Physics*, 26(2):599–630, 2019. doi: 10.4208/cicp.OA-2018-0152.
- Christophe Berthon. Why the MUSCL–Hancock scheme is L^1 -stable. *Numerische Mathematik*, 104(1):27–46, 2006. doi: 10.1007/s00211-006-0007-4.
- M. Carlsson and J. Leenaarts. Approximations for radiative cooling and heating in the solar chromosphere. *Astronomy & Astrophysics*, 539:A39, 2012. doi: 10.1051/0004-6361/201118366.
- M. Carlsson and R. F. Stein. Dynamic hydrogen ionization. *The Astrophysical Journal*, 572:626–635, 2002. doi: 10.1086/340293.
- Praveen Chandrashekar and Christian Klingenberg. A second order well-balanced finite volume scheme for Euler equations with gravity. *SIAM Journal on Scientific Computing*, 37(3):B382–B402, 2015. doi: 10.1137/140984373.
- D. G. Hummer. Total recombination and energy loss coefficients for hydrogenic ions at low density for $10 \leq T_e/Z^2 \leq 10^7$ K. *Monthly Notices of the Royal Astronomical Society*, 268:109–112, 1994. doi: 10.1093/mnras/268.1.109.
- C. D. Johnston, S. J. Bradshaw, P. J. Cargill, and A. W. Hood. Modeling the solar transition region using an adaptive conduction method. *The Astrophysical Journal Letters*, 873:L22, 2019. doi: 10.3847/2041-8213/ab0c1f.
- C. D. Johnston, P. J. Cargill, A. W. Hood, I. De Moortel, S. J. Bradshaw, and A. C. Vaseekar. Modelling the solar transition region using an adaptive conduction method. *Astronomy & Astrophysics*, 635:A168, 2020. doi: 10.1051/0004-6361/201936979.
- R. Käppeli and S. Mishra. Well-balanced schemes for the Euler equations with gravitation. *Journal of Computational Physics*, 259:199–219, 2014. doi: 10.1016/j.jcp.2013.11.028.

- E. Khomenko, M. Collados, A. Díaz, and N. Vitas. Fluid description of multi-component solar partially ionized plasma. *Physics of Plasmas*, 21(9):092901, 2014. doi: 10.1063/1.4894106.
- J. Leenaarts. Radiation hydrodynamics in simulations of the solar atmosphere. *Living Reviews in Solar Physics*, 17:3, 2020. doi: 10.1007/s41116-020-0024-x.
- J. Leenaarts and S. Wedemeyer-Böhm. Time-dependent hydrogen ionisation in 3d simulations of the solar chromosphere. methods and first results. *Astronomy & Astrophysics*, 460:301–307, 2006. doi: 10.1051/0004-6361:20066123.
- J. Leenaarts, M. Carlsson, V. Hansteen, and R. J. Rutten. Non-equilibrium hydrogen ionization in 2D simulations of the solar atmosphere. *Astronomy & Astrophysics*, 473:625–632, 2007. doi: 10.1051/0004-6361:20078161.
- E. T. Meier and U. Shumlak. A general nonlinear fluid model for reacting plasma–neutral mixtures. *Physics of Plasmas*, 19(7):072508, 2012. doi: 10.1063/1.4736975.
- Igor V. Sokolov, Bart van der Holst, Ward B. Manchester, Doga Can Su Ozturk, Judit Szenté, Aleksandre Taktakishvili, Gabor Toth, Meng Jin, and Tamas I. Gombosi. Threaded-field-line model for the low solar corona powered by the Alfvén wave turbulence. *The Astrophysical Journal*, 908(2):172, 2021. doi: 10.3847/1538-4357/abc000.
- L. Spitzer. *Physics of Fully Ionized Gases*. Interscience Publishers, New York, 2nd edition, 1962.
- Eleuterio F. Toro. *Riemann Solvers and Numerical Methods for Fluid Dynamics: A Practical Introduction*. Springer, Berlin, Heidelberg, 3 edition, 2009. doi: 10.1007/b79761. Chapter 14, “The MUSCL-Hancock Method”.
- B. van der Holst, IV Manchester, W. B., R. A. Frazin, A. M. Vásquez, G. Tóth, and T. I. Gombosi. A data-driven, two-temperature solar wind model with Alfvén waves. *The Astrophysical Journal*, 725(1):1373–1383, 2010. doi: 10.1088/0004-637X/725/1/1373.
- Bram van Leer. Towards the ultimate conservative difference scheme. V. A second-order sequel to Godunov’s method. *Journal of Computational Physics*, 32(1):101–136, 1979. doi: 10.1016/0021-9991(79)90145-1.
- G. S. Voronov. A practical fit formula for ionization rate coefficients of atoms and ions by electron impact: $Z=1-28$. *Atomic Data and Nuclear Data Tables*, 65(1):1–35, 1997. doi: 10.1006/adnd.1997.0732.
- J. Vranjes and P. S. Krstic. Collisions, magnetization, and transport coefficients in the lower solar atmosphere. *Astronomy and Astrophysics*, 554:A22, 2013. doi: 10.1051/0004-6361/201220738.